



GDB Internals

A guide to the internals of the GNU debugger

John Gilmore
Cygnus Solutions
Second Edition:
Stan Shebs
Cygnus Solutions

Cygnus Solutions
Revision: 1.142
T_EXinfo 2000-05-28.15



Copyright © 1990,1991,1992,1993,1994,1996,1998,1999,2000,2001 Free Software Foundation, Inc.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.1 or any later version published by the Free Software Foundation; with no Invariant Sections, with the Front-Cover Texts being “A GNU Manual,” and with the Back-Cover Texts as in (a) below.

(a) The FSF’s Back-Cover Text is: “You have freedom to copy and modify this GNU Manual, like GNU software. Copies published by the Free Software Foundation raise funds for GNU development.”

Table of Contents

Scope of this Document	1
1 Requirements	1
2 Overall Structure	1
2.1 The Symbol Side	2
2.2 The Target Side	2
2.3 Configurations	2
3 Algorithms	2
3.1 Frames	3
3.2 Breakpoint Handling	3
3.3 Single Stepping	4
3.4 Signal Handling	4
3.5 Thread Handling	4
3.6 Inferior Function Calls	4
3.7 Longjmp Support	4
3.8 Watchpoints	4
3.8.1 x86 Watchpoints	7
4 User Interface	9
4.1 Command Interpreter	9
4.2 UI-Independent Output—the <code>ui_out</code> Functions	10
4.2.1 Overview and Terminology	10
4.2.2 General Conventions	11
4.2.3 Table, Tuple and List Functions	11
4.2.4 Item Output Functions	13
4.2.5 Utility Output Functions	14
4.2.6 Examples of Use of <code>ui_out</code> functions	15
4.3 Console Printing	19
4.4 TUI	19
5 libgdb	19
5.1 libgdb 1.0	19
5.2 libgdb 2.0	19
5.3 The libgdb Model	20
5.4 CLI support	20
5.5 libgdb components	20

6	Symbol Handling	21
6.1	Symbol Reading	21
6.2	Partial Symbol Tables	22
6.3	Types	23
	Fundamental Types (e.g., FT_VOID, FT_BOOLEAN)	23
	Type Codes (e.g., TYPE_CODE_PTR, TYPE_CODE_ARRAY)	24
	Builtin Types (e.g., builtin_type_void, builtin_type_char)	24
6.4	Object File Formats	24
	6.4.1 a.out	24
	6.4.2 COFF	24
	6.4.3 ECOFF	24
	6.4.4 XCOFF	25
	6.4.5 PE	25
	6.4.6 ELF	25
	6.4.7 SOM	25
	6.4.8 Other File Formats	25
6.5	Debugging File Formats	25
	6.5.1 stabs	25
	6.5.2 COFF	26
	6.5.3 Mips debug (Third Eye)	26
	6.5.4 DWARF 1	26
	6.5.5 DWARF 2	26
	6.5.6 SOM	26
6.6	Adding a New Symbol Reader to GDB	26
7	Language Support	27
7.1	Adding a Source Language to GDB	27
8	Host Definition	29
8.1	Adding a New Host	29
8.2	Host Conditionals	30
9	Target Architecture Definition	33
9.1	Registers and Memory	33
9.2	Pointers Are Not Always Addresses	34
9.3	Using Different Register and Memory Data Representations	36
9.4	Frame Interpretation	37
9.5	Inferior Call Setup	37
9.6	Compiler Characteristics	38
9.7	Target Conditionals	38
9.8	Adding a New Target	51

10	Target Vector Definition	52
10.1	File Targets.....	52
10.2	Standard Protocol and Remote Stubs.....	52
10.3	ROM Monitor Interface.....	53
10.4	Custom Protocols	53
10.5	Transport Layer.....	53
10.6	Builtin Simulator	53
11	Native Debugging	53
11.1	Native core file Support.....	54
11.2	ptrace.....	55
11.3	/proc	55
11.4	win32	55
11.5	shared libraries	55
11.6	Native Conditionals	55
12	Support Libraries	58
12.1	BFD	58
12.2	opcodes	58
12.3	readline	58
12.4	mmalloc.....	58
12.5	libiberty	58
12.6	gnu-regex.....	59
12.7	include	59
13	Coding.....	59
13.1	Cleanups	59
13.2	Wrapping Output Lines	60
13.3	GDB Coding Standards	60
13.3.1	ISO-C	60
13.3.2	Memory Management	60
13.3.3	Compiler Warnings	61
13.3.4	Formatting.....	62
13.3.5	Comments	62
13.3.6	C Usage	63
13.3.7	Function Prototypes	63
13.3.8	Internal Error Recovery	63
13.3.9	File Names.....	64
13.3.10	Include Files	64
13.3.11	Clean Design and Portable Implementation....	64
14	Porting GDB.....	66
14.1	Configuring GDB for Release.....	67

15	Testsuite	67
15.1	Using the Testsuite	68
15.2	Testsuite Organization	68
15.3	Writing Tests	69
16	Hints	70
16.1	Getting Started	70
16.2	Debugging GDB with itself	71
16.3	Submitting Patches	71
16.4	Obsolete Conditionals	72
	Index	73

Scope of this Document

This document documents the internals of the GNU debugger, GDB. It includes description of GDB's key algorithms and operations, as well as the mechanisms that adapt GDB to specific hosts and targets.

1 Requirements

Before diving into the internals, you should understand the formal requirements and other expectations for GDB. Although some of these may seem obvious, there have been proposals for GDB that have run counter to these requirements.

First of all, GDB is a debugger. It's not designed to be a front panel for embedded systems. It's not a text editor. It's not a shell. It's not a programming environment.

GDB is an interactive tool. Although a batch mode is available, GDB's primary role is to interact with a human programmer.

GDB should be responsive to the user. A programmer hot on the trail of a nasty bug, and operating under a looming deadline, is going to be very impatient of everything, including the response time to debugger commands.

GDB should be relatively permissive, such as for expressions. While the compiler should be picky (or have the option to be made picky), since source code lives for a long time usually, the programmer doing debugging shouldn't be spending time figuring out to mollify the debugger.

GDB will be called upon to deal with really large programs. Executable sizes of 50 to 100 megabytes occur regularly, and we've heard reports of programs approaching 1 gigabyte in size.

GDB should be able to run everywhere. No other debugger is available for even half as many configurations as GDB supports.

2 Overall Structure

GDB consists of three major subsystems: user interface, symbol handling (the *symbol side*), and target system handling (the *target side*).

The user interface consists of several actual interfaces, plus supporting code.

The symbol side consists of object file readers, debugging info interpreters, symbol table management, source language expression parsing, type and value printing.

The target side consists of execution control, stack frame analysis, and physical target manipulation.

The target side/symbol side division is not formal, and there are a number of exceptions. For instance, core file support involves symbolic elements (the basic core file reader is in BFD) and target elements (it supplies the contents of memory and the values of registers). Instead, this division is useful for understanding how the minor subsystems should fit together.

2.1 The Symbol Side

The symbolic side of GDB can be thought of as “everything you can do in GDB without having a live program running”. For instance, you can look at the types of variables, and evaluate many kinds of expressions.

2.2 The Target Side

The target side of GDB is the “bits and bytes manipulator”. Although it may make reference to symbolic info here and there, most of the target side will run with only a stripped executable available—or even no executable at all, in remote debugging cases.

Operations such as disassembly, stack frame crawls, and register display, are able to work with no symbolic info at all. In some cases, such as disassembly, GDB will use symbolic info to present addresses relative to symbols rather than as raw numbers, but it will work either way.

2.3 Configurations

Host refers to attributes of the system where GDB runs. *Target* refers to the system where the program being debugged executes. In most cases they are the same machine, in which case a third type of *Native* attributes come into play.

Defines and include files needed to build on the host are host support. Examples are `tty` support, system defined types, host byte order, host float format.

Defines and information needed to handle the target format are target dependent. Examples are the stack frame format, instruction set, breakpoint instruction, registers, and how to set up and tear down the stack to call a function.

Information that is only needed when the host and target are the same, is native dependent. One example is Unix child process support; if the host and target are not the same, doing a fork to start the target process is a bad idea. The various macros needed for finding the registers in the `upage`, running `ptrace`, and such are all in the native-dependent files.

Another example of native-dependent code is support for features that are really part of the target environment, but which require `#include` files that are only available on the host system. Core file handling and `setjmp` handling are two common cases.

When you want to make GDB work “native” on a particular machine, you have to include all three kinds of information.

3 Algorithms

GDB uses a number of debugging-specific algorithms. They are often not very complicated, but get lost in the thicket of special cases and real-world issues. This chapter describes the basic algorithms and mentions some of the specific target definitions that they use.

3.1 Frames

A frame is a construct that GDB uses to keep track of calling and called functions.

`FRAME_FP` in the machine description has no meaning to the machine-independent part of GDB, except that it is used when setting up a new frame from scratch, as follows:

```
create_new_frame (read_register (FP_REGNUM), read_pc ());
```

Other than that, all the meaning imparted to `FP_REGNUM` is imparted by the machine-dependent code. So, `FP_REGNUM` can have any value that is convenient for the code that creates new frames. (`create_new_frame` calls `INIT_EXTRA_FRAME_INFO` if it is defined; that is where you should use the `FP_REGNUM` value, if your frames are nonstandard.)

Given a GDB frame, define `FRAME_CHAIN` to determine the address of the calling function's frame. This will be used to create a new GDB frame struct, and then `INIT_EXTRA_FRAME_INFO` and `INIT_FRAME_PC` will be called for the new frame.

3.2 Breakpoint Handling

In general, a breakpoint is a user-designated location in the program where the user wants to regain control if program execution ever reaches that location.

There are two main ways to implement breakpoints; either as “hardware” breakpoints or as “software” breakpoints.

Hardware breakpoints are sometimes available as a builtin debugging features with some chips. Typically these work by having dedicated register into which the breakpoint address may be stored. If the PC (shorthand for *program counter*) ever matches a value in a breakpoint registers, the CPU raises an exception and reports it to GDB.

Another possibility is when an emulator is in use; many emulators include circuitry that watches the address lines coming out from the processor, and force it to stop if the address matches a breakpoint's address.

A third possibility is that the target already has the ability to do breakpoints somehow; for instance, a ROM monitor may do its own software breakpoints. So although these are not literally “hardware breakpoints”, from GDB's point of view they work the same; GDB need not do nothing more than set the breakpoint and wait for something to happen.

Since they depend on hardware resources, hardware breakpoints may be limited in number; when the user asks for more, GDB will start trying to set software breakpoints. (On some architectures, notably the 32-bit x86 platforms, GDB cannot always know whether there's enough hardware resources to insert all the hardware breakpoints and watchpoints. On those platforms, GDB prints an error message only when the program being debugged is continued.)

Software breakpoints require GDB to do somewhat more work. The basic theory is that GDB will replace a program instruction with a trap, illegal divide, or some other instruction that will cause an exception, and then when it's encountered, GDB will take the exception and stop the program. When the user says to continue, GDB will restore the original instruction, single-step, re-insert the trap, and continue on.

Since it literally overwrites the program being tested, the program area must be writable, so this technique won't work on programs in ROM. It can also distort the behavior of programs that examine themselves, although such a situation would be highly unusual.

Also, the software breakpoint instruction should be the smallest size of instruction, so it doesn't overwrite an instruction that might be a jump target, and cause disaster when the program jumps into the middle of the breakpoint instruction. (Strictly speaking, the breakpoint must be no larger than the smallest interval between instructions that may be jump targets; perhaps there is an architecture where only even-numbered instructions may be jumped to.) Note that it's possible for an instruction set not to have any instructions usable for a software breakpoint, although in practice only the ARC has failed to define such an instruction.

The basic definition of the software breakpoint is the macro `BREAKPOINT`.

Basic breakpoint object handling is in `'breakpoint.c'`. However, much of the interesting breakpoint action is in `'infrun.c'`.

3.3 Single Stepping

3.4 Signal Handling

3.5 Thread Handling

3.6 Inferior Function Calls

3.7 Longjmp Support

GDB has support for figuring out that the target is doing a `longjmp` and for stopping at the target of the jump, if we are stepping. This is done with a few specialized internal breakpoints, which are visible in the output of the `'maint info breakpoint'` command.

To make this work, you need to define a macro called `GET_LONGJMP_TARGET`, which will examine the `jmp_buf` structure and extract the `longjmp` target address. Since `jmp_buf` is target specific, you will need to define it in the appropriate `'tm-target.h'` file. Look in `'tm-sun4os4.h'` and `'sparc-tdep.c'` for examples of how to do this.

3.8 Watchpoints

Watchpoints are a special kind of breakpoints (see [Chapter 3 \[Algorithms\], page 2](#)) which break when data is accessed rather than when some instruction is executed. When you have data which changes without your knowing what code does that, watchpoints are the silver bullet to hunt down and kill such bugs.

Watchpoints can be either hardware-assisted or not; the latter type is known as "software watchpoints." GDB always uses hardware-assisted watchpoints if they are available, and falls back on software watchpoints otherwise. Typical situations where GDB will use software watchpoints are:

- The watched memory region is too large for the underlying hardware watchpoint support. For example, each x86 debug register can watch up to 4 bytes of memory, so trying to watch data structures whose size is more than 16 bytes will cause GDB to use software watchpoints.
- The value of the expression to be watched depends on data held in registers (as opposed to memory).
- Too many different watchpoints requested. (On some architectures, this situation is impossible to detect until the debugged program is resumed.) Note that x86 debug registers are used both for hardware breakpoints and for watchpoints, so setting too many hardware breakpoints might cause watchpoint insertion to fail.
- No hardware-assisted watchpoints provided by the target implementation.

Software watchpoints are very slow, since GDB needs to single-step the program being debugged and test the value of the watched expression(s) after each instruction. The rest of this section is mostly irrelevant for software watchpoints.

GDB uses several macros and primitives to support hardware watchpoints:

TARGET_HAS_HARDWARE_WATCHPOINTS

If defined, the target supports hardware watchpoints.

TARGET_CAN_USE_HARDWARE_WATCHPOINT (*type*, *count*, *other*)

Return the number of hardware watchpoints of type *type* that are possible to be set. The value is positive if *count* watchpoints of this type can be set, zero if setting watchpoints of this type is not supported, and negative if *count* is more than the maximum number of watchpoints of type *type* that can be set. *other* is non-zero if other types of watchpoints are currently enabled (there are architectures which cannot set watchpoints of different types at the same time).

TARGET_REGION_OK_FOR_HW_WATCHPOINT (*addr*, *len*)

Return non-zero if hardware watchpoints can be used to watch a region whose address is *addr* and whose length in bytes is *len*.

TARGET_REGION_SIZE_OK_FOR_HW_WATCHPOINT (*size*)

Return non-zero if hardware watchpoints can be used to watch a region whose size is *size*. GDB only uses this macro as a fall-back, in case TARGET_REGION_OK_FOR_HW_WATCHPOINT is not defined.

TARGET_DISABLE_HW_WATCHPOINTS (*pid*)

Disables watchpoints in the process identified by *pid*. This is used, e.g., on HP-UX which provides operations to disable and enable the page-level memory protection that implements hardware watchpoints on that platform.

TARGET_ENABLE_HW_WATCHPOINTS (*pid*)

Enables watchpoints in the process identified by *pid*. This is used, e.g., on HP-UX which provides operations to disable and enable the page-level memory protection that implements hardware watchpoints on that platform.

TARGET_RANGE_PROFITABLE_FOR_HW_WATCHPOINT (*pid*, *start*, *len*)

Some addresses may not be profitable to use hardware to watch, or may be difficult to understand when the addressed object is out of scope, and hence should

not be watched with hardware watchpoints. On some targets, this may have severe performance penalties, such that we might as well use regular watchpoints, and save (possibly precious) hardware watchpoints for other locations.

`target_insert_watchpoint (addr, len, type)`

`target_remove_watchpoint (addr, len, type)`

Insert or remove a hardware watchpoint starting at *addr*, for *len* bytes. *type* is the watchpoint type, one of the possible values of the enumerated data type `target_hw_bp_type`, defined by ‘`breakpoint.h`’ as follows:

```
enum target_hw_bp_type
{
    hw_write   = 0, /* Common (write) HW watchpoint */
    hw_read    = 1, /* Read      HW watchpoint */
    hw_access  = 2, /* Access (read or write) HW watchpoint */
    hw_execute = 3  /* Execute HW breakpoint */
};
```

These two macros should return 0 for success, non-zero for failure.

`target_remove_hw_breakpoint (addr, shadow)`

`target_insert_hw_breakpoint (addr, shadow)`

Insert or remove a hardware-assisted breakpoint at address *addr*. Returns zero for success, non-zero for failure. *shadow* is the real contents of the byte where the breakpoint has been inserted; it is generally not valid when hardware breakpoints are used, but since no other code touches these values, the implementations of the above two macros can use them for their internal purposes.

`target_stopped_data_address ()`

If the inferior has some watchpoint that triggered, return the address associated with that watchpoint. Otherwise, return zero.

`DECR_PC_AFTER_HW_BREAK`

If defined, GDB decrements the program counter by the value of `DECR_PC_AFTER_HW_BREAK` after a hardware break-point. This overrides the value of `DECR_PC_AFTER_BREAK` when a breakpoint that breaks is a hardware-assisted breakpoint.

`HAVE_STEPPABLE_WATCHPOINT`

If defined to a non-zero value, it is not necessary to disable a watchpoint to step over it.

`HAVE_NONSTEPPABLE_WATCHPOINT`

If defined to a non-zero value, GDB should disable a watchpoint to step the inferior over it.

`HAVE_CONTINUABLE_WATCHPOINT`

If defined to a non-zero value, it is possible to continue the inferior after a watchpoint has been hit.

`CANNOT_STEP_HW_WATCHPOINTS`

If this is defined to a non-zero value, GDB will remove all watchpoints before stepping the inferior.

STOPPED_BY_WATCHPOINT (*wait_status*)

Return non-zero if stopped by a watchpoint. *wait_status* is of the type `struct target_waitstatus`, defined by ‘`target.h`’.

3.8.1 x86 Watchpoints

The 32-bit Intel x86 (a.k.a. ia32) processors feature special debug registers designed to facilitate debugging. GDB provides a generic library of functions that x86-based ports can use to implement support for watchpoints and hardware-assisted breakpoints. This subsection documents the x86 watchpoint facilities in GDB.

To use the generic x86 watchpoint support, a port should do the following:

- Define the macro `I386_USE_GENERIC_WATCHPOINTS` somewhere in the target-dependent headers.
- Include the ‘`config/i386/nm-i386.h`’ header file *after* defining `I386_USE_GENERIC_WATCHPOINTS`.
- Add ‘`i386-nat.o`’ to the value of the Make variable `NATDEPFILES` (see [Chapter 11 \[Native Debugging\]](#), page 53) or `TDEPFILES` (see [Chapter 9 \[Target Architecture Definition\]](#), page 33).
- Provide implementations for the `I386_DR_LOW_*` macros described below. Typically, each macro should call a target-specific function which does the real work.

The x86 watchpoint support works by maintaining mirror images of the debug registers. Values are copied between the mirror images and the real debug registers via a set of macros which each target needs to provide:

`I386_DR_LOW_SET_CONTROL` (*val*)

Set the Debug Control (DR7) register to the value *val*.

`I386_DR_LOW_SET_ADDR` (*idx*, *addr*)

Put the address *addr* into the debug register number *idx*.

`I386_DR_LOW_RESET_ADDR` (*idx*)

Reset (i.e. zero out) the address stored in the debug register number *idx*.

`I386_DR_LOW_GET_STATUS`

Return the value of the Debug Status (DR6) register. This value is used immediately after it is returned by `I386_DR_LOW_GET_STATUS`, so as to support per-thread status register values.

For each one of the 4 debug registers (whose indices are from 0 to 3) that store addresses, a reference count is maintained by GDB, to allow sharing of debug registers by several watchpoints. This allows users to define several watchpoints that watch the same expression, but with different conditions and/or commands, without wasting debug registers which are in short supply. GDB maintains the reference counts internally, targets don’t have to do anything to use this feature.

The x86 debug registers can each watch a region that is 1, 2, or 4 bytes long. The ia32 architecture requires that each watched region be appropriately aligned: 2-byte region on 2-byte boundary, 4-byte region on 4-byte boundary. However, the x86 watchpoint support

in GDB can watch unaligned regions and regions larger than 4 bytes (up to 16 bytes) by allocating several debug registers to watch a single region. This allocation of several registers per a watched region is also done automatically without target code intervention.

The generic x86 watchpoint support provides the following API for the GDB's application code:

`i386_region_ok_for_watchpoint (addr, len)`

The macro `TARGET_REGION_OK_FOR_HW_WATCHPOINT` is set to call this function. It counts the number of debug registers required to watch a given region, and returns a non-zero value if that number is less than 4, the number of debug registers available to x86 processors.

`i386_stopped_data_address (void)`

The macros `STOPPED_BY_WATCHPOINT` and `target_stopped_data_address` are set to call this function. The argument passed to `STOPPED_BY_WATCHPOINT` is ignored. This function examines the breakpoint condition bits in the DR6 Debug Status register, as returned by the `I386_DR_LOW_GET_STATUS` macro, and returns the address associated with the first bit that is set in DR6.

`i386_insert_watchpoint (addr, len, type)`

`i386_remove_watchpoint (addr, len, type)`

Insert or remove a watchpoint. The macros `target_insert_watchpoint` and `target_remove_watchpoint` are set to call these functions. `i386_insert_watchpoint` first looks for a debug register which is already set to watch the same region for the same access types; if found, it just increments the reference count of that debug register, thus implementing debug register sharing between watchpoints. If no such register is found, the function looks for a vacant debug register, sets its mirrored value to `addr`, sets the mirrored value of DR7 Debug Control register as appropriate for the `len` and `type` parameters, and then passes the new values of the debug register and DR7 to the inferior by calling `I386_DR_LOW_SET_ADDR` and `I386_DR_LOW_SET_CONTROL`. If more than one debug register is required to cover the given region, the above process is repeated for each debug register.

`i386_remove_watchpoint` does the opposite: it resets the address in the mirrored value of the debug register and its read/write and length bits in the mirrored value of DR7, then passes these new values to the inferior via `I386_DR_LOW_RESET_ADDR` and `I386_DR_LOW_SET_CONTROL`. If a register is shared by several watchpoints, each time a `i386_remove_watchpoint` is called, it decrements the reference count, and only calls `I386_DR_LOW_RESET_ADDR` and `I386_DR_LOW_SET_CONTROL` when the count goes to zero.

`i386_insert_hw_breakpoint (addr, shadow)`

`i386_remove_hw_breakpoint (addr, shadow)`

These functions insert and remove hardware-assisted breakpoints. The macros `target_insert_hw_breakpoint` and `target_remove_hw_breakpoint` are set to call these functions. These functions work like `i386_insert_watchpoint` and `i386_remove_watchpoint`, respectively, except that they set up the debug registers to watch instruction execution, and each hardware-assisted breakpoint always requires exactly one debug register.

i386_stopped_by_hwbp (void)

This function returns non-zero if the inferior has some watchpoint or hardware breakpoint that triggered. It works like `i386_stopped_data_address`, except that it doesn't return the address whose watchpoint triggered.

i386_cleanup_dregs (void)

This function clears all the reference counts, addresses, and control bits in the mirror images of the debug registers. It doesn't affect the actual debug registers in the inferior process.

Notes:

1. x86 processors support setting watchpoints on I/O reads or writes. However, since no target supports this (as of March 2001), and since `enum target_hw_bp_type` doesn't even have an enumeration for I/O watchpoints, this feature is not yet available to GDB running on x86.
2. x86 processors can enable watchpoints locally, for the current task only, or globally, for all the tasks. For each debug register, there's a bit in the DR7 Debug Control register that determines whether the associated address is watched locally or globally. The current implementation of x86 watchpoint support in GDB always sets watchpoints to be locally enabled, since global watchpoints might interfere with the underlying OS and are probably unavailable in many platforms.

4 User Interface

GDB has several user interfaces. Although the command-line interface is the most common and most familiar, there are others.

4.1 Command Interpreter

The command interpreter in GDB is fairly simple. It is designed to allow for the set of commands to be augmented dynamically, and also has a recursive subcommand capability, where the first argument to a command may itself direct a lookup on a different command list.

For instance, the `'set'` command just starts a lookup on the `setlist` command list, while `'set thread'` recurses to the `set_thread_cmd_list`.

To add commands in general, use `add_cmd`. `add_com` adds to the main command list, and should be used for those commands. The usual place to add commands is in the `_initialize_xyz` routines at the ends of most source files.

Before removing commands from the command set it is a good idea to deprecate them for some time. Use `deprecate_cmd` on commands or aliases to set the deprecated flag. `deprecate_cmd` takes a `struct cmd_list_element` as its first argument. You can use the return value from `add_com` or `add_cmd` to deprecate the command immediately after it is created.

The first time a command is used the user will be warned and offered a replacement (if one exists). Note that the replacement string passed to `deprecate_cmd` should be the full name of the command, i.e. the entire string the user should type at the command line.

4.2 UI-Independent Output—the `ui_out` Functions

The `ui_out` functions present an abstraction level for the GDB output code. They hide the specifics of different user interfaces supported by GDB, and thus free the programmer from the need to write several versions of the same code, one each for every UI, to produce output.

4.2.1 Overview and Terminology

In general, execution of each GDB command produces some sort of output, and can even generate an input request.

Output can be generated for the following purposes:

- to display a *result* of an operation;
- to convey *info* or produce side-effects of a requested operation;
- to provide a *notification* of an asynchronous event (including progress indication of a prolonged asynchronous operation);
- to display *error messages* (including warnings);
- to show *debug data*;
- to *query* or prompt a user for input (a special case).

This section mainly concentrates on how to build result output, although some of it also applies to other kinds of output.

Generation of output that displays the results of an operation involves one or more of the following:

- output of the actual data
- formatting the output as appropriate for console output, to make it easily readable by humans
- machine oriented formatting—a more terse formatting to allow for easy parsing by programs which read GDB's output
- annotation, whose purpose is to help legacy GUIs to identify interesting parts in the output

The `ui_out` routines take care of the first three aspects. Annotations are provided by separate annotation routines. Note that use of annotations for an interface between a GUI and GDB is deprecated.

Output can be in the form of a single item, which we call a *field*; a *list* consisting of identical fields; a *tuple* consisting of non-identical fields; or a *table*, which is a tuple consisting of a header and a body. In a BNF-like form:

```

<table> ⇨
    <header> <body>
<header> ⇨
    { <column> }
<column> ⇨
    <width> <alignment> <title>
<body> ⇨ {<row>}

```


4.2.2 General Conventions

Most `ui_out` routines are of type `void`, the exceptions are `ui_out_stream_new` (which returns a pointer to the newly created object) and the `make_cleanup` routines.

The first parameter is always the `ui_out` vector object, a pointer to a `struct ui_out`.

The *format* parameter is like in `printf` family of functions. When it is present, there must also be a variable list of arguments sufficient used to satisfy the `%` specifiers in the supplied format.

When a character string argument is not used in a `ui_out` function call, a `NULL` pointer has to be supplied instead.

4.2.3 Table, Tuple and List Functions

This section introduces `ui_out` routines for building lists, tuples and tables. The routines to output the actual data items (fields) are presented in the next section.

To recap: A *tuple* is a sequence of *fields*, each field containing information about an object; a *list* is a sequence of fields where each field describes an identical object.

Use the *table* functions when your output consists of a list of rows (tuples) and the console output should include a heading. Use this even when you are listing just one object but you still want the header.

Tables can not be nested. Tuples and lists can be nested up to a maximum of five levels.

The overall structure of the table output code is something like this:

```
ui_out_table_begin
  ui_out_table_header
  ...
  ui_out_table_body
    ui_out_tuple_begin
      ui_out_field_*
      ...
    ui_out_tuple_end
  ...
ui_out_table_end
```

Here is the description of table-, tuple- and list-related `ui_out` functions:

void <code>ui_out_table_begin</code> (<code>struct ui_out *uiout</code> , <code>int nbrofcols</code> , <code>int nr_rows</code> , <code>const char *tblid</code>)	Function
--	----------

The function `ui_out_table_begin` marks the beginning of the output of a table. It should always be called before any other `ui_out` function for a given table. *nbrofcols* is the number of columns in the table. *nr_rows* is the number of rows in the table. *tblid* is an optional string identifying the table. The string pointed to by *tblid* is copied by the implementation of `ui_out_table_begin`, so the application can free the string if it was `malloced`.

The companion function `ui_out_table_end`, described below, marks the end of the table's output.

void ui_out_table_header (struct ui_out *uiout, int width, enum ui_align alignment, const char *colhdr) Function

`ui_out_table_header` provides the header information for a single table column. You call this function several times, one each for every column of the table, after `ui_out_table_begin`, but before `ui_out_table_body`.

The value of *width* gives the column width in characters. The value of *alignment* is one of `left`, `center`, and `right`, and it specifies how to align the header: left-justify, center, or right-justify it. *colhdr* points to a string that specifies the column header; the implementation copies that string, so column header strings in `malloced` storage can be freed after the call.

void ui_out_table_body (struct ui_out *uiout) Function

This function delimits the table header from the table body.

void ui_out_table_end (struct ui_out *uiout) Function

This function signals the end of a table's output. It should be called after the table body has been produced by the list and field output functions.

There should be exactly one call to `ui_out_table_end` for each call to `ui_out_table_begin`, otherwise the `ui_out` functions will signal an internal error.

The output of the tuples that represent the table rows must follow the call to `ui_out_table_body` and precede the call to `ui_out_table_end`. You build a tuple by calling `ui_out_tuple_begin` and `ui_out_tuple_end`, with suitable calls to functions which actually output fields between them.

void ui_out_tuple_begin (struct ui_out *uiout, const char *id) Function

This function marks the beginning of a tuple output. *id* points to an optional string that identifies the tuple; it is copied by the implementation, and so strings in `malloced` storage can be freed after the call.

void ui_out_tuple_end (struct ui_out *uiout) Function

This function signals an end of a tuple output. There should be exactly one call to `ui_out_tuple_end` for each call to `ui_out_tuple_begin`, otherwise an internal GDB error will be signaled.

struct cleanup *make_cleanup_ui_out_tuple_begin_end (struct ui_out *uiout, const char *id) Function

This function first opens the tuple and then establishes a cleanup (see [Chapter 13 \[Coding\]](#), page 59) to close the tuple. It provides a convenient and correct implementation of the non-portable¹ code sequence:

```
struct cleanup *old_cleanup;
ui_out_tuple_begin (uiout, "...");
old_cleanup = make_cleanup ((void(*) (void *)) ui_out_tuple_end,
                             uiout);
```

¹ The function cast is not portable ISO-C.

void ui_out_list_begin (struct ui_out *uiout, const char *id) Function
 This function marks the beginning of a list output. *id* points to an optional string that identifies the list; it is copied by the implementation, and so strings in **malloced** storage can be freed after the call.

void ui_out_list_end (struct ui_out *uiout) Function
 This function signals an end of a list output. There should be exactly one call to **ui_out_list_end** for each call to **ui_out_list_begin**, otherwise an internal GDB error will be signaled.

struct cleanup *make_cleanup_ui_out_list_begin_end (struct ui_out *uiout, const char *id) Function
 Similar to **make_cleanup_ui_out_tuple_begin_end**, this function opens a list and then establishes cleanup (see [Chapter 13 \[Coding\], page 59](#)) that will close the list.

4.2.4 Item Output Functions

The functions described below produce output for the actual data items, or fields, which contain information about the object.

Choose the appropriate function accordingly to your particular needs.

void ui_out_field_fmt (struct ui_out *uiout, char *fldname, char *format, ...) Function
 This is the most general output function. It produces the representation of the data in the variable-length argument list according to formatting specifications in *format*, a **printf**-like format string. The optional argument *fldname* supplies the name of the field. The data items themselves are supplied as additional arguments after *format*. This generic function should be used only when it is not possible to use one of the specialized versions (see below).

void ui_out_field_int (struct ui_out *uiout, const char *fldname, int value) Function
 This function outputs a value of an **int** variable. It uses the **"%d"** output conversion specification. *fldname* specifies the name of the field.

void ui_out_field_core_addr (struct ui_out *uiout, const char *fldname, CORE_ADDR address) Function
 This function outputs an address.

void ui_out_field_string (struct ui_out *uiout, const char *fldname, const char *string) Function
 This function outputs a string using the **"%s"** conversion specification.

Sometimes, there's a need to compose your output piece by piece using functions that operate on a stream, such as **value_print** or **fprintf_symbol_filtered**. These functions accept an argument of the type **struct ui_file ***, a pointer to a **ui_file** object used to

store the data stream used for the output. When you use one of these functions, you need a way to pass their results stored in a `ui_file` object to the `ui_out` functions. To this end, you first create a `ui_stream` object by calling `ui_out_stream_new`, pass the `stream` member of that `ui_stream` object to `value_print` and similar functions, and finally call `ui_out_field_stream` to output the field you constructed. When the `ui_stream` object is no longer needed, you should destroy it and free its memory by calling `ui_out_stream_delete`.

struct ui_stream *ui_out_stream_new (struct ui_out *uiout) Function

This function creates a new `ui_stream` object which uses the same output methods as the `ui_out` object whose pointer is passed in *uiout*. It returns a pointer to the newly created `ui_stream` object.

void ui_out_stream_delete (struct ui_stream *streambuf) Function

This functions destroys a `ui_stream` object specified by *streambuf*.

void ui_out_field_stream (struct ui_out *uiout, const char *fieldname, struct ui_stream *streambuf) Function

This function consumes all the data accumulated in `streambuf->stream` and outputs it like `ui_out_field_string` does. After a call to `ui_out_field_stream`, the accumulated data no longer exists, but the stream is still valid and may be used for producing more fields.

Important: If there is any chance that your code could bail out before completing output generation and reaching the point where `ui_out_stream_delete` is called, it is necessary to set up a cleanup, to avoid leaking memory and other resources. Here's a skeleton code to do that:

```
struct ui_stream *mybuf = ui_out_stream_new (uiout);
struct cleanup *old = make_cleanup (ui_out_stream_delete, mybuf);
...
do_cleanups (old);
```

If the function already has the old cleanup chain set (for other kinds of cleanups), you just have to add your cleanup to it:

```
mybuf = ui_out_stream_new (uiout);
make_cleanup (ui_out_stream_delete, mybuf);
```

Note that with cleanups in place, you should not call `ui_out_stream_delete` directly, or you would attempt to free the same buffer twice.

4.2.5 Utility Output Functions

void ui_out_field_skip (struct ui_out *uiout, const char *fldname) Function

This function skips a field in a table. Use it if you have to leave an empty field without disrupting the table alignment. The argument *fldname* specifies a name for the (missing) field.

void ui_out_text (struct ui_out *uiout, const char *string) Function

This function outputs the text in *string* in a way that makes it easy to be read by humans. For example, the console implementation of this method filters the text through a built-in pager, to prevent it from scrolling off the visible portion of the screen.

Use this function for printing relatively long chunks of text around the actual field data: the text it produces is not aligned according to the table's format. Use **ui_out_field_string** to output a string field, and use **ui_out_message**, described below, to output short messages.

void ui_out_spaces (struct ui_out *uiout, int nspaces) Function

This function outputs *nspaces* spaces. It is handy to align the text produced by **ui_out_text** with the rest of the table or list.

void ui_out_message (struct ui_out *uiout, int verbosity, const char *format, ...) Function

This function produces a formatted message, provided that the current verbosity level is at least as large as given by *verbosity*. The current verbosity level is specified by the user with the 'set verbositylevel' command.²

void ui_out_wrap_hint (struct ui_out *uiout, char *indent) Function

This function gives the console output filter (a paging filter) a hint of where to break lines which are too long. Ignored for all other output consumers. *indent*, if non-NULL, is the string to be printed to indent the wrapped text on the next line; it must remain accessible until the next call to **ui_out_wrap_hint**, or until an explicit newline is produced by one of the other functions. If *indent* is NULL, the wrapped text will not be indented.

void ui_out_flush (struct ui_out *uiout) Function

This function flushes whatever output has been accumulated so far, if the UI buffers output.

4.2.6 Examples of Use of ui_out functions

This section gives some practical examples of using the **ui_out** functions to generalize the old console-oriented code in GDB. The examples all come from functions defined on the 'breakpoints.c' file.

This example, from the **breakpoint_1** function, shows how to produce a table.

The original code was:

```
if (!found_a_breakpoint++)
{
    annotate_breakpoints_headers ();
```

² As of this writing (April 2001), setting verbosity level is not yet implemented, and is always returned as zero. So calling **ui_out_message** with a *verbosity* argument more than zero will cause the message to never be printed.

```

    annotate_field (0);
    printf_filtered ("Num ");
    annotate_field (1);
    printf_filtered ("Type          ");
    annotate_field (2);
    printf_filtered ("Disp ");
    annotate_field (3);
    printf_filtered ("Enb ");
    if (addressprint)
    {
        annotate_field (4);
        printf_filtered ("Address    ");
    }
    annotate_field (5);
    printf_filtered ("What\n");

    annotate_breakpoints_table ();
}

```

Here's the new version:

```

nr_printable_breakpoints = ...;

if (addressprint)
    ui_out_table_begin (ui, 6, nr_printable_breakpoints, "BreakpointTable");
else
    ui_out_table_begin (ui, 5, nr_printable_breakpoints, "BreakpointTable");

if (nr_printable_breakpoints > 0)
    annotate_breakpoints_headers ();
if (nr_printable_breakpoints > 0)
    annotate_field (0);
ui_out_table_header (uiout, 3, ui_left, "number", "Num"); /* 1 */
if (nr_printable_breakpoints > 0)
    annotate_field (1);
ui_out_table_header (uiout, 14, ui_left, "type", "Type"); /* 2 */
if (nr_printable_breakpoints > 0)
    annotate_field (2);
ui_out_table_header (uiout, 4, ui_left, "disp", "Disp"); /* 3 */
if (nr_printable_breakpoints > 0)
    annotate_field (3);
ui_out_table_header (uiout, 3, ui_left, "enabled", "Enb"); /* 4 */
if (addressprint)
{
    if (nr_printable_breakpoints > 0)
        annotate_field (4);
    if (TARGET_ADDR_BIT <= 32)
        ui_out_table_header (uiout, 10, ui_left, "addr", "Address"); /* 5 */
    else
        ui_out_table_header (uiout, 18, ui_left, "addr", "Address"); /* 5 */
}

```

```

    }
    if (nr_printable_breakpoints > 0)
        annotate_field (5);
    ui_out_table_header (uiout, 40, ui_noalign, "what", "What"); /* 6 */
    ui_out_table_body (uiout);
    if (nr_printable_breakpoints > 0)
        annotate_breakpoints_table ();

```

This example, from the `print_one_breakpoint` function, shows how to produce the actual data for the table whose structure was defined in the above example. The original code was:

```

    annotate_record ();
    annotate_field (0);
    printf_filtered ("% -3d ", b->number);
    annotate_field (1);
    if ((int)b->type > (sizeof(bptypes)/sizeof(bptypes[0]))
        || ((int) b->type != bptypes[(int) b->type].type))
        internal_error ("bptypes table does not describe type %#d.",
                        (int)b->type);
    printf_filtered ("% -14s ", bptypes[(int)b->type].description);
    annotate_field (2);
    printf_filtered ("% -4s ", bpdisps[(int)b->disposition]);
    annotate_field (3);
    printf_filtered ("% -3c ", bpenables[(int)b->enable]);
    ...

```

This is the new version:

```

    annotate_record ();
    ui_out_tuple_begin (uiout, "bkpt");
    annotate_field (0);
    ui_out_field_int (uiout, "number", b->number);
    annotate_field (1);
    if (((int) b->type > (sizeof (bptypes) / sizeof (bptypes[0])))
        || ((int) b->type != bptypes[(int) b->type].type))
        internal_error ("bptypes table does not describe type %#d.",
                        (int) b->type);
    ui_out_field_string (uiout, "type", bptypes[(int)b->type].description);
    annotate_field (2);
    ui_out_field_string (uiout, "disp", bpdisps[(int)b->disposition]);
    annotate_field (3);
    ui_out_field_fmt (uiout, "enabled", "%c", bpenables[(int)b->enable]);
    ...

```

This example, also from `print_one_breakpoint`, shows how to produce a complicated output field using the `print_expression` functions which requires a stream to be passed. It also shows how to automate stream destruction with cleanups. The original code was:

```

    annotate_field (5);
    print_expression (b->exp, gdb_stdout);

```

The new version is:

```

    struct ui_stream *stb = ui_out_stream_new (uiout);

```

```

struct cleanup *old_chain = make_cleanup_ui_out_stream_delete (stb);
...
annotate_field (5);
print_expression (b->exp, stb->stream);
ui_out_field_stream (uiout, "what", local_stream);

```

This example, also from `print_one_breakpoint`, shows how to use `ui_out_text` and `ui_out_field_string`. The original code was:

```

annotate_field (5);
if (b->dll_pathname == NULL)
    printf_filtered ("<any library> ");
else
    printf_filtered ("library \"%s\" ", b->dll_pathname);

```

It became:

```

annotate_field (5);
if (b->dll_pathname == NULL)
{
    ui_out_field_string (uiout, "what", "<any library>");
    ui_out_spaces (uiout, 1);
}
else
{
    ui_out_text (uiout, "library \"");
    ui_out_field_string (uiout, "what", b->dll_pathname);
    ui_out_text (uiout, "\" ");
}

```

The following example from `print_one_breakpoint` shows how to use `ui_out_field_int` and `ui_out_spaces`. The original code was:

```

annotate_field (5);
if (b->forked_inferior_pid != 0)
    printf_filtered ("process %d ", b->forked_inferior_pid);

```

It became:

```

annotate_field (5);
if (b->forked_inferior_pid != 0)
{
    ui_out_text (uiout, "process ");
    ui_out_field_int (uiout, "what", b->forked_inferior_pid);
    ui_out_spaces (uiout, 1);
}

```

Here's an example of using `ui_out_field_string`. The original code was:

```

annotate_field (5);
if (b->exec_pathname != NULL)
    printf_filtered ("program \"%s\" ", b->exec_pathname);

```

It became:

```

annotate_field (5);
if (b->exec_pathname != NULL)
{

```



```

    ui_out_text (uiout, "program \"");
    ui_out_field_string (uiout, "what", b->exec_pathname);
    ui_out_text (uiout, "\" ");
}

```

Finally, here's an example of printing an address. The original code:

```

annotate_field (4);
printf_filtered ("%s ",
    local_hex_string_custom ((unsigned long) b->address, "081"));

```

It became:

```

annotate_field (4);
ui_out_field_core_addr (uiout, "Address", b->address);

```

4.3 Console Printing

4.4 TUI

5 libgdb

5.1 libgdb 1.0

libgdb 1.0 was an abortive project of years ago. The theory was to provide an API to GDB's functionality.

5.2 libgdb 2.0

libgdb 2.0 is an ongoing effort to update GDB so that is better able to support graphical and other environments.

Since libgdb development is on-going, its architecture is still evolving. The following components have so far been identified:

- Observer - 'gdb-events.h'.
- Builder - 'ui-out.h'
- Event Loop - 'event-loop.h'
- Library - 'gdb.h'

The model that ties these components together is described below.

5.3 The libgdb Model

A client of `libgdb` interacts with the library in two ways.

- As an observer (using ‘`gdb-events`’) receiving notifications from `libgdb` of any internal state changes (break point changes, run state, etc).
- As a client querying `libgdb` (using the ‘`ui-out`’ builder) to obtain various status values from GDB.

Since `libgdb` could have multiple clients (e.g. a GUI supporting the existing GDB CLI), those clients must co-operate when controlling `libgdb`. In particular, a client must ensure that `libgdb` is idle (i.e. no other client is using `libgdb`) before responding to a ‘`gdb-event`’ by making a query.

5.4 CLI support

At present GDB’s CLI is very much entangled in with the core of `libgdb`. Consequently, a client wishing to include the CLI in their interface needs to carefully co-ordinate its own and the CLI’s requirements.

It is suggested that the client set `libgdb` up to be bi-modal (alternate between CLI and client query modes). The notes below sketch out the theory:

- The client registers itself as an observer of `libgdb`.
- The client create and install `cli-out` builder using its own versions of the `ui-file` `gdb_stderr`, `gdb_stdtdarg` and `gdb_stdout` streams.
- The client creates a separate custom `ui-out` builder that is only used while making direct queries to `libgdb`.

When the client receives input intended for the CLI, it simply passes it along. Since the `cli-out` builder is installed by default, all the CLI output in response to that command is routed (pronounced rooted) through to the client controlled `gdb_stdout` et. al. streams. At the same time, the client is kept abreast of internal changes by virtue of being a `libgdb` observer.

The only restriction on the client is that it must wait until `libgdb` becomes idle before initiating any queries (using the client’s custom builder).

5.5 libgdb components

Observer - ‘`gdb-events.h`’

‘`gdb-events`’ provides the client with a very raw mechanism that can be used to implement an observer. At present it only allows for one observer and that observer must, internally, handle the need to delay the processing of any event notifications until after `libgdb` has finished the current command.

Builder - ‘`ui-out.h`’

‘`ui-out`’ provides the infrastructure necessary for a client to create a builder. That builder is then passed down to `libgdb` when doing any queries.

Event Loop - ‘event-loop.h’

‘event-loop’, currently non-re-entrant, provides a simple event loop. A client would need to either plug its self into this loop or, implement a new event-loop that GDB would use.

The event-loop will eventually be made re-entrant. This is so that [No value for “GDB”] can better handle the problem of some commands blocking instead of returning.

Library - ‘gdb.h’

‘libgdb’ is the most obvious component of this system. It provides the query interface. Each function is parameterized by a ui-out builder. The result of the query is constructed using that builder before the query function returns.

6 Symbol Handling

Symbols are a key part of GDB’s operation. Symbols include variables, functions, and types.

6.1 Symbol Reading

GDB reads symbols from *symbol files*. The usual symbol file is the file containing the program which GDB is debugging. GDB can be directed to use a different file for symbols (with the ‘symbol-file’ command), and it can also read more symbols via the ‘add-file’ and ‘load’ commands, or while reading symbols from shared libraries.

Symbol files are initially opened by code in ‘symfile.c’ using the BFD library (see [Chapter 12 \[Support Libraries\], page 58](#)). BFD identifies the type of the file by examining its header. `find_sym_fns` then uses this identification to locate a set of symbol-reading functions.

Symbol-reading modules identify themselves to GDB by calling `add_symtab_fns` during their module initialization. The argument to `add_symtab_fns` is a `struct sym_fns` which contains the name (or name prefix) of the symbol format, the length of the prefix, and pointers to four functions. These functions are called at various times to process symbol files whose identification matches the specified prefix.

The functions supplied by each module are:

`xyz_symfile_init(struct sym_fns *sf)`

Called from `symbol_file_add` when we are about to read a new symbol file. This function should clean up any internal state (possibly resulting from half-read previous files, for example) and prepare to read a new symbol file. Note that the symbol file which we are reading might be a new “main” symbol file, or might be a secondary symbol file whose symbols are being added to the existing symbol table.

The argument to `xyz_symfile_init` is a newly allocated `struct sym_fns` whose `bfd` field contains the BFD for the new symbol file being read. Its `private` field has been zeroed, and can be modified as desired. Typically, a

struct of private information will be `malloc`'d, and a pointer to it will be placed in the `private` field.

There is no result from `xyz_symfile_init`, but it can call `error` if it detects an unavoidable problem.

`xyz_new_init()`

Called from `symbol_file_add` when discarding existing symbols. This function needs only handle the symbol-reading module's internal state; the symbol table data structures visible to the rest of GDB will be discarded by `symbol_file_add`. It has no arguments and no result. It may be called after `xyz_symfile_init`, if a new symbol table is being read, or may be called alone if all symbols are simply being discarded.

`xyz_symfile_read(struct sym_fns *sf, CORE_ADDR addr, int mainline)`

Called from `symbol_file_add` to actually read the symbols from a symbol-file into a set of `psymtabs` or `symtabs`.

`sf` points to the `struct sym_fns` originally passed to `xyz_sym_init` for possible initialization. `addr` is the offset between the file's specified start address and its true address in memory. `mainline` is 1 if this is the main symbol table being read, and 0 if a secondary symbol file (e.g. shared library or dynamically loaded file) is being read.

In addition, if a symbol-reading module creates `psymtabs` when `xyz_symfile_read` is called, these `psymtabs` will contain a pointer to a function `xyz_psymtab_to_symtab`, which can be called from any point in the GDB symbol-handling code.

`xyz_psymtab_to_symtab (struct partial_symtab *pst)`

Called from `psymtab_to_symtab` (or the `PSYMTAB_TO_SYMTAB` macro) if the `psymtab` has not already been read in and had its `pst->symtab` pointer set. The argument is the `psymtab` to be fleshed-out into a `symtab`. Upon return, `pst->readin` should have been set to 1, and `pst->symtab` should contain a pointer to the new corresponding `symtab`, or zero if there were no symbols in that part of the symbol file.

6.2 Partial Symbol Tables

GDB has three types of symbol tables:

- Full symbol tables (*symtabs*). These contain the main information about symbols and addresses.
- Partial symbol tables (*psymtabs*). These contain enough information to know when to read the corresponding part of the full symbol table.
- Minimal symbol tables (*msymtabs*). These contain information gleaned from non-debugging symbols.

This section describes partial symbol tables.

A `psymtab` is constructed by doing a very quick pass over an executable file's debugging information. Small amounts of information are extracted—enough to identify which parts of the symbol table will need to be re-read and fully digested later, when the user needs

the information. The speed of this pass causes GDB to start up very quickly. Later, as the detailed rereading occurs, it occurs in small pieces, at various times, and the delay therefrom is mostly invisible to the user.

The symbols that show up in a file's psymtab should be, roughly, those visible to the debugger's user when the program is not running code from that file. These include external symbols and types, static symbols and types, and `enum` values declared at file scope.

The psymtab also contains the range of instruction addresses that the full symbol table would represent.

The idea is that there are only two ways for the user (or much of the code in the debugger) to reference a symbol:

- By its address (e.g. execution stops at some address which is inside a function in this file). The address will be noticed to be in the range of this psymtab, and the full symtab will be read in. `find_pc_function`, `find_pc_line`, and other `find_pc_...` functions handle this.
- By its name (e.g. the user asks to print a variable, or set a breakpoint on a function). Global names and file-scope names will be found in the psymtab, which will cause the symtab to be pulled in. Local names will have to be qualified by a global name, or a file-scope name, in which case we will have already read in the symtab as we evaluated the qualifier. Or, a local symbol can be referenced when we are "in" a local scope, in which case the first case applies. `lookup_symbol` does most of the work here.

The only reason that psymtabs exist is to cause a symtab to be read in at the right moment. Any symbol that can be elided from a psymtab, while still causing that to happen, should not appear in it. Since psymtabs don't have the idea of scope, you can't put local symbols in them anyway. Psymtabs don't have the idea of the type of a symbol, either, so types need not appear, unless they will be referenced by name.

It is a bug for GDB to behave one way when only a psymtab has been read, and another way if the corresponding symtab has been read in. Such bugs are typically caused by a psymtab that does not contain all the visible symbols, or which has the wrong instruction address ranges.

The psymtab for a particular section of a symbol file (objfile) could be thrown away after the symtab has been read in. The symtab should always be searched before the psymtab, so the psymtab will never be used (in a bug-free environment). Currently, psymtabs are allocated on an obstack, and all the symbols themselves are allocated in a pair of large arrays on an obstack, so there is little to be gained by trying to free them unless you want to do a lot more work.

6.3 Types

Fundamental Types (e.g., `FT_VOID`, `FT_BOOLEAN`).

These are the fundamental types that GDB uses internally. Fundamental types from the various debugging formats (stabs, ELF, etc) are mapped into one of these. They are basically a union of all fundamental types that GDB knows about for all the languages that GDB knows about.

Type Codes (e.g., `TYPE_CODE_PTR`, `TYPE_CODE_ARRAY`).

Each time GDB builds an internal type, it marks it with one of these types. The type may be a fundamental type, such as `TYPE_CODE_INT`, or a derived type, such as `TYPE_CODE_PTR` which is a pointer to another type. Typically, several `FT_*` types map to one `TYPE_CODE_*` type, and are distinguished by other members of the type struct, such as whether the type is signed or unsigned, and how many bits it uses.

Builtin Types (e.g., `builtin_type_void`, `builtin_type_char`).

These are instances of type structs that roughly correspond to fundamental types and are created as global types for GDB to use for various ugly historical reasons. We eventually want to eliminate these. Note for example that `builtin_type_int` initialized in `'gdbtypes.c'` is basically the same as a `TYPE_CODE_INT` type that is initialized in `'c-lang.c'` for an `FT_INTEGER` fundamental type. The difference is that the `builtin_type` is not associated with any particular objfile, and only one instance exists, while `'c-lang.c'` builds as many `TYPE_CODE_INT` types as needed, with each one associated with some particular objfile.

6.4 Object File Formats

6.4.1 a.out

The `a.out` format is the original file format for Unix. It consists of three sections: `text`, `data`, and `bss`, which are for program code, initialized data, and uninitialized data, respectively.

The `a.out` format is so simple that it doesn't have any reserved place for debugging information. (Hey, the original Unix hackers used `'adb'`, which is a machine-language debugger!) The only debugging format for `a.out` is stabs, which is encoded as a set of normal symbols with distinctive attributes.

The basic `a.out` reader is in `'dbxread.c'`.

6.4.2 COFF

The COFF format was introduced with System V Release 3 (SVR3) Unix. COFF files may have multiple sections, each prefixed by a header. The number of sections is limited.

The COFF specification includes support for debugging. Although this was a step forward, the debugging information was woefully limited. For instance, it was not possible to represent code that came from an included file.

The COFF reader is in `'coffread.c'`.

6.4.3 ECOFF

ECOFF is an extended COFF originally introduced for Mips and Alpha workstations.

The basic ECOFF reader is in `'mipsread.c'`.

6.4.4 XCOFF

The IBM RS/6000 running AIX uses an object file format called XCOFF. The COFF sections, symbols, and line numbers are used, but debugging symbols are `dbx`-style stabs whose strings are located in the `.debug` section (rather than the string table). For more information, see [section “Top” in *The Stabs Debugging Format*](#).

The shared library scheme has a clean interface for figuring out what shared libraries are in use, but the catch is that everything which refers to addresses (symbol tables and breakpoints at least) needs to be relocated for both shared libraries and the main executable. At least using the standard mechanism this can only be done once the program has been run (or the core file has been read).

6.4.5 PE

Windows 95 and NT use the PE (*Portable Executable*) format for their executables. PE is basically COFF with additional headers.

While BFD includes special PE support, GDB needs only the basic COFF reader.

6.4.6 ELF

The ELF format came with System V Release 4 (SVR4) Unix. ELF is similar to COFF in being organized into a number of sections, but it removes many of COFF's limitations.

The basic ELF reader is in `'elfread.c'`.

6.4.7 SOM

SOM is HP's object file and debug format (not to be confused with IBM's SOM, which is a cross-language ABI).

The SOM reader is in `'hpread.c'`.

6.4.8 Other File Formats

Other file formats that have been supported by GDB include Netware Loadable Modules (`'nlmread.c'`).

6.5 Debugging File Formats

This section describes characteristics of debugging information that are independent of the object file format.

6.5.1 stabs

`stabs` started out as special symbols within the `a.out` format. Since then, it has been encapsulated into other file formats, such as COFF and ELF.

While `'dbxread.c'` does some of the basic stab processing, including for encapsulated versions, `'stabsread.c'` does the real work.

6.5.2 COFF

The basic COFF definition includes debugging information. The level of support is minimal and non-extensible, and is not often used.

6.5.3 Mips debug (Third Eye)

ECOFF includes a definition of a special debug format.

The file `'mdebugread.c'` implements reading for this format.

6.5.4 DWARF 1

DWARF 1 is a debugging format that was originally designed to be used with ELF in SVR4 systems.

The DWARF 1 reader is in `'dwarfread.c'`.

6.5.5 DWARF 2

DWARF 2 is an improved but incompatible version of DWARF 1.

The DWARF 2 reader is in `'dwarf2read.c'`.

6.5.6 SOM

Like COFF, the SOM definition includes debugging information.

6.6 Adding a New Symbol Reader to GDB

If you are using an existing object file format (`a.out`, COFF, ELF, etc), there is probably little to be done.

If you need to add a new object file format, you must first add it to BFD. This is beyond the scope of this document.

You must then arrange for the BFD code to provide access to the debugging symbols. Generally GDB will have to call swapping routines from BFD and a few other BFD internal routines to locate the debugging information. As much as possible, GDB should not depend on the BFD internal data structures.

For some targets (e.g., COFF), there is a special transfer vector used to call swapping routines, since the external data structures on various platforms have different sizes and layouts. Specialized routines that will only ever be implemented by one object file format may be called directly. This interface should be described in a file `'bfd/libxyz.h'`, which is included by GDB.

7 Language Support

GDB's language support is mainly driven by the symbol reader, although it is possible for the user to set the source language manually.

GDB chooses the source language by looking at the extension of the file recorded in the debug info; `.c` means C, `.f` means Fortran, etc. It may also use a special-purpose language identifier if the debug format supports it, like with DWARF.

7.1 Adding a Source Language to GDB

To add other languages to GDB's expression parser, follow the following steps:

Create the expression parser.

This should reside in a file `'lang-exp.y'`. Routines for building parsed expressions into a union `exp_element` list are in `'parse.c'`.

Since we can't depend upon everyone having Bison, and YACC produces parsers that define a bunch of global names, the following lines **must** be included at the top of the YACC parser, to prevent the various parsers from defining the same global names:

```
#define yyparse      lang_parse
#define yylex        lang_lex
#define yyerror      lang_error
#define yylval       lang_lval
#define yychar       lang_char
#define yydebug      lang_debug
#define yypact       lang_pact
#define yyr1         lang_r1
#define yyr2         lang_r2
#define yydef        lang_def
#define yychk        lang_chk
#define yypgo        lang_pgo
#define yyact        lang_act
#define yyexca       lang_exca
#define yyerrflag    lang_errflag
#define yynerrs      lang_nerrs
```

At the bottom of your parser, define a `struct language_defn` and initialize it with the right values for your language. Define an `initialize_lang` routine and have it call `'add_language(lang_language_defn)'` to tell the rest of GDB that your language exists. You'll need some other supporting variables and functions, which will be used via pointers from your `lang_language_defn`. See the declaration of `struct language_defn` in `'language.h'`, and the other `'*-exp.y'` files, for more information.

Add any evaluation routines, if necessary

If you need new opcodes (that represent the operations of the language), add them to the enumerated type in `'expression.h'`. Add support code for these operations in the `evaluate_subexp` function defined in the file `'eval.c'`. Add cases for new opcodes in two functions from `'parse.c'`: `prefixify_subexp` and

`length_of_subexp`. These compute the number of `exp_elements` that a given operation takes up.

Update some existing code

Add an enumerated identifier for your language to the enumerated type `enum language` in `'defs.h'`.

Update the routines in `'language.c'` so your language is included. These routines include type predicates and such, which (in some cases) are language dependent. If your language does not appear in the switch statement, an error is reported.

Also included in `'language.c'` is the code that updates the variable `current_language`, and the routines that translate the `language_lang` enumerated identifier into a printable string.

Update the function `_initialize_language` to include your language. This function picks the default language upon startup, so is dependent upon which languages that GDB is built for.

Update `allocate_symtab` in `'symfile.c'` and/or symbol-reading code so that the language of each symtab (source file) is set properly. This is used to determine the language to use at each stack frame level. Currently, the language is set based upon the extension of the source file. If the language can be better inferred from the symbol information, please set the language of the symtab in the symbol-reading code.

Add helper code to `print_subexp` (in `'expprint.c'`) to handle any new expression opcodes you have added to `'expression.h'`. Also, add the printed representations of your operators to `op_print_tab`.

Add a place of call

Add a call to `lang_parse()` and `lang_error` in `parse_exp_1` (defined in `'parse.c'`).

Use macros to trim code

The user has the option of building GDB for some or all of the languages. If the user decides to build GDB for the language `lang`, then every file dependent on `'language.h'` will have the macro `_LANG_lang` defined in it. Use `#ifdefs` to leave out large routines that the user won't need if he or she is not using your language.

Note that you do not need to do this in your YACC parser, since if GDB is not build for `lang`, then `'lang-exp.tab.o'` (the compiled form of your parser) is not linked into GDB at all.

See the file `'configure.in'` for how GDB is configured for different languages.

Edit 'Makefile.in'

Add dependencies in `'Makefile.in'`. Make sure you update the macro variables such as `HFILES` and `OBJS`, otherwise your code may not get linked in, or, worse yet, it may not get tarred into the distribution!

8 Host Definition

Maintainer's note: In theory, new targets no longer need to use the host framework described below. Instead it should be possible to handle everything using autoconf. Patches eliminating this framework welcome.

With the advent of Autoconf, it's rarely necessary to have host definition machinery anymore.

8.1 Adding a New Host

Most of GDB's host configuration support happens via Autoconf. New host-specific definitions should be rarely needed. GDB still uses the host-specific definitions and files listed below, but these mostly exist for historical reasons, and should eventually disappear.

Several files control GDB's configuration for host systems:

`'gdb/config/arch/xyz.mh'`

Specifies Makefile fragments needed when hosting on machine xyz. In particular, this lists the required machine-dependent object files, by defining `'XDEPFILES=...'`. Also specifies the header file which describes host xyz, by defining `XM_FILE= xm-xyz.h`. You can also define `CC`, `SYSV_DEFINE`, `XM_CFLAGS`, `XM_ADD_FILES`, `XM_CLIBS`, `XM_CDEPS`, etc.; see `'Makefile.in'`.

`'gdb/config/arch/xm-xyz.h'`

(`'xm.h'` is a link to this file, created by `configure`). Contains C macro definitions describing the host system environment, such as byte order, host C compiler and library.

`'gdb/xyz-xdep.c'`

Contains any miscellaneous C code required for this machine as a host. On most machines it doesn't exist at all. If it does exist, put `'xyz-xdep.o'` into the `XDEPFILES` line in `'gdb/config/arch/xyz.mh'`.

Generic Host Support Files

There are some "generic" versions of routines that can be used by various systems. These can be customized in various ways by macros defined in your `'xm-xyz.h'` file. If these routines work for the xyz host, you can just include the generic file's name (with `' .o'`, not `' .c'`) in `XDEPFILES`.

Otherwise, if your machine needs custom support routines, you will need to write routines that perform the same functions as the generic file. Put them into `xyz-xdep.c`, and put `xyz-xdep.o` into `XDEPFILES`.

`'ser-unix.c'`

This contains serial line support for Unix systems. This is always included, via the makefile variable `SER_HARDWARE`; override this variable in the `' .mh'` file to avoid it.

`'ser-go32.c'`

This contains serial line support for 32-bit programs running under DOS, using the DJGPP (a.k.a. GO32) execution environment.

`'ser-tcp.c'`

This contains generic TCP support using sockets.

8.2 Host Conditionals

When GDB is configured and compiled, various macros are defined or left undefined, to control compilation based on the attributes of the host system. These macros and their meanings (or if the meaning is not documented here, then one of the source files where they are used is indicated) are:

GDBINIT_FILENAME

The default name of GDB's initialization file (normally `'gdbinit'`).

MEM_FNS_DECLARED

Your host config file defines this if it includes declarations of `memcpy` and `memset`. Define this to avoid conflicts between the native include files and the declarations in `'defs.h'`.

NO_STD_REGS

This macro is deprecated.

NO_SYS_FILE

Define this if your system does not have a `<sys/file.h>`.

SIGWINCH_HANDLER

If your host defines `SIGWINCH`, you can define this to be the name of a function to be called if `SIGWINCH` is received.

SIGWINCH_HANDLER_BODY

Define this to expand into code that will define the function named by the expansion of `SIGWINCH_HANDLER`.

ALIGN_STACK_ON_STARTUP

Define this if your system is of a sort that will crash in `tgetent` if the stack happens not to be longword-aligned when `main` is called. This is a rare situation, but is known to occur on several different types of systems.

CRLF_SOURCE_FILES

Define this if host files use `\r\n` rather than `\n` as a line terminator. This will cause source file listings to omit `\r` characters when printing and it will allow `\r\n` line endings of files which are "sourced" by gdb. It must be possible to open files in binary mode using `O_BINARY` or, for `fopen`, `"rb"`.

DEFAULT_PROMPT

The default value of the prompt string (normally `"(gdb) "`).

DEV_TTY

The name of the generic TTY device, defaults to `"/dev/tty"`.

FCLOSE_PROVIDED

Define this if the system declares `fclose` in the headers included in `defs.h`. This isn't needed unless your compiler is unusually anal.

FOPEN_RB Define this if binary files are opened the same way as text files.

GETENV_PROVIDED

Define this if the system declares `getenv` in its headers included in `defs.h`. This isn't needed unless your compiler is unusually anal.

HAVE_MMAP

In some cases, use the system call `mmap` for reading symbol tables. For some machines this allows for sharing and quick updates.

HAVE_SIGSETMASK

Define this if the host system has job control, but does not define `sigsetmask`. Currently, this is only true of the RS/6000.

HAVE_TERMIO

Define this if the host system has `termio.h`.

HOST_BYTE_ORDER

The ordering of bytes in the host. This must be defined to be either `BIG_ENDIAN` or `LITTLE_ENDIAN`.

INT_MAX**INT_MIN****LONG_MAX****UINT_MAX****ULONG_MAX**

Values for host-side constants.

ISATTY Substitute for `isatty`, if not available.

LONGEST This is the longest integer type available on the host. If not defined, it will default to `long long` or `long`, depending on `CC_HAS_LONG_LONG`.

CC_HAS_LONG_LONG

Define this if the host C compiler supports `long long`. This is set by the `configure` script.

PRINTF_HAS_LONG_LONG

Define this if the host can handle printing of `long long` integers via the `printf` format conversion specifier `ll`. This is set by the `configure` script.

HAVE_LONG_DOUBLE

Define this if the host C compiler supports `long double`. This is set by the `configure` script.

PRINTF_HAS_LONG_DOUBLE

Define this if the host can handle printing of `long double` float-point numbers via the `printf` format conversion specifier `Lg`. This is set by the `configure` script.

SCANF_HAS_LONG_DOUBLE

Define this if the host can handle the parsing of long double float-point numbers via the `scanf` format conversion specifier `Lg`. This is set by the `configure` script.

LSEEK_NOT_LINEAR

Define this if `lseek (n)` does not necessarily move to byte number `n` in the file. This is only used when reading source files. It is normally faster to define `CRLF_SOURCE_FILES` when possible.

L_SET

This macro is used as the argument to `lseek` (or, most commonly, `bfd_seek`). `FIXME`, should be replaced by `SEEK_SET` instead, which is the POSIX equivalent.

MALLOC_INCOMPATIBLE

Define this if the system's prototype for `malloc` differs from the ANSI definition.

MMAP_BASE_ADDRESS

When using `HAVE_MMAP`, the first mapping should go at this address.

MMAP_INCREMENT

when using `HAVE_MMAP`, this is the increment between mappings.

NORETURN

If defined, this should be one or more tokens, such as `volatile`, that can be used in both the declaration and definition of functions to indicate that they never return. The default is already set correctly if compiling with GCC. This will almost never need to be defined.

ATTR_NORETURN

If defined, this should be one or more tokens, such as `__attribute__((noreturn))`, that can be used in the declarations of functions to indicate that they never return. The default is already set correctly if compiling with GCC. This will almost never need to be defined.

USE_GENERIC_DUMMY_FRAMES

Define this to 1 if the target is using the generic inferior function call code. See `blockframe.c` for more information.

USE_MMALLOC

GDB will use the `mmalloc` library for memory allocation for symbol reading if this symbol is defined. Be careful defining it since there are systems on which `mmalloc` does not work for some reason. One example is the DECstation, where its RPC library can't cope with our redefinition of `malloc` to call `mmalloc`. When defining `USE_MMALLOC`, you will also have to set `MMALLOC` in the Makefile, to point to the `mmalloc` library. This define is set when you configure with `'--with-mmalloc'`.

NO_MMCHECK

Define this if you are using `mmalloc`, but don't want the overhead of checking the heap with `mmcheck`. Note that on some systems, the C runtime makes calls to `malloc` prior to calling `main`, and if `free` is ever called with these pointers

after calling `mmcheck` to enable checking, a memory corruption abort is certain to occur. These systems can still use `mmalloc`, but must define `NO_MMCHECK`.

`MMCHECK_FORCE`

Define this to 1 if the C runtime allocates memory prior to `mmcheck` being called, but that memory is never freed so we don't have to worry about it triggering a memory corruption abort. The default is 0, which means that `mmcheck` will only install the heap checking functions if there has not yet been any memory allocation calls, and if it fails to install the functions, GDB will issue a warning. This is currently defined if you configure using `'--with-mmalloc'`.

`NO_SIGINTERRUPT`

Define this to indicate that `siginterrupt` is not available.

`R_OK` Define if this is not in a system header file (typically, `'unistd.h'`).

`SEEK_CUR`

`SEEK_SET` Define these to appropriate value for the system `lseek`, if not already defined.

`STOP_SIGNAL`

This is the signal for stopping GDB. Defaults to `SIGTSTP`. (Only redefined for the Convex.)

`USE_O_NOCTTY`

Define this if the interior's tty should be opened with the `O_NOCTTY` flag. (FIXME: This should be a native-only flag, but `'inflow.c'` is always linked in.)

`USG` Means that System V (prior to SVR4) include files are in use. (FIXME: This symbol is abused in `'infrun.c'`, `'regex.c'`, `'remote-nindy.c'`, and `'utils.c'` for other things, at the moment.)

`lint` Define this to help placate `lint` in some situations.

`volatile` Define this to override the defaults of `__volatile__` or `/**/`.

9 Target Architecture Definition

GDB's target architecture defines what sort of machine-language programs GDB can work with, and how it works with them.

The target architecture object is implemented as the C structure `struct gdbarch *`. The structure, and its methods, are generated using the Bourn shell script `'gdbarch.sh'`.

9.1 Registers and Memory

GDB's model of the target machine is rather simple. GDB assumes the machine includes a bank of registers and a block of memory. Each register may have a different size.

GDB does not have a magical way to match up with the compiler's idea of which registers are which; however, it is critical that they do match up accurately. The only way to make

this work is to get accurate information about the order that the compiler uses, and to reflect that in the `REGISTER_NAME` and related macros.

GDB can handle big-endian, little-endian, and bi-endian architectures.

9.2 Pointers Are Not Always Addresses

On almost all 32-bit architectures, the representation of a pointer is indistinguishable from the representation of some fixed-length number whose value is the byte address of the object pointed to. On such machines, the words “pointer” and “address” can be used interchangeably. However, architectures with smaller word sizes are often cramped for address space, so they may choose a pointer representation that breaks this identity, and allows a larger code address space.

For example, the Mitsubishi D10V is a 16-bit VLIW processor whose instructions are 32 bits long³. If the D10V used ordinary byte addresses to refer to code locations, then the processor would only be able to address 64kb of instructions. However, since instructions must be aligned on four-byte boundaries, the low two bits of any valid instruction’s byte address are always zero—byte addresses waste two bits. So instead of byte addresses, the D10V uses word addresses—byte addresses shifted right two bits—to refer to code. Thus, the D10V can use 16-bit words to address 256kb of code space.

However, this means that code pointers and data pointers have different forms on the D10V. The 16-bit word `0xC020` refers to byte address `0xC020` when used as a data address, but refers to byte address `0x30080` when used as a code address.

(The D10V also uses separate code and data address spaces, which also affects the correspondence between pointers and addresses, but we’re going to ignore that here; this example is already too long.)

To cope with architectures like this—the D10V is not the only one!—GDB tries to distinguish between *addresses*, which are byte numbers, and *pointers*, which are the target’s representation of an address of a particular type of data. In the example above, `0xC020` is the pointer, which refers to one of the addresses `0xC020` or `0x30080`, depending on the type imposed upon it. GDB provides functions for turning a pointer into an address and vice versa, in the appropriate way for the current architecture.

Unfortunately, since addresses and pointers are identical on almost all processors, this distinction tends to bit-rot pretty quickly. Thus, each time you port GDB to an architecture which does distinguish between pointers and addresses, you’ll probably need to clean up some architecture-independent code.

Here are functions which convert between pointers and addresses:

CORE_ADDR extract_typed_address (<code>void *buf</code> , <code>struct type *type</code>)	Function
<p>Treat the bytes at <i>buf</i> as a pointer or reference of type <i>type</i>, and return the address it represents, in a manner appropriate for the current architecture. This yields an address GDB can use to read target memory, disassemble, etc. Note that <i>buf</i> refers to a buffer in GDB’s memory, not the inferior’s.</p>	

³ Some D10V instructions are actually pairs of 16-bit sub-instructions. However, since you can’t jump into the middle of such a pair, code addresses can only refer to full 32 bit instructions, which is what matters in this explanation.

For example, if the current architecture is the Intel x86, this function extracts a little-endian integer of the appropriate length from *buf* and returns it. However, if the current architecture is the D10V, this function will return a 16-bit integer extracted from *buf*, multiplied by four if *type* is a pointer to a function.

If *type* is not a pointer or reference type, then this function will signal an internal error.

CORE_ADDR store_typed_address (void **buf*, struct type **type*,
CORE_ADDR *addr*) Function

Store the address *addr* in *buf*, in the proper format for a pointer of type *type* in the current architecture. Note that *buf* refers to a buffer in GDB's memory, not the inferior's.

For example, if the current architecture is the Intel x86, this function stores *addr* unmodified as a little-endian integer of the appropriate length in *buf*. However, if the current architecture is the D10V, this function divides *addr* by four if *type* is a pointer to a function, and then stores it in *buf*.

If *type* is not a pointer or reference type, then this function will signal an internal error.

CORE_ADDR value_as_pointer (value_ptr *val*) Function

Assuming that *val* is a pointer, return the address it represents, as appropriate for the current architecture.

This function actually works on integral values, as well as pointers. For pointers, it performs architecture-specific conversions as described above for **extract_typed_address**.

CORE_ADDR value_from_pointer (struct type **type*, CORE_ADDR
addr) Function

Create and return a value representing a pointer of type *type* to the address *addr*, as appropriate for the current architecture. This function performs architecture-specific conversions as described above for **store_typed_address**.

GDB also provides functions that do the same tasks, but assume that pointers are simply byte addresses; they aren't sensitive to the current architecture, beyond knowing the appropriate endianness.

CORE_ADDR extract_address (void **addr*, int *len*) Function

Extract a *len*-byte number from *addr* in the appropriate endianness for the current architecture, and return it. Note that *addr* refers to GDB's memory, not the inferior's.

This function should only be used in architecture-specific code; it doesn't have enough information to turn bits into a true address in the appropriate way for the current architecture. If you can, use **extract_typed_address** instead.

void store_address (void **addr*, int *len*, LONGEST *val*) Function

Store *val* at *addr* as a *len*-byte integer, in the appropriate endianness for the current architecture. Note that *addr* refers to a buffer in GDB's memory, not the inferior's.

This function should only be used in architecture-specific code; it doesn't have enough information to turn a true address into bits in the appropriate way for the current architecture. If you can, use `store_typed_address` instead.

Here are some macros which architectures can define to indicate the relationship between pointers and addresses. These have default definitions, appropriate for architectures on which all pointers are simple byte addresses.

CORE_ADDR POINTER_TO_ADDRESS (struct type *type, char *buf) Target Macro

Assume that *buf* holds a pointer of type *type*, in the appropriate format for the current architecture. Return the byte address the pointer refers to.

This function may safely assume that *type* is either a pointer or a C++ reference type.

void ADDRESS_TO_POINTER (struct type *type, char *buf, CORE_ADDR addr) Target Macro

Store in *buf* a pointer of type *type* representing the address *addr*, in the appropriate format for the current architecture.

This function may safely assume that *type* is either a pointer or a C++ reference type.

9.3 Using Different Register and Memory Data Representations

Maintainer's note: The way GDB manipulates registers is undergoing significant change. Many of the macros and functions referred to in the sections below are likely to be made obsolete. See the file 'TODO' for more up-to-date information.

Some architectures use one representation for a value when it lives in a register, but use a different representation when it lives in memory. In GDB's terminology, the *raw* representation is the one used in the target registers, and the *virtual* representation is the one used in memory, and within GDB `struct value` objects.

For almost all data types on almost all architectures, the virtual and raw representations are identical, and no special handling is needed. However, they do occasionally differ. For example:

- The x86 architecture supports an 80-bit `long double` type. However, when we store those values in memory, they occupy twelve bytes: the floating-point number occupies the first ten, and the final two bytes are unused. This keeps the values aligned on four-byte boundaries, allowing more efficient access. Thus, the x86 80-bit floating-point type is the raw representation, and the twelve-byte loosely-packed arrangement is the virtual representation.
- Some 64-bit MIPS targets present 32-bit registers to GDB as 64-bit registers, with garbage in their upper bits. GDB ignores the top 32 bits. Thus, the 64-bit form, with garbage in the upper 32 bits, is the raw representation, and the trimmed 32-bit representation is the virtual representation.

In general, the raw representation is determined by the architecture, or GDB's interface to the architecture, while the virtual representation can be chosen for GDB's convenience.

GDB's register file, `registers`, holds the register contents in raw format, and the GDB remote protocol transmits register values in raw format.

Your architecture may define the following macros to request conversions between the raw and virtual format:

int REGISTER_CONVERTIBLE (int *reg*) Target Macro

Return non-zero if register number *reg*'s value needs different raw and virtual formats.

You should not use `REGISTER_CONVERT_TO_VIRTUAL` for a register unless this macro returns a non-zero value for that register.

int REGISTER_RAW_SIZE (int *reg*) Target Macro

The size of register number *reg*'s raw value. This is the number of bytes the register will occupy in `registers`, or in a GDB remote protocol packet.

int REGISTER_VIRTUAL_SIZE (int *reg*) Target Macro

The size of register number *reg*'s value, in its virtual format. This is the size a `struct` value's buffer will have, holding that register's value.

struct type *REGISTER_VIRTUAL_TYPE (int *reg*) Target Macro

This is the type of the virtual representation of register number *reg*. Note that there is no need for a macro giving a type for the register's raw form; once the register's value has been obtained, GDB always uses the virtual form.

void REGISTER_CONVERT_TO_VIRTUAL (int *reg*, Target Macro
 struct type **type*, char **from*, char **to*)

Convert the value of register number *reg* to *type*, which should always be `REGISTER_VIRTUAL_TYPE (reg)`. The buffer at *from* holds the register's value in raw format; the macro should convert the value to virtual format, and place it at *to*.

Note that `REGISTER_CONVERT_TO_VIRTUAL` and `REGISTER_CONVERT_TO_RAW` take their *reg* and *type* arguments in different orders.

You should only use `REGISTER_CONVERT_TO_VIRTUAL` with registers for which the `REGISTER_CONVERTIBLE` macro returns a non-zero value.

void REGISTER_CONVERT_TO_RAW (struct type Target Macro
 **type*, int *reg*, char **from*, char **to*)

Convert the value of register number *reg* to *type*, which should always be `REGISTER_VIRTUAL_TYPE (reg)`. The buffer at *from* holds the register's value in raw format; the macro should convert the value to virtual format, and place it at *to*.

Note that `REGISTER_CONVERT_TO_VIRTUAL` and `REGISTER_CONVERT_TO_RAW` take their *reg* and *type* arguments in different orders.

9.4 Frame Interpretation

9.5 Inferior Call Setup

9.6 Compiler Characteristics

9.7 Target Conditionals

This section describes the macros that you can use to define the target machine.

ADDITIONAL_OPTIONS

ADDITIONAL_OPTION_CASES

ADDITIONAL_OPTION_HANDLER

ADDITIONAL_OPTION_HELP

These are a set of macros that allow the addition of additional command line options to GDB. They are currently used only for the unsupported i960 Nindy target, and should not be used in any other configuration.

ADDR_BITS_REMOVE (*addr*)

If a raw machine instruction address includes any bits that are not really part of the address, then define this macro to expand into an expression that zeroes those bits in *addr*. This is only used for addresses of instructions, and even then not in all contexts.

For example, the two low-order bits of the PC on the Hewlett-Packard PA 2.0 architecture contain the privilege level of the corresponding instruction. Since instructions must always be aligned on four-byte boundaries, the processor masks out these bits to generate the actual address of the instruction. ADDR_BITS_REMOVE should filter out these bits with an expression such as `((addr) & ~3)`.

ADDRESS_TO_POINTER (*type*, *buf*, *addr*)

Store in *buf* a pointer of type *type* representing the address *addr*, in the appropriate format for the current architecture. This macro may safely assume that *type* is either a pointer or a C++ reference type. See [Chapter 9 \[Pointers Are Not Always Addresses\]](#), page 33.

BEFORE_MAIN_LOOP_HOOK

Define this to expand into any code that you want to execute before the main loop starts. Although this is not, strictly speaking, a target conditional, that is how it is currently being used. Note that if a configuration were to define it one way for a host and a different way for the target, GDB will probably not compile, let alone run correctly. This macro is currently used only for the unsupported i960 Nindy target, and should not be used in any other configuration.

BELIEVE_PCC_PROMOTION

Define if the compiler promotes a `short` or `char` parameter to an `int`, but still reports the parameter as its original type, rather than the promoted type.

BELIEVE_PCC_PROMOTION_TYPE

Define this if GDB should believe the type of a `short` argument when compiled by `pcc`, but look within a full `int` space to get its value. Only defined for Sun-3 at present.

BITS_BIG_ENDIAN

Define this if the numbering of bits in the targets does **not** match the endianness of the target byte order. A value of 1 means that the bits are numbered in a big-endian bit order, 0 means little-endian.

BREAKPOINT

This is the character array initializer for the bit pattern to put into memory where a breakpoint is set. Although it's common to use a trap instruction for a breakpoint, it's not required; for instance, the bit pattern could be an invalid instruction. The breakpoint must be no longer than the shortest instruction of the architecture.

BREAKPOINT has been deprecated in favor of **BREAKPOINT_FROM_PC**.

BIG_BREAKPOINT**LITTLE_BREAKPOINT**

Similar to **BREAKPOINT**, but used for bi-endian targets.

BIG_BREAKPOINT and **LITTLE_BREAKPOINT** have been deprecated in favor of **BREAKPOINT_FROM_PC**.

REMOTE_BREAKPOINT**LITTLE_REMOTE_BREAKPOINT****BIG_REMOTE_BREAKPOINT**

Similar to **BREAKPOINT**, but used for remote targets.

BIG_REMOTE_BREAKPOINT and **LITTLE_REMOTE_BREAKPOINT** have been deprecated in favor of **BREAKPOINT_FROM_PC**.

BREAKPOINT_FROM_PC (*pcptr*, *lenptr*)

Use the program counter to determine the contents and size of a breakpoint instruction. It returns a pointer to a string of bytes that encode a breakpoint instruction, stores the length of the string to **lenptr*, and adjusts *pc* (if necessary) to point to the actual memory location where the breakpoint should be inserted.

Although it is common to use a trap instruction for a breakpoint, it's not required; for instance, the bit pattern could be an invalid instruction. The breakpoint must be no longer than the shortest instruction of the architecture.

Replaces all the other **BREAKPOINT** macros.

MEMORY_INSERT_BREAKPOINT (*addr*, *contents.cache*)**MEMORY_REMOVE_BREAKPOINT** (*addr*, *contents.cache*)

Insert or remove memory based breakpoints. Reasonable defaults (**default_memory_insert_breakpoint** and **default_memory_remove_breakpoint** respectively) have been provided so that it is not necessary to define these for most architectures. Architectures which may want to define **MEMORY_INSERT_BREAKPOINT** and **MEMORY_REMOVE_BREAKPOINT** will likely have instructions that are oddly sized or are not stored in a conventional manner.

It may also be desirable (from an efficiency standpoint) to define custom breakpoint insertion and removal routines if **BREAKPOINT_FROM_PC** needs to read the target's memory for some reason.

CALL_DUMMY_P

A C expression that is non-zero when the target supports inferior function calls.

CALL_DUMMY_WORDS

Pointer to an array of **LONGEST** words of data containing host-byte-ordered **REGISTER_BYTES** sized values that partially specify the sequence of instructions needed for an inferior function call.

Should be deprecated in favor of a macro that uses target-byte-ordered data.

SIZEOF_CALL_DUMMY_WORDS

The size of **CALL_DUMMY_WORDS**. When **CALL_DUMMY_P** this must return a positive value. See also **CALL_DUMMY_LENGTH**.

CALL_DUMMY

A static initializer for **CALL_DUMMY_WORDS**. Deprecated.

CALL_DUMMY_LOCATION

See the file ‘**inferior.h**’.

CALL_DUMMY_STACK_ADJUST

Stack adjustment needed when performing an inferior function call.

Should be deprecated in favor of something like **STACK_ALIGN**.

CALL_DUMMY_STACK_ADJUST_P

Predicate for use of **CALL_DUMMY_STACK_ADJUST**.

Should be deprecated in favor of something like **STACK_ALIGN**.

CANNOT_FETCH_REGISTER (*regno*)

A C expression that should be nonzero if *regno* cannot be fetched from an inferior process. This is only relevant if **FETCH_INFERIOR_REGISTERS** is not defined.

CANNOT_STORE_REGISTER (*regno*)

A C expression that should be nonzero if *regno* should not be written to the target. This is often the case for program counters, status words, and other special registers. If this is not defined, GDB will assume that all registers may be written.

DO_DEFERRED_STORES**CLEAR_DEFERRED_STORES**

Define this to execute any deferred stores of registers into the inferior, and to cancel any deferred stores.

Currently only implemented correctly for native Sparc configurations?

COERCE_FLOAT_TO_DOUBLE (*formal*, *actual*)

If we are calling a function by hand, and the function was declared (according to the debug info) without a prototype, should we automatically promote **floats** to **doubles**? This macro must evaluate to non-zero if we should, or zero if we should leave the value alone.

The argument *actual* is the type of the value we want to pass to the function. The argument *formal* is the type of this argument, as it appears in the function’s

definition. Note that *formal* may be zero if we have no debugging information for the function, or if we're passing more arguments than are officially declared (for example, varargs). This macro is never invoked if the function definitely has a prototype.

The default behavior is to promote only when we have no type information for the formal parameter. This is different from the obvious behavior, which would be to promote whenever we have no prototype, just as the compiler does. It's annoying, but some older targets rely on this. If you want GDB to follow the typical compiler behavior—to always promote when there is no prototype in scope—your `gdbarch_init` function can call `set_gdbarch_coerce_float_to_double` and select the `standard_coerce_float_to_double` function.

CPLUS_MARKER

Define this to expand into the character that G++ uses to distinguish compiler-generated identifiers from programmer-specified identifiers. By default, this expands into '\$'. Most System V targets should define this to '.'.

DBX_PARM_SYMBOL_CLASS

Hook for the `SYMBOL_CLASS` of a parameter when decoding DBX symbol information. In the i960, parameters can be stored as locals or as args, depending on the type of the debug record.

DECR_PC_AFTER_BREAK

Define this to be the amount by which to decrement the PC after the program encounters a breakpoint. This is often the number of bytes in `BREAKPOINT`, though not always. For most targets this value will be 0.

DECR_PC_AFTER_HW_BREAK

Similarly, for hardware breakpoints.

DISABLE_UNSETTABLE_BREAK (*addr*)

If defined, this should evaluate to 1 if *addr* is in a shared library in which breakpoints cannot be set and so should be disabled.

DO_REGISTERS_INFO

If defined, use this to print the value of a register or all registers.

DWARF_REG_TO_REGNUM

Convert DWARF register number into GDB regnum. If not defined, no conversion will be performed.

DWARF2_REG_TO_REGNUM

Convert DWARF2 register number into GDB regnum. If not defined, no conversion will be performed.

ECOFF_REG_TO_REGNUM

Convert ECOFF register number into GDB regnum. If not defined, no conversion will be performed.

END_OF_TEXT_DEFAULT

This is an expression that should designate the end of the text section.

EXTRACT_RETURN_VALUE(*type*, *regbuf*, *valbuf*)

Define this to extract a function's return value of type *type* from the raw register state *regbuf* and copy that, in virtual format, into *valbuf*.

EXTRACT_STRUCT_VALUE_ADDRESS(*regbuf*)

When defined, extract from the array *regbuf* (containing the raw register state) the **CORE_ADDR** at which a function should return its structure value.

If not defined, **EXTRACT_RETURN_VALUE** is used.

EXTRACT_STRUCT_VALUE_ADDRESS_P()

Predicate for **EXTRACT_STRUCT_VALUE_ADDRESS**.

FLOAT_INFO

If defined, then the '**info float**' command will print information about the processor's floating point unit.

FP_REGNUM

If the virtual frame pointer is kept in a register, then define this macro to be the number (greater than or equal to zero) of that register.

This should only need to be defined if **TARGET_READ_FP** and **TARGET_WRITE_FP** are not defined.

FRAMELESS_FUNCTION_INVOCATION(*fi*)

Define this to an expression that returns 1 if the function invocation represented by *fi* does not have a stack frame associated with it. Otherwise return 0.

FRAME_ARGS_ADDRESS_CORRECT

See '**stack.c**'.

FRAME_CHAIN(*frame*)

Given *frame*, return a pointer to the calling frame.

FRAME_CHAIN_COMBINE(*chain*, *frame*)

Define this to take the frame chain pointer and the frame's nominal address and produce the nominal address of the caller's frame. Presently only defined for HP PA.

FRAME_CHAIN_VALID(*chain*, *thisframe*)

Define this to be an expression that returns zero if the given frame is an outermost frame, with no caller, and nonzero otherwise. Several common definitions are available:

- **file_frame_chain_valid** is nonzero if the chain pointer is nonzero and given frame's PC is not inside the startup file (such as '**crt0.o**').
- **func_frame_chain_valid** is nonzero if the chain pointer is nonzero and the given frame's PC is not in **main** or a known entry point function (such as **_start**).
- **generic_file_frame_chain_valid** and **generic_func_frame_chain_valid** are equivalent implementations for targets using generic dummy frames.

FRAME_INIT_SAVED_REGS(*frame*)

See ‘*frame.h*’. Determines the address of all registers in the current stack frame storing each in *frame->saved_regs*. Space for *frame->saved_regs* shall be allocated by **FRAME_INIT_SAVED_REGS** using either *frame_saved_regs_zalloc* or *frame_obstack_alloc*.

FRAME_FIND_SAVED_REGS and **EXTRA_FRAME_INFO** are deprecated.

FRAME_NUM_ARGS (*fi*)

For the frame described by *fi* return the number of arguments that are being passed. If the number of arguments is not known, return -1.

FRAME_SAVED_PC(*frame*)

Given *frame*, return the pc saved there. This is the return address.

FUNCTION_EPILOGUE_SIZE

For some COFF targets, the *x_sym.x_misc.x_fsize* field of the function end symbol is 0. For such targets, you must define **FUNCTION_EPILOGUE_SIZE** to expand into the standard size of a function’s epilogue.

FUNCTION_START_OFFSET

An integer, giving the offset in bytes from a function’s address (as used in the values of symbols, function pointers, etc.), and the function’s first genuine instruction.

This is zero on almost all machines: the function’s address is usually the address of its first instruction. However, on the VAX, for example, each function starts with two bytes containing a bitmask indicating which registers to save upon entry to the function. The VAX *call* instructions check this value, and save the appropriate registers automatically. Thus, since the offset from the function’s address to its first instruction is two bytes, **FUNCTION_START_OFFSET** would be 2 on the VAX.

GCC_COMPILED_FLAG_SYMBOL**GCC2_COMPILED_FLAG_SYMBOL**

If defined, these are the names of the symbols that GDB will look for to detect that GCC compiled the file. The default symbols are *gcc_compiled.* and *gcc2_compiled.*, respectively. (Currently only defined for the Delta 68.)

GDB_MULTI_ARCH

If defined and non-zero, enables support for multiple architectures within GDB. This support can be enabled at two levels. At level one, only definitions for previously undefined macros are provided; at level two, a multi-arch definition of all architecture dependant macros will be defined.

GDB_TARGET_IS_HPPA

This determines whether horrible kludge code in ‘*dbxread.c*’ and ‘*partial-stab.h*’ is used to mangle multiple-symbol-table files from HPPA’s. This should all be ripped out, and a scheme like ‘*elfread.c*’ used instead.

GET_LONGJMP_TARGET

For most machines, this is a target-dependent parameter. On the DECstation and the Iris, this is a native-dependent parameter, since the header file ‘*setjmp.h*’ is needed to define it.

This macro determines the target PC address that `longjmp` will jump to, assuming that we have just stopped at a `longjmp` breakpoint. It takes a `CORE_ADDR *` as argument, and stores the target PC value through this pointer. It examines the current state of the machine as needed.

GET_SAVED_REGISTER

Define this if you need to supply your own definition for the function `get_saved_register`.

HAVE_REGISTER_WINDOWS

Define this if the target has register windows.

REGISTER_IN_WINDOW_P (*regnum*)

Define this to be an expression that is 1 if the given register is in the window.

IBM6000_TARGET

Shows that we are configured for an IBM RS/6000 target. This conditional should be eliminated (FIXME) and replaced by feature-specific macros. It was introduced in a haste and we are repenting at leisure.

I386_USE_GENERIC_WATCHPOINTS

An x86-based target can define this to use the generic x86 watchpoint support; see [Chapter 3 \[Algorithms\]](#), page 2.

SYMBOLS_CAN_START_WITH_DOLLAR

Some systems have routines whose names start with '\$'. Giving this macro a non-zero value tells GDB's expression parser to check for such routines when parsing tokens that begin with '\$'.

On HP-UX, certain system routines (millicode) have names beginning with '\$' or '\$\$'. For example, `$$dyncall` is a millicode routine that handles inter-space procedure calls on PA-RISC.

IEEE_FLOAT

Define this if the target system uses IEEE-format floating point numbers.

INIT_EXTRA_FRAME_INFO (*fromleaf*, *frame*)

If additional information about the frame is required this should be stored in `frame->extra_info`. Space for `frame->extra_info` is allocated using `frame_obstack_alloc`.

INIT_FRAME_PC (*fromleaf*, *prev*)

This is a C statement that sets the pc of the frame pointed to by *prev*. [By default...]

INNER_THAN (*lhs*, *rhs*)

Returns non-zero if stack address *lhs* is inner than (nearer to the stack top) stack address *rhs*. Define this as `lhs < rhs` if the target's stack grows downward in memory, or `lhs > rhs` if the stack grows upward.

IN_SIGTRAMP (*pc*, *name*)

Define this to return non-zero if the given *pc* and/or *name* indicates that the current function is a `sigtramp`.

SIGTRAMP_START (*pc*)

SIGTRAMP_END (*pc*)

Define these to be the start and end address of the `sigtramp` for the given *pc*. On machines where the address is just a compile time constant, the macro expansion will typically just ignore the supplied *pc*.

IN_SOLIB_CALL_TRAMPOLINE (*pc*, *name*)

Define this to evaluate to nonzero if the program is stopped in the trampoline that connects to a shared library.

IN_SOLIB_RETURN_TRAMPOLINE (*pc*, *name*)

Define this to evaluate to nonzero if the program is stopped in the trampoline that returns from a shared library.

IN_SOLIB_DYNSYM_RESOLVE_CODE (*pc*)

Define this to evaluate to nonzero if the program is stopped in the dynamic linker.

SKIP_SOLIB_RESOLVER (*pc*)

Define this to evaluate to the (nonzero) address at which execution should continue to get past the dynamic linker's symbol resolution function. A zero value indicates that it is not important or necessary to set a breakpoint to get through the dynamic linker and that single stepping will suffice.

IS_TRAPPED_INTERNALVAR (*name*)

This is an ugly hook to allow the specification of special actions that should occur as a side-effect of setting the value of a variable internal to GDB. Currently only used by the h8500. Note that this could be either a host or target conditional.

NEED_TEXT_START_END

Define this if GDB should determine the start and end addresses of the text section. (Seems dubious.)

NO_HIF_SUPPORT

(Specific to the a29k.)

POINTER_TO_ADDRESS (*type*, *buf*)

Assume that *buf* holds a pointer of type *type*, in the appropriate format for the current architecture. Return the byte address the pointer refers to. See [Chapter 9 \[Pointers Are Not Always Addresses\]](#), page 33.

REGISTER_CONVERTIBLE (*reg*)

Return non-zero if *reg* uses different raw and virtual formats. See [Chapter 9 \[Using Different Register and Memory Data Representations\]](#), page 33.

REGISTER_RAW_SIZE (*reg*)

Return the raw size of *reg*. See [Chapter 9 \[Using Different Register and Memory Data Representations\]](#), page 33.

REGISTER_VIRTUAL_SIZE (*reg*)

Return the virtual size of *reg*. See [Chapter 9 \[Using Different Register and Memory Data Representations\]](#), page 33.

REGISTER_VIRTUAL_TYPE(*reg*)

Return the virtual type of *reg*. See Chapter 9 [Using Different Register and Memory Data Representations], page 33.

REGISTER_CONVERT_TO_VIRTUAL(*reg*, *type*, *from*, *to*)

Convert the value of register *reg* from its raw form to its virtual form. See Chapter 9 [Using Different Register and Memory Data Representations], page 33.

REGISTER_CONVERT_TO_RAW(*type*, *reg*, *from*, *to*)

Convert the value of register *reg* from its virtual form to its raw form. See Chapter 9 [Using Different Register and Memory Data Representations], page 33.

RETURN_VALUE_ON_STACK(*type*)

Return non-zero if values of type *TYPE* are returned on the stack, using the “struct convention” (i.e., the caller provides a pointer to a buffer in which the callee should store the return value). This controls how the ‘*finish*’ command finds a function’s return value, and whether an inferior function call reserves space on the stack for the return value.

The full logic GDB uses here is kind of odd.

- If the type being returned by value is not a structure, union, or array, and **RETURN_VALUE_ON_STACK** returns zero, then GDB concludes the value is not returned using the struct convention.
- Otherwise, GDB calls **USE_STRUCT_CONVENTION** (see below). If that returns non-zero, GDB assumes the struct convention is in use.

In other words, to indicate that a given type is returned by value using the struct convention, that type must be either a struct, union, array, or something **RETURN_VALUE_ON_STACK** likes, *and* something that **USE_STRUCT_CONVENTION** likes.

Note that, in C and C++, arrays are never returned by value. In those languages, these predicates will always see a pointer type, never an array type. All the references above to arrays being returned by value apply only to other languages.

SOFTWARE_SINGLE_STEP_P()

Define this as 1 if the target does not have a hardware single-step mechanism. The macro **SOFTWARE_SINGLE_STEP** must also be defined.

SOFTWARE_SINGLE_STEP(*signal*, *insert-breapoints-p*)

A function that inserts or removes (depending on *insert-breapoints-p*) breakpoints at each possible destinations of the next instruction. See ‘*sparc-tdep.c*’ and ‘*rs6000-tdep.c*’ for examples.

SOFUN_ADDRESS_MAYBE_MISSING

Somebody clever observed that, the more actual addresses you have in the debug information, the more time the linker has to spend relocating them. So whenever there’s some other way the debugger could find the address it needs, you should omit it from the debug info, to make linking faster.

SOFUN_ADDRESS_MAYBE_MISSING indicates that a particular set of hacks of this sort are in use, affecting **N_SO** and **N_FUN** entries in stabs-format debugging in-

formation. `N_SO` stabs mark the beginning and ending addresses of compilation units in the text segment. `N_FUN` stabs mark the starts and ends of functions.

`SOFUN_ADDRESS_MAYBE_MISSING` means two things:

- `N_FUN` stabs have an address of zero. Instead, you should find the addresses where the function starts by taking the function name from the stab, and then looking that up in the minsyms (the linker/assembler symbol table). In other words, the stab has the name, and the linker/assembler symbol table is the only place that carries the address.
- `N_SO` stabs have an address of zero, too. You just look at the `N_FUN` stabs that appear before and after the `N_SO` stab, and guess the starting and ending addresses of the compilation unit from them.

`PCC_SOL_BROKEN`

(Used only in the Convex target.)

`PC_IN_CALL_DUMMY`

See ‘`inferior.h`’.

`PC_LOAD_SEGMENT`

If defined, print information about the load segment for the program counter. (Defined only for the RS/6000.)

`PC_REGNUM`

If the program counter is kept in a register, then define this macro to be the number (greater than or equal to zero) of that register.

This should only need to be defined if `TARGET_READ_PC` and `TARGET_WRITE_PC` are not defined.

`NPC_REGNUM`

The number of the “next program counter” register, if defined.

`NNPC_REGNUM`

The number of the “next next program counter” register, if defined. Currently, this is only defined for the Motorola 88K.

`PARAM_BOUNDARY`

If non-zero, round arguments to a boundary of this many bits before pushing them on the stack.

`PRINT_REGISTER_HOOK` (*regno*)

If defined, this must be a function that prints the contents of the given register to standard output.

`PRINT_TYPELESS_INTEGER`

This is an obscure substitute for `print_longest` that seems to have been defined for the Convex target.

`PROCESS_LINENUMBER_HOOK`

A hook defined for XCOFF reading.

`PROLOGUE_FIRSTLINE_OVERLAP`

(Only used in unsupported Convex configuration.)

PS_REGNUM

If defined, this is the number of the processor status register. (This definition is only used in generic code when parsing "\$ps".)

POP_FRAME

Used in 'call_function_by_hand' to remove an artificial stack frame and in 'return_command' to remove a real stack frame.

PUSH_ARGUMENTS (*nargs, args, sp, struct_return, struct_addr*)

Define this to push arguments onto the stack for inferior function call. Returns the updated stack pointer value.

PUSH_DUMMY_FRAME

Used in 'call_function_by_hand' to create an artificial stack frame.

REGISTER_BYTES

The total amount of space needed to store GDB's copy of the machine's register state.

REGISTER_NAME(*i*)

Return the name of register *i* as a string. May return NULL or NUL to indicate that register *i* is not valid.

REGISTER_NAMES

Deprecated in favor of REGISTER_NAME.

REG_STRUCT_HAS_ADDR (*gcc_p, type*)

Define this to return 1 if the given type will be passed by pointer rather than directly.

SAVE_DUMMY_FRAME_TOS (*sp*)

Used in 'call_function_by_hand' to notify the target dependent code of the top-of-stack value that will be passed to the the inferior code. This is the value of the SP after both the dummy frame and space for parameters/results have been allocated on the stack.

SDB_REG_TO_REGNUM

Define this to convert sdb register numbers into GDB regnums. If not defined, no conversion will be done.

SHIFT_INST_REGS

(Only used for m88k targets.)

SKIP_PERMANENT_BREAKPOINT

Advance the inferior's PC past a permanent breakpoint. GDB normally steps over a breakpoint by removing it, stepping one instruction, and re-inserting the breakpoint. However, permanent breakpoints are hardwired into the inferior, and can't be removed, so this strategy doesn't work. Calling SKIP_PERMANENT_BREAKPOINT adjusts the processor's state so that execution will resume just after the breakpoint. This macro does the right thing even when the breakpoint is in the delay slot of a branch or jump.

SKIP_PROLOGUE (*pc*)

A C expression that returns the address of the "real" code beyond the function entry prologue found at *pc*.

SKIP_PROLOGUE_FRAMELESS_P

A C expression that should behave similarly, but that can stop as soon as the function is known to have a frame. If not defined, `SKIP_PROLOGUE` will be used instead.

SKIP_TRAMPOLINE_CODE (*pc*)

If the target machine has trampoline code that sits between callers and the functions being called, then define this macro to return a new PC that is at the start of the real function.

SP_REGNUM

If the stack-pointer is kept in a register, then define this macro to be the number (greater than or equal to zero) of that register.

This should only need to be defined if `TARGET_WRITE_SP` and `TARGET_WRITE_SP` are not defined.

STAB_REG_TO_REGNUM

Define this to convert stab register numbers (as gotten from ‘r’ declarations) into GDB regnums. If not defined, no conversion will be done.

STACK_ALIGN (*addr*)

Define this to adjust the address to the alignment required for the processor’s stack.

STEP_SKIPS_DELAY (*addr*)

Define this to return true if the address is of an instruction with a delay slot. If a breakpoint has been placed in the instruction’s delay slot, GDB will single-step over that instruction before resuming normally. Currently only defined for the Mips.

STORE_RETURN_VALUE (*type*, *valbuf*)

A C expression that stores a function return value of type *type*, where *valbuf* is the address of the value to be stored.

SUN_FIXED_LBRAC_BUG

(Used only for Sun-3 and Sun-4 targets.)

SYMBOL_RELOADING_DEFAULT

The default value of the “symbol-reloading” variable. (Never defined in current sources.)

TARGET_BYTE_ORDER_DEFAULT

The ordering of bytes in the target. This must be either `BIG_ENDIAN` or `LITTLE_ENDIAN`. This macro replaces `TARGET_BYTE_ORDER` which is deprecated.

TARGET_BYTE_ORDER_SELECTABLE_P

Non-zero if the target has both `BIG_ENDIAN` and `LITTLE_ENDIAN` variants. This macro replaces `TARGET_BYTE_ORDER_SELECTABLE` which is deprecated.

TARGET_CHAR_BIT

Number of bits in a char; defaults to 8.

TARGET_COMPLEX_BIT

Number of bits in a complex number; defaults to `2 * TARGET_FLOAT_BIT`.

At present this macro is not used.

TARGET_DOUBLE_BIT

Number of bits in a double float; defaults to $8 * \text{TARGET_CHAR_BIT}$.

TARGET_DOUBLE_COMPLEX_BIT

Number of bits in a double complex; defaults to $2 * \text{TARGET_DOUBLE_BIT}$.

At present this macro is not used.

TARGET_FLOAT_BIT

Number of bits in a float; defaults to $4 * \text{TARGET_CHAR_BIT}$.

TARGET_INT_BIT

Number of bits in an integer; defaults to $4 * \text{TARGET_CHAR_BIT}$.

TARGET_LONG_BIT

Number of bits in a long integer; defaults to $4 * \text{TARGET_CHAR_BIT}$.

TARGET_LONG_DOUBLE_BIT

Number of bits in a long double float; defaults to $2 * \text{TARGET_DOUBLE_BIT}$.

TARGET_LONG_LONG_BIT

Number of bits in a long long integer; defaults to $2 * \text{TARGET_LONG_BIT}$.

TARGET_PTR_BIT

Number of bits in a pointer; defaults to TARGET_INT_BIT .

TARGET_SHORT_BIT

Number of bits in a short integer; defaults to $2 * \text{TARGET_CHAR_BIT}$.

TARGET_READ_PC

TARGET_WRITE_PC (*val*, *pid*)

TARGET_READ_SP

TARGET_WRITE_SP

TARGET_READ_FP

TARGET_WRITE_FP

These change the behavior of `read_pc`, `write_pc`, `read_sp`, `write_sp`, `read_fp` and `write_fp`. For most targets, these may be left undefined. GDB will call the read and write register functions with the relevant `_REGNUM` argument.

These macros are useful when a target keeps one of these registers in a hard to get at place; for example, part in a segment register and part in an ordinary register.

TARGET_VIRTUAL_FRAME_POINTER(*pc*, *regp*, *offsetp*)

Returns a (`register`, `offset`) pair representing the virtual frame pointer in use at the code address *pc*. If virtual frame pointers are not used, a default definition simply returns `FP_REGNUM`, with an offset of zero.

TARGET_HAS_HARDWARE_WATCHPOINTS

If non-zero, the target has support for hardware-assisted watchpoints. See [Chapter 3 \[Algorithms\]](#), [page 2](#), for more details and other related macros.

USE_STRUCT_CONVENTION (*gcc-p*, *type*)

If defined, this must be an expression that is nonzero if a value of the given *type* being returned from a function must have space allocated for it on the stack.

gcc-p is true if the function being considered is known to have been compiled by GCC; this is helpful for systems where GCC is known to use different calling convention than other compilers.

VARIABLES_INSIDE_BLOCK (*desc*, *gcc-p*)

For dbx-style debugging information, if the compiler puts variable declarations inside LBRAC/RBRAC blocks, this should be defined to be nonzero. *desc* is the value of *n_desc* from the *N_RBRAC* symbol, and *gcc-p* is true if GDB has noticed the presence of either the *GCC_COMPILED_SYMBOL* or the *GCC2_COMPILED_SYMBOL*. By default, this is 0.

OS9K_VARIABLES_INSIDE_BLOCK (*desc*, *gcc-p*)

Similarly, for OS/9000. Defaults to 1.

Motorola M68K target conditionals.

BPT_VECTOR

Define this to be the 4-bit location of the breakpoint trap vector. If not defined, it will default to 0xf.

REMOTE_BPT_VECTOR

Defaults to 1.

9.8 Adding a New Target

The following files add a target to GDB:

‘gdb/config/arch/ttt.mt’

Contains a Makefile fragment specific to this target. Specifies what object files are needed for target *ttt*, by defining ‘*TDEPFILES=...*’ and ‘*TDEPLIBS=...*’. Also specifies the header file which describes *ttt*, by defining ‘*TM_FILE=tm-ttt.h*’.

You can also define ‘*TM_CFLAGS*’, ‘*TM_CLIBS*’, ‘*TM_CDEPS*’, but these are now deprecated, replaced by *autoconf*, and may go away in future versions of GDB.

‘gdb/ttt-tdep.c’

Contains any miscellaneous code required for this target machine. On some machines it doesn’t exist at all. Sometimes the macros in ‘*tm-ttt.h*’ become very complicated, so they are implemented as functions here instead, and the macro is simply defined to call the function. This is vastly preferable, since it is easier to understand and debug.

‘gdb/arch-tdep.c’

‘gdb/arch-tdep.h’

This often exists to describe the basic layout of the target machine’s processor chip (registers, stack, etc.). If used, it is included by ‘*ttt-tdep.h*’. It can be shared among many targets that use the same processor.

‘gdb/config/arch/tm-ttt.h’

(‘*tm.h*’ is a link to this file, created by *configure*). Contains macro definitions about the target machine’s registers, stack frame format and instructions.

New targets do not need this file and should not create it.

`'gdb/config/arch/tm-arch.h'`

This often exists to describe the basic layout of the target machine's processor chip (registers, stack, etc.). If used, it is included by `'tm-ttt.h'`. It can be shared among many targets that use the same processor.

New targets do not need this file and should not create it.

If you are adding a new operating system for an existing CPU chip, add a `'config/tm-os.h'` file that describes the operating system facilities that are unusual (extra symbol table info; the breakpoint instruction needed; etc.). Then write a `'arch/tm-os.h'` that just `#includes` `'tm-arch.h'` and `'config/tm-os.h'`.

10 Target Vector Definition

The target vector defines the interface between GDB's abstract handling of target systems, and the nitty-gritty code that actually exercises control over a process or a serial port. GDB includes some 30-40 different target vectors; however, each configuration of GDB includes only a few of them.

10.1 File Targets

Both executables and core files have target vectors.

10.2 Standard Protocol and Remote Stubs

GDB's file `'remote.c'` talks a serial protocol to code that runs in the target system. GDB provides several sample *stubs* that can be integrated into target programs or operating systems for this purpose; they are named `'*-stub.c'`.

The GDB user's manual describes how to put such a stub into your target code. What follows is a discussion of integrating the SPARC stub into a complicated operating system (rather than a simple program), by Stu Grossman, the author of this stub.

The trap handling code in the stub assumes the following upon entry to `trap_low`:

1. `%l1` and `%l2` contain pc and npc respectively at the time of the trap;
2. traps are disabled;
3. you are in the correct trap window.

As long as your trap handler can guarantee those conditions, then there is no reason why you shouldn't be able to "share" traps with the stub. The stub has no requirement that it be jumped to directly from the hardware trap vector. That is why it calls `exceptionHandler()`, which is provided by the external environment. For instance, this could set up the hardware traps to actually execute code which calls the stub first, and then transfers to its own trap handler.

For the most part, there probably won't be much of an issue with "sharing" traps, as the traps we use are usually not used by the kernel, and often indicate unrecoverable error conditions. Anyway, this is all controlled by a table, and is trivial to modify. The most important trap for us is for `ta 1`. Without that, we can't single step or do breakpoints. Everything else is unnecessary for the proper operation of the debugger/stub.

From reading the stub, it's probably not obvious how breakpoints work. They are simply done by deposit/examine operations from GDB.

10.3 ROM Monitor Interface

10.4 Custom Protocols

10.5 Transport Layer

10.6 Builtin Simulator

11 Native Debugging

Several files control GDB's configuration for native support:

`'gdb/config/arch/xyz.mh'`

Specifies Makefile fragments needed when hosting *or native* on machine xyz. In particular, this lists the required native-dependent object files, by defining `'NATDEPFILES=...'`. Also specifies the header file which describes native support on xyz, by defining `'NAT_FILE= nm-xyz.h'`. You can also define `'NAT_CFLAGS'`, `'NAT_ADD_FILES'`, `'NAT_CLIBS'`, `'NAT_CDEPS'`, etc.; see `'Makefile.in'`.

`'gdb/config/arch/nm-xyz.h'`

(`'nm.h'` is a link to this file, created by `configure`). Contains C macro definitions describing the native system environment, such as child process control and core file support.

`'gdb/xyz-nat.c'`

Contains any miscellaneous C code required for this native support of this machine. On some machines it doesn't exist at all.

There are some “generic” versions of routines that can be used by various systems. These can be customized in various ways by macros defined in your `'nm-xyz.h'` file. If these routines work for the xyz host, you can just include the generic file's name (with `' .o'`, not `' .c'`) in `NATDEPFILES`.

Otherwise, if your machine needs custom support routines, you will need to write routines that perform the same functions as the generic file. Put them into `'xyz-nat.c'`, and put `'xyz-nat.o'` into `NATDEPFILES`.

`'inftarg.c'`

This contains the *target_ops* vector that supports Unix child processes on systems which use `ptrace` and wait to control the child.

`'procfs.c'`

This contains the *target_ops* vector that supports Unix child processes on systems which use `/proc` to control the child.

`'fork-child.c'`

This does the low-level grunge that uses Unix system calls to do a “fork and exec” to start up a child process.

`'infptrace.c'`

This is the low level interface to inferior processes for systems using the Unix `ptrace` call in a vanilla way.

11.1 Native core file Support

`'core-aout.c::fetch_core_registers()'`

Support for reading registers out of a core file. This routine calls `register_addr()`, see below. Now that BFD is used to read core files, virtually all machines should use `core-aout.c`, and should just provide `fetch_core_registers` in `xyz-nat.c` (or `REGISTER_U_ADDR` in `nm-xyz.h`).

`'core-aout.c::register_addr()'`

If your `nm-xyz.h` file defines the macro `REGISTER_U_ADDR(addr, blockend, regno)`, it should be defined to set `addr` to the offset within the ‘user’ struct of GDB register number `regno`. `blockend` is the offset within the “upage” of `u.u_ar0`. If `REGISTER_U_ADDR` is defined, ‘core-aout.c’ will define the `register_addr()` function and use the macro in it. If you do not define `REGISTER_U_ADDR`, but you are using the standard `fetch_core_registers()`, you will need to define your own version of `register_addr()`, put it into your `xyz-nat.c` file, and be sure `xyz-nat.o` is in the `NATDEPFILES` list. If you have your own `fetch_core_registers()`, you may not need a separate `register_addr()`. Many custom `fetch_core_registers()` implementations simply locate the registers themselves.

When making GDB run native on a new operating system, to make it possible to debug core files, you will need to either write specific code for parsing your OS’s core files, or customize ‘`bfd/trad-core.c`’. First, use whatever `#include` files your machine uses to define the struct of registers that is accessible (possibly in the u-area) in a core file (rather than ‘`machine/reg.h`’), and an include file that defines whatever header exists on a core file (e.g. the u-area or a `struct core`). Then modify `trad_unix_core_file_p` to use these values to set up the section information for the data segment, stack segment, any other segments in the core file (perhaps shared library contents or control information), “registers” segment, and if there are two discontinuous sets of registers (e.g. integer and float), the “reg2” segment. This section information basically delimits areas in the core file in a standard way, which the section-reading routines in BFD know how to seek around in.

Then back in GDB, you need a matching routine called `fetch_core_registers`. If you can use the generic one, it’s in ‘`core-aout.c`’; if not, it’s in your ‘`xyz-nat.c`’ file. It will be passed a char pointer to the entire “registers” segment, its length, and a zero; or a char pointer to the entire “regs2” segment, its length, and a 2. The routine should suck out the supplied register values and install them into GDB’s “registers” array.

If your system uses `/proc` to control processes, and uses ELF format core files, then you may be able to use the same routines for reading the registers out of processes and out of core files.

11.2 ptrace

11.3 /proc

11.4 win32

11.5 shared libraries

11.6 Native Conditionals

When GDB is configured and compiled, various macros are defined or left undefined, to control compilation when the host and target systems are the same. These macros should be defined (or left undefined) in `nm-system.h`.

ATTACH_DETACH

If defined, then GDB will include support for the `attach` and `detach` commands.

CHILD_PREPARE_TO_STORE

If the machine stores all registers at once in the child process, then define this to ensure that all values are correct. This usually entails a read from the child. [Note that this is incorrectly defined in `xm-system.h` files currently.]

FETCH_INFERIOR_REGISTERS

Define this if the native-dependent code will provide its own routines `fetch_inferior_registers` and `store_inferior_registers` in `host-nat.c`. If this symbol is *not* defined, and `infptrace.c` is included in this configuration, the default routines in `infptrace.c` are used for these functions.

FILES_INFO_HOOK

(Only defined for Convex.)

FPO_REGNUM

This macro is normally defined to be the number of the first floating point register, if the machine has such registers. As such, it would appear only in target-specific code. However, `/proc` support uses this to decide whether floats are in use on this target.

GET_LONGJMP_TARGET

For most machines, this is a target-dependent parameter. On the DECstation and the Iris, this is a native-dependent parameter, since `setjmp.h` is needed to define it.

This macro determines the target PC address that `longjmp` will jump to, assuming that we have just stopped at a `longjmp` breakpoint. It takes a `CORE_ADDR *` as argument, and stores the target PC value through this pointer. It examines the current state of the machine as needed.

I386_USE_GENERIC_WATCHPOINTS

An x86-based machine can define this to use the generic x86 watchpoint support; see [Chapter 3 \[Algorithms\]](#), page 2.

KERNEL_U_ADDR

Define this to the address of the `u` structure (the “user struct”, also known as the “u-page”) in kernel virtual memory. GDB needs to know this so that it can subtract this address from absolute addresses in the upage, that are obtained via `ptrace` or from core files. On systems that don’t need this value, set it to zero.

KERNEL_U_ADDR_BSD

Define this to cause GDB to determine the address of `u` at runtime, by using Berkeley-style `nlist` on the kernel’s image in the root directory.

KERNEL_U_ADDR_HPUX

Define this to cause GDB to determine the address of `u` at runtime, by using HP-style `nlist` on the kernel’s image in the root directory.

ONE_PROCESS_WRITETEXT

Define this to be able to, when a breakpoint insertion fails, warn the user that another process may be running with the same executable.

PREPARE_TO_PROCEED (*select_it*)

This (ugly) macro allows a native configuration to customize the way the `proceed` function in ‘`infrun.c`’ deals with switching between threads.

In a multi-threaded task we may select another thread and then continue or step. But if the old thread was stopped at a breakpoint, it will immediately cause another breakpoint stop without any execution (i.e. it will report a breakpoint hit incorrectly). So GDB must step over it first.

If defined, `PREPARE_TO_PROCEED` should check the current thread against the thread that reported the most recent event. If a step-over is required, it returns `TRUE`. If *select_it* is non-zero, it should reselect the old thread.

PROC_NAME_FMT

Defines the format for the name of a ‘`/proc`’ device. Should be defined in ‘`nm.h`’ *only* in order to override the default definition in ‘`procfs.c`’.

PTRACE_FP_BUG

See ‘`mach386-xdep.c`’.

PTRACE_ARG3_TYPE

The type of the third argument to the `ptrace` system call, if it exists and is different from `int`.

REGISTER_U_ADDR

Defines the offset of the registers in the “u area”.

SHELL_COMMAND_CONCAT

If defined, is a string to prefix on the shell command used to start the inferior.

SHELL_FILE

If defined, this is the name of the shell to use to run the inferior. Defaults to `"/bin/sh"`.

SOLIB_ADD (*filename*, *from_tty*, *targ*)

Define this to expand into an expression that will cause the symbols in *filename* to be added to GDB's symbol table.

SOLIB_CREATE_INFERIOR_HOOK

Define this to expand into any shared-library-relocation code that you want to be run just after the child process has been forked.

START_INFERIOR_TRAPS_EXPECTED

When starting an inferior, GDB normally expects to trap twice; once when the shell execs, and once when the program itself execs. If the actual number of traps is something other than 2, then define this macro to expand into the number expected.

SVR4_SHARED_LIBS

Define this to indicate that SVR4-style shared libraries are in use.

USE_PROC_FS

This determines whether small routines in `'*-tdep.c'`, which translate register values between GDB's internal representation and the `'/proc'` representation, are compiled.

U_REGS_OFFSET

This is the offset of the registers in the upage. It need only be defined if the generic ptrace register access routines in `'infptrace.c'` are being used (that is, `'infptrace.c'` is configured in, and `FETCH_INFERIOR_REGISTERS` is not defined). If the default value from `'infptrace.c'` is good enough, leave it undefined.

The default value means that `u.u_ar0` *points to* the location of the registers. I'm guessing that `#define U_REGS_OFFSET 0` means that `u.u_ar0` *is* the location of the registers.

CLEAR_SOLIB

See `'objfiles.c'`.

DEBUG_PTRACE

Define this to debug `ptrace` calls.

12 Support Libraries

12.1 BFD

BFD provides support for GDB in several ways:

identifying executable and core files

BFD will identify a variety of file types, including a.out, coff, and several variants thereof, as well as several kinds of core files.

access to sections of files

BFD parses the file headers to determine the names, virtual addresses, sizes, and file locations of all the various named sections in files (such as the text section or the data section). GDB simply calls BFD to read or write section *x* at byte offset *y* for length *z*.

specialized core file support

BFD provides routines to determine the failing command name stored in a core file, the signal with which the program failed, and whether a core file matches (i.e. could be a core dump of) a particular executable file.

locating the symbol information

GDB uses an internal interface of BFD to determine where to find the symbol information in an executable file or symbol-file. GDB itself handles the reading of symbols, since BFD does not “understand” debug symbols, but GDB uses BFD’s cached information to find the symbols, string table, etc.

12.2 opcodes

The opcodes library provides GDB’s disassembler. (It’s a separate library because it’s also used in binutils, for ‘objdump’).

12.3 readline

12.4 mmalloc

12.5 libiberty

12.6 gnu-regex

Regex conditionals.

C_ALLOCA

NFAILURES

RE_NREGS

SIGN_EXTEND_CHAR

SWITCH_ENUM_BUG

SYNTAX_TABLE

Sword

sparc

12.7 include

13 Coding

This chapter covers topics that are lower-level than the major algorithms of GDB.

13.1 Cleanups

Cleanups are a structured way to deal with things that need to be done later. When your code does something (like `malloc` some memory, or open a file) that needs to be undone later (e.g., free the memory or close the file), it can make a cleanup. The cleanup will be done at some future point: when the command is finished, when an error occurs, or when your code decides it's time to do cleanups.

You can also discard cleanups, that is, throw them away without doing what they say. This is only done if you ask that it be done.

Syntax:

```
struct cleanup *old_chain;
```

Declare a variable which will hold a cleanup chain handle.

```
old_chain = make_cleanup (function, arg);
```

Make a cleanup which will cause *function* to be called with *arg* (a `char *`) later. The result, *old_chain*, is a handle that can be passed to `do_cleanups` or `discard_cleanups` later. Unless you are going to call `do_cleanups` or `discard_cleanups` yourself, you can ignore the result from `make_cleanup`.

```
do_cleanups (old_chain);
```

Perform all cleanups done since `make_cleanup` returned *old_chain*. E.g.:

```
make_cleanup (a, 0);
old = make_cleanup (b, 0);
do_cleanups (old);
```

will call `b()` but will not call `a()`. The cleanup that calls `a()` will remain in the cleanup chain, and will be done later unless otherwise discarded.

```
discard_cleanups (old_chain);
```

Same as `do_cleanups` except that it just removes the cleanups from the chain and does not call the specified functions.

Some functions, e.g. `fputs_filtered()` or `error()`, specify that they “should not be called when cleanups are not in place”. This means that any actions you need to reverse in the case of an error or interruption must be on the cleanup chain before you call these functions, since they might never return to your code (they ‘`longjmp`’ instead).

13.2 Wrapping Output Lines

Output that goes through `printf_filtered` or `fputs_filtered` or `fputs_demangled` needs only to have calls to `wrap_here` added in places that would be good breaking points. The utility routines will take care of actually wrapping if the line width is exceeded.

The argument to `wrap_here` is an indentation string which is printed *only* if the line breaks there. This argument is saved away and used later. It must remain valid until the next call to `wrap_here` or until a newline has been printed through the `*_filtered` functions. Don't pass in a local variable and then return!

It is usually best to call `wrap_here` after printing a comma or space. If you call it before printing a space, make sure that your indentation properly accounts for the leading space that will print if the line wraps there.

Any function or set of functions that produce filtered output must finish by printing a newline, to flush the wrap buffer, before switching to unfiltered (`printf`) output. Symbol reading routines that print warnings are a good example.

13.3 GDB Coding Standards

GDB follows the GNU coding standards, as described in '`etc/standards.texi`'. This file is also available for anonymous FTP from GNU archive sites. GDB takes a strict interpretation of the standard; in general, when the GNU standard recommends a practice but does not require it, GDB requires it.

GDB follows an additional set of coding standards specific to GDB, as described in the following sections.

13.3.1 ISO-C

GDB assumes an ISO-C compliant compiler.

GDB does not assume an ISO-C or POSIX compliant C library.

13.3.2 Memory Management

GDB does not use the functions `malloc`, `realloc`, `calloc`, `free` and `asprintf`.

GDB uses the functions `xmalloc`, `xrealloc` and `xcalloc` when allocating memory. Unlike `malloc` et.al. these functions do not return when the memory pool is empty. Instead, they unwind the stack using cleanups. These functions return `NULL` when requested to allocate a chunk of memory of size zero.

Pragmatics: By using these functions, the need to check every memory allocation is removed. These functions provide portable behavior.

GDB does not use the function `free`.

GDB uses the function `xfree` to return memory to the memory pool. Consistent with ISO-C, this function ignores a request to free a `NULL` pointer.

Pragmatics: On some systems `free` fails when passed a `NULL` pointer.

GDB can use the non-portable function `alloca` for the allocation of small temporary values (such as strings).

Pragmatics: This function is very non-portable. Some systems restrict the memory being allocated to no more than a few kilobytes.

GDB uses the string function `xstrdup` and the print function `xasprintf`.

Pragmatics: `asprintf` and `strdup` can fail. Print functions such as `sprintf` are very prone to buffer overflow errors.

13.3.3 Compiler Warnings

With few exceptions, developers should include the configuration option `--enable-gdb-build-warnings=-Werror` when building GDB. The exceptions are listed in the file `'gdb/MAINTAINERS'`.

This option causes GDB (when built using GCC) to be compiled with a carefully selected list of compiler warning flags. Any warnings from those flags being treated as errors.

The current list of warning flags includes:

`'-Wimplicit'`

Since GDB coding standard requires all functions to be declared using a prototype, the flag has the side effect of ensuring that prototyped functions are always visible with out resorting to `'-Wstrict-prototypes'`.

`'-Wreturn-type'`

Such code often appears to work except on instruction set architectures that use register windows.

`'-Wcomment'`

`'-Wtrigraphs'`

`'-Wformat'`

Since GDB uses the `format printf` attribute on all `printf` like functions this checks not just `printf` calls but also calls to functions such as `fprintf_unfiltered`.

`'-Wparentheses'`

This warning includes uses of the assignment operator within an `if` statement.

`'-Wpointer-arith'`

`'-Wuninitialized'`

Pragmatics: Due to the way that GDB is implemented most functions have unused parameters. Consequently the warning `'-Wunused-parameter'` is precluded from the list. The macro `ATTRIBUTE_UNUSED` is not used as it leads to false negatives — it is not an error to have `ATTRIBUTE_UNUSED` on a parameter that is being used. The options `'-Wall'` and `'-Wunused'` are also precluded because they both include `'-Wunused-parameter'`.

Pragmatics: GDB has not simply accepted the warnings enabled by `'-Wall -Werror -W...'`. Instead it is selecting warnings when and where their benefits can be demonstrated.

13.3.4 Formatting

The standard GNU recommendations for formatting must be followed strictly.

A function declaration should not have its name in column zero. A function definition should have its name in column zero.

```

/* Declaration */
static void foo (void);
/* Definition */
void
foo (void)
{
}

```

Pragmatics: This simplifies scripting. Function definitions can be found using ‘function-name’.

There must be a space between a function or macro name and the opening parenthesis of its argument list (except for macro definitions, as required by C). There must not be a space after an open paren/bracket or before a close paren/bracket.

While additional whitespace is generally helpful for reading, do not use more than one blank line to separate blocks, and avoid adding whitespace after the end of a program line (as of 1/99, some 600 lines had whitespace after the semicolon). Excess whitespace causes difficulties for `diff` and `patch` utilities.

Pointers are declared using the traditional K&R C style:

```
void *foo;
```

and not:

```
void * foo;
void* foo;
```

13.3.5 Comments

The standard GNU requirements on comments must be followed strictly.

Block comments must appear in the following form, with no `/*-` or `*/-` only lines, and no leading `*`:

```

/* Wait for control to return from inferior to debugger.  If inferior
   gets a signal, we may decide to start it up again instead of
   returning.  That is why there is a loop in this function.  When
   this function actually returns it means the inferior should be left
   stopped and GDB should read more commands.  */

```

(Note that this format is encouraged by Emacs; tabbing for a multi-line comment works correctly, and `M-q` fills the block consistently.)

Put a blank line between the block comments preceding function or variable definitions, and the definition itself.

In general, put function-body comments on lines by themselves, rather than trying to fit them into the 20 characters left at the end of a line, since either the comment or the code will inevitably get longer than will fit, and then somebody will have to move it anyhow.

13.3.6 C Usage

Code must not depend on the sizes of C data types, the format of the host’s floating point numbers, the alignment of anything, or the order of evaluation of expressions.

Use functions freely. There are only a handful of compute-bound areas in GDB that might be affected by the overhead of a function call, mainly in symbol reading. Most of GDB's performance is limited by the target interface (whether serial line or system call).

However, use functions with moderation. A thousand one-line functions are just as hard to understand as a single thousand-line function.

Macros are bad, M'kay. (But if you have to use a macro, make sure that the macro arguments are protected with parentheses.)

Declarations like `'struct foo *'` should be used in preference to declarations like `'typedef struct foo { ... } *foo_ptr'`.

13.3.7 Function Prototypes

Prototypes must be used when both *declaring* and *defining* a function. Prototypes for GDB functions must include both the argument type and name, with the name matching that used in the actual function definition.

All external functions should have a declaration in a header file that callers include, except for `_initialize_*` functions, which must be external so that `'init.c'` construction works, but shouldn't be visible to random source files.

Where a source file needs a forward declaration of a static function, that declaration must appear in a block near the top of the source file.

13.3.8 Internal Error Recovery

During its execution, GDB can encounter two types of errors. User errors and internal errors. User errors include not only a user entering an incorrect command but also problems arising from corrupt object files and system errors when interacting with the target. Internal errors include situations where GDB has detected, at run time, a corrupt or erroneous situation.

When reporting an internal error, GDB uses `internal_error` and `gdb_assert`.

GDB must not call `abort` or `assert`.

Pragmatics: There is no `internal_warning` function. Either the code detected a user error, recovered from it and issued a `warning` or the code failed to correctly recover from the user error and issued an `internal_error`.

13.3.9 File Names

Any file used when building the core of GDB must be in lower case. Any file used when building the core of GDB must be 8.3 unique. These requirements apply to both source and generated files.

Pragmatics: The core of GDB must be buildable on many platforms including DJGPP and MacOS/HFS. Every time an unfriendly file is introduced to the build process both `'Makefile.in'` and `'configure.in'` need to be modified accordingly. Compare the convoluted conversion process needed to transform `'COPYING'` into `'copying.c'` with the conversion needed to transform `'version.in'` into `'version.c'`.

Any file non 8.3 compliant file (that is not used when building the core of GDB) must be added to `'gdb/config/djgpp/fnchange.lst'`.

Pragmatics: This is clearly a compromise.

When GDB has a local version of a system header file (ex ‘string.h’) the file name based on the POSIX header prefixed with ‘gdb_’ (‘gdb_string.h’).

For other files ‘-’ is used as the separator.

13.3.10 Include Files

All ‘.c’ files should include ‘defs.h’ first.

All ‘.c’ files should explicitly include the headers for any declarations they refer to. They should not rely on files being included indirectly.

With the exception of the global definitions supplied by ‘defs.h’, a header file should explicitly include the header declaring any `typedefs` et.al. it refers to.

`extern` declarations should never appear in .c files.

All include files should be wrapped in:

```
#ifndef INCLUDE_FILE_NAME_H
#define INCLUDE_FILE_NAME_H
header body
#endif
```

13.3.11 Clean Design and Portable Implementation

In addition to getting the syntax right, there’s the little question of semantics. Some things are done in certain ways in GDB because long experience has shown that the more obvious ways caused various kinds of trouble.

You can’t assume the byte order of anything that comes from a target (including *values*, object files, and instructions). Such things must be byte-swapped using `SWAP_TARGET_AND_HOST` in GDB, or one of the swap routines defined in ‘bfd.h’, such as `bfd_get_32`.

You can’t assume that you know what interface is being used to talk to the target system. All references to the target must go through the current `target_ops` vector.

You can’t assume that the host and target machines are the same machine (except in the “native” support modules). In particular, you can’t assume that the target machine’s header files will be available on the host machine. Target code must bring along its own header files – written from scratch or explicitly donated by their owner, to avoid copyright problems.

Insertion of new `#ifdef`’s will be frowned upon. It’s much better to write the code portably than to conditionalize it for various systems.

New `#ifdef`’s which test for specific compilers or manufacturers or operating systems are unacceptable. All `#ifdef`’s should test for features. The information about which configurations contain which features should be segregated into the configuration files. Experience has proven far too often that a feature unique to one particular system often creeps into other systems; and that a conditional based on some predefined macro for your current system will become worthless over time, as new versions of your system come out that behave differently with regard to this feature.

Adding code that handles specific architectures, operating systems, target interfaces, or hosts, is not acceptable in generic code.

One particularly notorious area where system dependencies tend to creep in is handling of file names. The mainline GDB code assumes Posix semantics of file names: absolute file names begin with a forward slash '/', slashes are used to separate leading directories, case-sensitive file names. These assumptions are not necessarily true on non-Posix systems such as MS-Windows. To avoid system-dependent code where you need to take apart or construct a file name, use the following portable macros:

HAVE_DOS_BASED_FILE_SYSTEM

This preprocessing symbol is defined to a non-zero value on hosts whose filesystems belong to the MS-DOS/MS-Windows family. Use this symbol to write conditional code which should only be compiled for such hosts.

IS_DIR_SEPARATOR (*c*)

Evaluates to a non-zero value if *c* is a directory separator character. On Unix and GNU/Linux systems, only a slash '/' is such a character, but on Windows, both '/' and '\' will pass.

IS_ABSOLUTE_PATH (*file*)

Evaluates to a non-zero value if *file* is an absolute file name. For Unix and GNU/Linux hosts, a name which begins with a slash '/' is absolute. On DOS and Windows, 'd:/foo' and 'x:\bar' are also absolute file names.

FILENAME_CMP (*f1*, *f2*)

Calls a function which compares file names *f1* and *f2* as appropriate for the underlying host filesystem. For Posix systems, this simply calls `strcmp`; on case-insensitive filesystems it will call `strcasecmp` instead.

DIRNAME_SEPARATOR

Evaluates to a character which separates directories in PATH-style lists, typically held in environment variables. This character is ':' on Unix, ';' on DOS and Windows.

SLASH_STRING

This evaluates to a constant string you should use to produce an absolute filename from leading directories and the file's basename. `SLASH_STRING` is "/" on most systems, but might be "\\\" for some Windows-based ports.

In addition to using these macros, be sure to use portable library functions whenever possible. For example, to extract a directory or a basename part from a file name, use the `dirname` and `basename` library functions (available in `libiberty` for platforms which don't provide them), instead of searching for a slash with `strchr`.

Another way to generalize GDB along a particular interface is with an attribute struct. For example, GDB has been generalized to handle multiple kinds of remote interfaces—not by `#ifdefs` everywhere, but by defining the `target_ops` structure and having a current target (as well as a stack of targets below it, for memory references). Whenever something needs to be done that depends on which remote interface we are using, a flag in the current `target_ops` structure is tested (e.g., `target_has_stack`), or a function is called through a pointer in the current `target_ops` structure. In this way, when a new remote interface is added, only one module needs to be touched—the one that actually implements the new remote interface. Other examples of attribute-structs are BFD access to multiple kinds of object file formats, or GDB's access to multiple source languages.

Please avoid duplicating code. For example, in GDB 3.x all the code interfacing between `ptrace` and the rest of GDB was duplicated in `*-dep.c`, and so changing something was very painful. In GDB 4.x, these have all been consolidated into `infptrace.c`. `infptrace.c` can deal with variations between systems the same way any system-independent file would (hooks, `#if defined`, etc.), and machines which are radically different don't need to use `infptrace.c` at all.

All debugging code must be controllable using the `'set debug module'` command. Do not use `printf` to print trace messages. Use `fprintf_unfiltered(gdb_stdlog,` Do not use `#ifdef DEBUG`.

14 Porting GDB

Most of the work in making GDB compile on a new machine is in specifying the configuration of the machine. This is done in a dizzying variety of header files and configuration scripts, which we hope to make more sensible soon. Let's say your new host is called an `xyz` (e.g., `'sun4'`), and its full three-part configuration name is `arch-xvend-xos` (e.g., `'sparc-sun-sunos4'`). In particular:

- In the top level directory, edit `'config.sub'` and add `arch`, `xvend`, and `xos` to the lists of supported architectures, vendors, and operating systems near the bottom of the file. Also, add `xyz` as an alias that maps to `arch-xvend-xos`. You can test your changes by running

```
./config.sub xyz
```

and

```
./config.sub arch-xvend-xos
```

which should both respond with `arch-xvend-xos` and no error messages.

You need to port BFD, if that hasn't been done already. Porting BFD is beyond the scope of this manual.

- To configure GDB itself, edit `'gdb/configure.host'` to recognize your system and set `gdb_host` to `xyz`, and (unless your desired target is already available) also edit `'gdb/configure.tgt'`, setting `gdb_target` to something appropriate (for instance, `xyz`).
- Finally, you'll need to specify and define GDB's host-, native-, and target-dependent `'h'` and `'c'` files used for your configuration.

14.1 Configuring GDB for Release

From the top level directory (containing `'gdb'`, `'bfd'`, `'libiberty'`, and so on):

```
make -f Makefile.in gdb.tar.gz
```

This will properly configure, clean, rebuild any files that are distributed pre-built (e.g. `'c-exp.tab.c'` or `'refcard.ps'`), and will then make a tarfile. (If the top level directory has already been configured, you can just do `make gdb.tar.gz` instead.)

This procedure requires:

- symbolic links;

- `makeinfo` (texinfo2 level);
- `TEX`;
- `dvips`;
- `yacc` or `bison`.

... and the usual slew of utilities (`sed`, `tar`, etc.).

TEMPORARY RELEASE PROCEDURE FOR DOCUMENTATION

‘`gdb.texinfo`’ is currently marked up using the texinfo-2 macros, which are not yet a default for anything (but we have to start using them sometime).

For making paper, the only thing this implies is the right generation of ‘`texinfo.tex`’ needs to be included in the distribution.

For making info files, however, rather than duplicating the texinfo2 distribution, generate ‘`gdb-all.texinfo`’ locally, and include the files ‘`gdb.info*`’ in the distribution. Note the plural; `makeinfo` will split the document into one overall file and five or so included files.

15 Testsuite

The testsuite is an important component of the GDB package. While it is always worthwhile to encourage user testing, in practice this is rarely sufficient; users typically use only a small subset of the available commands, and it has proven all too common for a change to cause a significant regression that went unnoticed for some time.

The GDB testsuite uses the DejaGNU testing framework. DejaGNU is built using `Tcl` and `expect`. The tests themselves are calls to various `Tcl` procs; the framework runs all the procs and summarizes the passes and fails.

15.1 Using the Testsuite

To run the testsuite, simply go to the GDB object directory (or to the testsuite’s `objdir`) and type `make check`. This just sets up some environment variables and invokes DejaGNU’s `runtest` script. While the testsuite is running, you’ll get mentions of which test file is in use, and a mention of any unexpected passes or fails. When the testsuite is finished, you’ll get a summary that looks like this:

```
=== gdb Summary ===
```

```
# of expected passes      6016
# of unexpected failures   58
# of unexpected successes  5
# of expected failures     183
# of unresolved testcases  3
# of untested testcases   5
```

The ideal test run consists of expected passes only; however, reality conspires to keep us from this ideal. Unexpected failures indicate real problems, whether in GDB or in the testsuite. Expected failures are still failures, but ones which have been decided are too hard

to deal with at the time; for instance, a test case might work everywhere except on AIX, and there is no prospect of the AIX case being fixed in the near future. Expected failures should not be added lightly, since you may be masking serious bugs in GDB. Unexpected successes are expected fails that are passing for some reason, while unresolved and untested cases often indicate some minor catastrophe, such as the compiler being unable to deal with a test program.

When making any significant change to GDB, you should run the testsuite before and after the change, to confirm that there are no regressions. Note that truly complete testing would require that you run the testsuite with all supported configurations and a variety of compilers; however this is more than really necessary. In many cases testing with a single configuration is sufficient. Other useful options are to test one big-endian (Sparc) and one little-endian (x86) host, a cross config with a builtin simulator (powerpc-eabi, mips-elf), or a 64-bit host (Alpha).

If you add new functionality to GDB, please consider adding tests for it as well; this way future GDB hackers can detect and fix their changes that break the functionality you added. Similarly, if you fix a bug that was not previously reported as a test failure, please add a test case for it. Some cases are extremely difficult to test, such as code that handles host OS failures or bugs in particular versions of compilers, and it's OK not to try to write tests for all of those.

15.2 Testsuite Organization

The testsuite is entirely contained in `'gdb/testsuite'`. While the testsuite includes some makefiles and configury, these are very minimal, and used for little besides cleaning up, since the tests themselves handle the compilation of the programs that GDB will run. The file `'testsuite/lib/gdb.exp'` contains common utility procs useful for all GDB tests, while the directory `'testsuite/config'` contains configuration-specific files, typically used for special-purpose definitions of procs like `gdb_load` and `gdb_start`.

The tests themselves are to be found in `'testsuite/gdb.*'` and subdirectories of those. The names of the test files must always end with `'.exp'`. DejaGNU collects the test files by wildcarding in the test directories, so both subdirectories and individual files get chosen and run in alphabetical order.

The following table lists the main types of subdirectories and what they are for. Since DejaGNU finds test files no matter where they are located, and since each test file sets up its own compilation and execution environment, this organization is simply for convenience and intelligibility.

`'gdb.base'`

This is the base testsuite. The tests in it should apply to all configurations of GDB (but generic native-only tests may live here). The test programs should be in the subset of C that is valid K&R, ANSI/ISO, and C++ (`#ifdefs` are allowed if necessary, for instance for prototypes).

`'gdb.lang'` Language-specific tests for any language *lang* besides C. Examples are `'gdb.c++'` and `'gdb.java'`.

‘gdb.platform’

Non-portable tests. The tests are specific to a specific configuration (host or target), such as HP-UX or eCos. Example is ‘gdb.hp’, for HP-UX.

‘gdb.compiler’

Tests specific to a particular compiler. As of this writing (June 1999), there aren’t currently any groups of tests in this category that couldn’t just as sensibly be made platform-specific, but one could imagine a ‘gdb.gcc’, for tests of GDB’s handling of GCC extensions.

‘gdb.subsystem’

Tests that exercise a specific GDB subsystem in more depth. For instance, ‘gdb.disasm’ exercises various disassemblers, while ‘gdb.stabs’ tests pathways through the stabs symbol reader.

15.3 Writing Tests

In many areas, the GDB tests are already quite comprehensive; you should be able to copy existing tests to handle new cases.

You should try to use `gdb_test` whenever possible, since it includes cases to handle all the unexpected errors that might happen. However, it doesn’t cost anything to add new test procedures; for instance, ‘gdb.base/exprs.exp’ defines a `test_expr` that calls `gdb_test` multiple times.

Only use `send_gdb` and `gdb_expect` when absolutely necessary, such as when GDB has several valid responses to a command.

The source language programs do *not* need to be in a consistent style. Since GDB is used to debug programs written in many different styles, it’s worth having a mix of styles in the testsuite; for instance, some GDB bugs involving the display of source lines would never manifest themselves if the programs used GNU coding style uniformly.

16 Hints

Check the ‘README’ file, it often has useful information that does not appear anywhere else in the directory.

16.1 Getting Started

GDB is a large and complicated program, and if you first starting to work on it, it can be hard to know where to start. Fortunately, if you know how to go about it, there are ways to figure out what is going on.

This manual, the GDB Internals manual, has information which applies generally to many parts of GDB.

Information about particular functions or data structures are located in comments with those functions or data structures. If you run across a function or a global variable which does not have a comment correctly explaining what it does, this can be thought of as a bug in GDB; feel free to submit a bug report, with a suggested comment if you can figure out

what the comment should say. If you find a comment which is actually wrong, be especially sure to report that.

Comments explaining the function of macros defined in host, target, or native dependent files can be in several places. Sometimes they are repeated every place the macro is defined. Sometimes they are where the macro is used. Sometimes there is a header file which supplies a default definition of the macro, and the comment is there. This manual also documents all the available macros.

Start with the header files. Once you have some idea of how GDB's internal symbol tables are stored (see `'syntab.h'`, `'gdbtypes.h'`), you will find it much easier to understand the code which uses and creates those symbol tables.

You may wish to process the information you are getting somehow, to enhance your understanding of it. Summarize it, translate it to another language, add some (perhaps trivial or non-useful) feature to GDB, use the code to predict what a test case would do and write the test case and verify your prediction, etc. If you are reading code and your eyes are starting to glaze over, this is a sign you need to use a more active approach.

Once you have a part of GDB to start with, you can find more specifically the part you are looking for by stepping through each function with the `next` command. Do not use `step` or you will quickly get distracted; when the function you are stepping through calls another function try only to get a big-picture understanding (perhaps using the comment at the beginning of the function being called) of what it does. This way you can identify which of the functions being called by the function you are stepping through is the one which you are interested in. You may need to examine the data structures generated at each stage, with reference to the comments in the header files explaining what the data structures are supposed to look like.

Of course, this same technique can be used if you are just reading the code, rather than actually stepping through it. The same general principle applies—when the code you are looking at calls something else, just try to understand generally what the code being called does, rather than worrying about all its details.

A good place to start when tracking down some particular area is with a command which invokes that feature. Suppose you want to know how single-stepping works. As a GDB user, you know that the `step` command invokes single-stepping. The command is invoked via command tables (see `'command.h'`); by convention the function which actually performs the command is formed by taking the name of the command and adding `'_command'`, or in the case of an `info` subcommand, `'_info'`. For example, the `step` command invokes the `step_command` function and the `info display` command invokes `display_info`. When this convention is not followed, you might have to use `grep` or `M-x tags-search` in emacs, or run GDB on itself and set a breakpoint in `execute_command`.

If all of the above fail, it may be appropriate to ask for information on `bug-gdb`. But *never* post a generic question like “I was wondering if anyone could give me some tips about understanding GDB”—if we had some magic secret we would put it in this manual. Suggestions for improving the manual are always welcome, of course.

16.2 Debugging GDB with itself

If GDB is limping on your machine, this is the preferred way to get it fully functional. Be warned that in some ancient Unix systems, like Ultrix 4.2, a program can't be running

in one process while it is being debugged in another. Rather than typing the command `./gdb ./gdb`, which works on Suns and such, you can copy 'gdb' to 'gdb2' and then type `./gdb ./gdb2`.

When you run GDB in the GDB source directory, it will read a '`.gdbinit`' file that sets up some simple things to make debugging gdb easier. The `info` command, when executed without a subcommand in a GDB being debugged by gdb, will pop you back up to the top level gdb. See '`.gdbinit`' for details.

If you use emacs, you will probably want to do a `make TAGS` after you configure your distribution; this will put the machine dependent routines for your local machine where they will be accessed first by `M-`.

Also, make sure that you've either compiled GDB with your local cc, or have run `fixincludes` if you are compiling with gcc.

16.3 Submitting Patches

Thanks for thinking of offering your changes back to the community of GDB users. In general we like to get well designed enhancements. Thanks also for checking in advance about the best way to transfer the changes.

The GDB maintainers will only install "cleanly designed" patches. This manual summarizes what we believe to be clean design for GDB.

If the maintainers don't have time to put the patch in when it arrives, or if there is any question about a patch, it goes into a large queue with everyone else's patches and bug reports.

The legal issue is that to incorporate substantial changes requires a copyright assignment from you and/or your employer, granting ownership of the changes to the Free Software Foundation. You can get the standard documents for doing this by sending mail to gnu@gnu.org and asking for it. We recommend that people write in "All programs owned by the Free Software Foundation" as "NAME OF PROGRAM", so that changes in many programs (not just GDB, but GAS, Emacs, GCC, etc) can be contributed with only one piece of legalese pushed through the bureaucracy and filed with the FSF. We can't start merging changes until this paperwork is received by the FSF (their rules, which we follow since we maintain it for them).

Technically, the easiest way to receive changes is to receive each feature as a small context diff or unidiff, suitable for `patch`. Each message sent to me should include the changes to C code and header files for a single feature, plus '`ChangeLog`' entries for each directory where files were modified, and diffs for any changes needed to the manuals ('`gdb/doc/gdb.texinfo`' or '`gdb/doc/gdbint.texinfo`'). If there are a lot of changes for a single feature, they can be split down into multiple messages.

In this way, if we read and like the feature, we can add it to the sources with a single patch command, do some testing, and check it in. If you leave out the '`ChangeLog`', we have to write one. If you leave out the doc, we have to puzzle out what needs documenting. Etc., etc.

The reason to send each change in a separate message is that we will not install some of the changes. They'll be returned to you with questions or comments. If we're doing our job correctly, the message back to you will say what you have to fix in order to make the

change acceptable. The reason to have separate messages for separate features is so that the acceptable changes can be installed while one or more changes are being reworked. If multiple features are sent in a single message, we tend to not put in the effort to sort out the acceptable changes from the unacceptable, so none of the features get installed until all are acceptable.

If this sounds painful or authoritarian, well, it is. But we get a lot of bug reports and a lot of patches, and many of them don't get installed because we don't have the time to finish the job that the bug reporter or the contributor could have done. Patches that arrive complete, working, and well designed, tend to get installed on the day they arrive. The others go into a queue and get installed as time permits, which, since the maintainers have many demands to meet, may not be for quite some time.

Please send patches directly to [the GDB maintainers](#).

16.4 Obsolete Conditionals

Fragments of old code in GDB sometimes reference or set the following configuration macros. They should not be used by new code, and old uses should be removed as those parts of the debugger are otherwise touched.

STACK_END_ADDR

This macro used to define where the end of the stack appeared, for use in interpreting core file formats that don't record this address in the core file itself. This information is now configured in BFD, and GDB gets the info portably from there. The values in GDB's configuration files should be moved into BFD configuration files (if needed there), and deleted from all of GDB's config files. Any 'foo-xdep.c' file that references STACK_END_ADDR is so old that it has never been converted to use BFD. Now that's old!

PYRAMID_CONTROL_FRAME_DEBUGGING

pyr-xdep.c

PYRAMID_CORE

pyr-xdep.c

PYRAMID_PTRACE

pyr-xdep.c

REG_STACK_SEGMENT

exec.c

Index

—	a.out format	24
_initialize_language	add_cmd	9
	add_com	9
	add_symtab_fns	21
	adding a new host	29
	adding a symbol-reading module	21
A		

adding a target	51	ADDITIONAL_OPTIONS	38
adding debugging info reader	26	ADDR_BITS_REMOVE	38
adding source language	27	address representation	34
ADDITIONAL_OPTION_CASES	38	address spaces, separate data and code	34
ADDITIONAL_OPTION_HANDLER	38	ADDRESS_TO_POINTER	36, 38
ADDITIONAL_OPTION_HELP	38	algorithms	2

ALIGN_STACK_ON_STARTUP	30
allocate_symtab	28
assumptions about targets	64
ATTACH_DETACH	55
ATTR_NORETURN	32

B

BEFORE_MAIN_LOOP_HOOK	38
BELIEVE_PCC_PROMOTION	38
BELIEVE_PCC_PROMOTION_TYPE	38
BFD library	58
BIG_BREAKPOINT	39

BIG_REMOTE_BREAKPOINT 39
BITS_BIG_ENDIAN 39
BPT_VECTOR 51
BREAKPOINT 4, 39
BREAKPOINT_FROM_PC 39
breakpoints 3

bug-gdb mailing list 71
byte order 31

C

C data types 63

call stack frame.....	3	CALL_DUMMY_WORDS.....	40
CALL_DUMMY.....	40	call_function_by_hand.....	48
CALL_DUMMY_LOCATION.....	40	CANNOT_FETCH_REGISTER.....	40
CALL_DUMMY_P.....	40	CANNOT_STEP_HW_WATCHPOINTS.....	6
CALL_DUMMY_STACK_ADJUST.....	40	CANNOT_STORE_REGISTER.....	40
CALL_DUMMY_STACK_ADJUST_P.....	40	CC_HAS_LONG_LONG.....	31

CHILD_PREPARE_TO_STORE	55	code pointers, word-addressed	34
cleanup	12, 13	coding standards	60
cleanups	59	COERCE_FLOAT_TO_DOUBLE	40
CLEAR_DEFERRED_STORES	40	COFF debugging info	26
CLEAR_SOLIB	58	COFF format	24
CLI	9	command implementation	70

command interpreter..... 9
comment formatting..... 62
compiler warnings..... 61
converting between pointers and addresses..... 34
CPLUS_MARKERz..... 41
create_new_frame..... 3

CRLF_SOURCE_FILES..... 30
current_language..... 28

D

D10V addresses..... 34

data output	13	DEFAULT_PROMPT	30
DBX_PARM_SYMBOL_CLASS	41	deprecate_cmd	9
DEBUG_PTRACE	58	deprecating commands	9
debugging GDB	71	design	64
DECR_PC_AFTER_BREAK	41	DEV_TTY	30
DECR_PC_AFTER_HW_BREAK	6, 41	DIRNAME_SEPARATOR	65

DISABLE_UNSETTABLE_BREAK	41	DWARF 1 debugging info	26
discard_cleanups	59	DWARF 2 debugging info	26
do_cleanups	59	DWARF_REG_TO_REGNUM	41
DO_DEFERRED_STORES	40	DWARF2_REG_TO_REGNUM	41
DO_REGISTERS_INFO	41		
DOS text files	30		

E

ECOFF debugging info.....	26	<code>evaluate_subexp</code>	27
ECOFF format	24	expression evaluation routines	27
ECOFF_REG_TO_REGNUM.....	41	expression parser	27
ELF format	25	<code>extract_address</code>	35
END_OF_TEXT_DEFAULT.....	41	<code>EXTRACT_RETURN_VALUE</code>	42
		<code>EXTRACT_STRUCT_VALUE_ADDRESS</code>	42

EXTRACT_STRUCT_VALUE_ADDRESS_P	42	fetch_core_registers	54
extract_typed_address	34	FETCH_INFERIOR_REGISTERS	55
F		field output functions	13
FCLOSE_PROVIDED	31	file names, portability	65
		FILENAME_CMP	65
		FILES_INFO_HOOK	56

<code>find_pc_function</code>	23	<code>FP_REGNUM</code>	42
<code>find_pc_line</code>	23	<code>FPO_REGNUM</code>	56
<code>find_sym_fns</code>	21	<code>frame</code>	3
<code>finding a symbol</code>	23	<code>frame chain</code>	3
<code>FLOAT_INFO</code>	42	<code>frame pointer register</code>	3
<code>FOPEN_RB</code>	31	<code>FRAME_ARGS_ADDRESS_CORRECT</code>	42

FRAME_CHAIN	42	FRAME_SAVED_PC	43
FRAME_CHAIN_COMBINE	42	FRAMELESS_FUNCTION_INVOCATION	42
FRAME_CHAIN_VALID	42	full symbol table	22
FRAME_FP	3	function prototypes	63
FRAME_INIT_SAVED_REGS	43	function usage	63
FRAME_NUM_ARGS	43	FUNCTION_EPILOGUE_SIZE	43

FUNCTION_START_OFFSET	43	GCC2_COMPILED_FLAG_SYMBOL	43
fundamental types	23	GDB_MULTI_ARCH	43
G		GDB_TARGET_IS_HPPA	43
GCC_COMPILED_FLAG_SYMBOL	43	GDBINIT_FILENAME	30
		generic dummy frames	32
		generic host support	29

GET_LONGJMP_TARGET 4, 43, 56
get_saved_register 44
GET_SAVED_REGISTER 44
GETENV_PROVIDED 31

H

hardware breakpoints 3
hardware watchpoints 4
HAVE_CONTINUABLE_WATCHPOINT 6
HAVE_DOS_BASED_FILE_SYSTEM 65
HAVE_LONG_DOUBLE 31

HAVE_MMAP	31	host	2
HAVE_NONSTEPPABLE_WATCHPOINT	6	host, adding	29
HAVE_REGISTER_WINDOWS	44	HOST_BYTE_ORDER	31
HAVE_SIGSETMASK	31		
HAVE_STEPPABLE_WATCHPOINT	6		
HAVE_TERMIO	31		

I

i386_cleanup_dregs	9	i386_insert_watchpoint	8
I386_DR_LOW_GET_STATUS	7	i386_region_ok_for_watchpoint	8
I386_DR_LOW_RESET_ADDR	7	i386_remove_hw_breakpoint	8
I386_DR_LOW_SET_ADDR	7	i386_remove_watchpoint	8
I386_DR_LOW_SET_CONTROL	7	i386_stopped_by_hwbp	8
i386_insert_hw_breakpoint	8	i386_stopped_data_address	8

I386_USE_GENERIC_WATCHPOINTS	7	IN_SOLIB_RETURN_TRAMPOLINE	45
IBM6000_TARGET	44	INIT_EXTRA_FRAME_INFO	44
IEEE_FLOAT	44	INIT_FRAME_PC	44
IN_SIGTRAMP	44	INNER_THAN	44
IN_SOLIB_CALL_TRAMPOLINE	45	insert or remove hardware breakpoint	6
IN_SOLIB_DYNSYM_RESOLVE_CODE	45	INT_MAX	31

INT_MIN	31
IS_ABSOLUTE_PATH	65
IS_DIR_SEPARATOR	65
IS_TRAPPED_INTERNALVAR	45
ISATTY	31
item output functions	13

K

KERNEL_U_ADDR	56
KERNEL_U_ADDR_BSD	56
KERNEL_U_ADDR_HPUX	56

L

<code>L_SET</code>	32
language parser	27
language support	27
legal papers for code contributions	71
<code>length_of_subexp</code>	27
<code>libgdb</code>	19
line wrap in output	60
<code>lint</code>	33
list output functions	11
<code>LITTLE_BREAKPOINT</code>	39
<code>LITTLE_REMOTE_BREAKPOINT</code>	39

long long data type	31
LONG_MAX	31
LONGEST	31
longjmp debugging	4
lookup_symbol	23
LSEEK_NOT_LINEAR	32

M

make_cleanup	59
making a distribution tarball	67
MALLOC_INCOMPATIBLE	32
MEM_FNS_DECLARED	30
MEMORY_INSERT_BREAKPOINT	39

MEMORY_REMOVE_BREAKPOINT	39	MMAP_INCREMENT	32
minimal symbol table	22	mmcheck	32
minsyntaxs	22	MMCHECK_FORCE	33
mmap	32		
mmap	31		
MMAP_BASE_ADDRESS	32		

N

NATDEPFILES	53	Netware Loadable Module format	25
native conditionals	55	NNPC_REGNUM	47
native core files	54	NO_HIF_SUPPORT	45
native debugging	53	NO_MM_CHECK	32
NEED_TEXT_START_END	45	NO_SIG_INTERRUPT	33
nesting level in <code>ui_out</code> functions	11	NO_STD_REGS	30

NO_SYS_FILE	30
NORETURN	32
NPC_REGNUM	47

O

object file formats	24
obsolete code	72
ONE_PROCESS_WRITETEXT	56
op_print_tab	28
opcodes library	58
OS9K_VARIABLES_INSIDE_BLOCK	51

P

PARM_BOUNDARY	47	PC_REGNUM	47
parse_exp_1	28	PCC_SOL_BROKEN	47
partial symbol table	22	PE-COFF format	25
PC_IN_CALL_DUMMY	47	pointer representation	34
PC_LOAD_SEGMENT	47	POINTER_TO_ADDRESS	36, 45
		POP_FRAME	48

portability	65	PRINT_REGISTER_HOOK	47
portable file name handling	65	print_subexp	28
porting to new machines	66	PRINT_TYPELESS_INTEGER	47
prefixify_subexp	27	PRINTF_HAS_LONG_DOUBLE	31
PREPARE_TO_PROCEED	56	PRINTF_HAS_LONG_LONG	31
preparing a release	67	PROC_NAME_FMT	57

PROCESS_LINENUMBER_HOOK	47	psymtabs	22
program counter	3	PTRACE_ARG3_TYPE	57
PROLOGUE_FIRSTLINE_OVERLAP	48	PTRACE_FP_BUG	57
promotion to double	40	PUSH_ARGUMENTS	48
prompt	30	PUSH_DUMMY_FRAME	48
PS_REGNUM	48		

R

R_OK	33	reading of symbols	21
raw representation	36	REG_STRUCT_HAS_ADDR	48
read_fp	50	register data formats, converting	36
read_pc	50	REGISTER_BYTES	48
read_sp	50	REGISTER_CONVERT_TO_RAW	37, 46
		REGISTER_CONVERT_TO_VIRTUAL	37, 46

REGISTER_CONVERTIBLE	37, 45	REGISTER_VIRTUAL_SIZE	37, 45
REGISTER_IN_WINDOW_P	44	REGISTER_VIRTUAL_TYPE	46
REGISTER_NAME	48	regular expressions library	59
REGISTER_NAMES	48	remote debugging support	29
REGISTER_RAW_SIZE	37, 45	REMOTE_BPT_VECTOR	51
REGISTER_U_ADDR	57	REMOTE_BREAKPOINT	39

representations, raw and virtual	36
requirements for GDB	1
<code>return_command</code>	48
<code>RETURN_VALUE_ON_STACK</code>	46
returning structures by value	46
running the test suite	68

S

<code>SAVE_DUMMY_FRAME_TOS</code>	48
<code>SCANF_HAS_LONG_DOUBLE</code>	32
<code>SDB_REG_TO_REGNUM</code>	48
secondary symbol file	21
<code>SEEK_CUR</code>	33

SEEK_SET	33	SHIFT_INST_REGS	48
separate data and code address spaces	34	siginterrupt	33
serial line support	29	sigsetmask	31
set_gdbarch_coerce_float_to_double	41	SIGTRAMP_END	45
SHELL_COMMAND_CONCAT	57	SIGTRAMP_START	45
SHELL_FILE	57	SIGWINCH_HANDLER	30

SIGWINCH_HANDLER_BODY	30	SKIP_TRAMPOLINE_CODE	49
SIZEOF_CALL_DUMMY_WORDS	40	SLASH_STRING	65
SKIP_PERMANENT_BREAKPOINT	48	software breakpoints	3
SKIP_PROLOGUE	49	software watchpoints	4
SKIP_PROLOGUE_FRAMELESS_P	49	SOFTWARE_SINGLE_STEP	46
SKIP_SOLIB_RESOLVER	45	SOFTWARE_SINGLE_STEP_P	46

SOFUN_ADDRESS_MAYBE_MISSING	46	SP_REGNUM	49
SOLIB_ADD	57	spaces, separate data and code address	34
SOLIB_CREATE_INFERIOR_HOOK	57	STAB_REG_TO_REGNUM	49
SOM debugging info	26	stabs debugging info	25
SOM format	25	stack alignment	30
source code formatting	62	STACK_ALIGN	49

<code>standard_coerce_float_to_double</code>	41	<code>STORE_RETURN_VALUE</code>	49
<code>START_INFERIOR_TRAPS_EXPECTED</code>	57	<code>store_typed_address</code>	35
<code>STEP_SKIPS_DELAY</code>	49	<code>struct value</code> , converting register contents to ..	36
<code>STOP_SIGNAL</code>	33	structures, returning by value.....	46
<code>STOPPED_BY_WATCHPOINT</code>	6	submitting patches.....	71
<code>store_address</code>	35	<code>SUN_FIXED_LBRAC_BUG</code>	49

SVR4_SHARED_LIBS	57	SYMBOLS_CAN_START_WITH_DOLLAR	44
sym_fns structure	21	symtabs	22
symbol files	21	system dependencies	65
symbol lookup	23		
symbol reading	21		
SYMBOL_RELOADING_DEFAULT	49		

T

table output functions.....	11	TARGET_CAN_USE_HARDWARE_WATCHPOINT.....	5
target.....	2	TARGET_CHAR_BIT.....	50
target architecture definition.....	33	TARGET_COMPLEX_BIT.....	50
target vector.....	52	TARGET_DISABLE_HW_WATCHPOINTS.....	5
TARGET_BYTE_ORDER_DEFAULT.....	49	TARGET_DOUBLE_BIT.....	50
TARGET_BYTE_ORDER_SELECTABLE_P.....	49	TARGET_DOUBLE_COMPLEX_BIT.....	50

TARGET_ENABLE_HW_WATCHPOINTS	5	TARGET_LONG_BIT	50
TARGET_FLOAT_BIT	50	TARGET_LONG_DOUBLE_BIT	50
TARGET_HAS_HARDWARE_WATCHPOINTS	5	TARGET_LONG_LONG_BIT	50
target_insert_hw_breakpoint	6	TARGET_PTR_BIT	50
target_insert_watchpoint	6	TARGET_RANGE_PROFITABLE_FOR_HW_WATCHPOINT	
TARGET_INT_BIT	50	5

TARGET_READ_FP	50	target_remove_watchpoint	6
TARGET_READ_PC	50	TARGET_SHORT_BIT	50
TARGET_READ_SP	50	target_stopped_data_address	6
TARGET_REGION_OK_FOR_HW_WATCHPOINT	5	TARGET_VIRTUAL_FRAME_POINTER	50
TARGET_REGION_SIZE_OK_FOR_HW_WATCHPOINT ...	5	TARGET_WRITE_FP	50
target_remove_hw_breakpoint	6	TARGET_WRITE_PC	50

TARGET_WRITE_SP	50	trimming language-dependent code	28
TCP remote support	30	tuple output functions	11
TDEPFILES	51	type	37
terminal device	30	type codes	24
test suite	67	types	63
test suite organization	68		

U

U_REGS_OFFSET	57	ui_out_field_int	13
ui_out functions	10	ui_out_field_skip	14
ui_out functions, usage examples	15	ui_out_field_stream	14
ui_out_field_core_addr	13	ui_out_field_string	13
ui_out_field_fmt	13	ui_out_flush	15
		ui_out_list_begin	13

<code>ui_out_list_end</code>	13	<code>ui_out_table_end</code>	12
<code>ui_out_message</code>	15	<code>ui_out_table_header</code>	12
<code>ui_out_spaces</code>	15	<code>ui_out_text</code>	15
<code>ui_out_stream_delete</code>	14	<code>ui_out_tuple_begin</code>	12
<code>ui_out_table_begin</code>	11	<code>ui_out_tuple_end</code>	12
<code>ui_out_table_body</code>	12	<code>ui_out_wrap_hint</code>	15

<code>ui_stream</code>	14	<code>USE_PROC_FS</code>	57
<code>UINT_MAX</code>	31	<code>USE_STRUCT_CONVENTION</code>	51
<code>ULONG_MAX</code>	31	<code>USG</code>	33
<code>USE_GENERIC_DUMMY_FRAMES</code>	32	using <code>ui_out</code> functions	15
<code>USE_MMALLOC</code>	32		
<code>USE_O_NOCTTY</code>	33		

V

<code>value_as_pointer</code>	35
<code>value_from_pointer</code>	35
<code>VARIABLES_INSIDE_BLOCK</code>	51
virtual representation	36
<code>volatile</code>	33

W

watchpoints	4
watchpoints, on x86	7
word-addressed machines	34
<code>wrap_here</code>	60
<code>write_fp</code>	50

`write_pc`..... 50
`write_sp`..... 50
writing tests..... 69

X

x86 debug registers 7
XCOFF format..... 25
`XDEPFILES` 29