

Redis as a Time Series DB

Josiah Carlson - @dr_josiah

Agenda

- Who are you?
- What is Redis? (3 minutes, optional)
- What is a time series database?
- Combining structures for success
- Analyzing/segmenting events
- Next steps
- Questions

Who are you?

- A guy who does a lot of stuff with Redis and Python
 - Book:
 - Redis in Action <http://manning.com/carlson>
 - Read for free <http://bit.ly/ria-free>
 - Libraries: <https://github.com/josiahcarlson>
 - rom, RPQueue, lua_call, parse-crontab, ...
 - Mailing list: <https://groups.google.com/forum/#!forum/redis-db> (not active here recently)
 - Blog: dr-josiah.com
 - Twitter: @dr_josiah

What is Redis?

- In-memory key/data structure server
 - Strings, lists, hashes, sets, sorted sets (plus integers and floats)
 - HyperLogLog, lexicographic prefix scans, geo index/search api (unreleased 3.2)
 - Pubsub, Lua scripting (like a stored procedure)
- Ops features:
 - Master/slave replication + sentinel for HA
 - Alternative “cluster” mode with sharding/slaving/failover
 - On-disk persistence; point-in-time or incremental

What is a time series database?

- Not about the database, but about the data and operations
 - Discrete events or samples of at least one value or metric over time
 - Sometimes the *event* is the value/metric
 - If it has a timestamp (or one is implied), it is part of a time series
- Examples of data:
 - Data gathered from sensors embedded inside IoT devices
 - Nest thermostat - measures temperature, analyzes your preferences...
 - Sell prices and volumes of traded stocks
 - Values and delivery locations of orders placed at an online retailer
 - Actions of users in a video game and their outcome

Combining structures for success

- Consider a log of user actions on a web site like:
 - `{'ts': 1458710360.679, 'user': 218946, 'type': 'login', 'ip': '172.32.5.6'}`
- We can get an easily sliced time series with individual events stored in hashes, and a sorted set for an index of timestamps:
 - Increment the **actions:** key, which becomes the event **id**
 - Store each event individually in its own HASH at key: **actions:<event_id>**
 - Add an (**<event_id>: <ts>**) member/score pair to a ZSET at key: **actions:ts**
- To use:
 - Scan over **actions:ts** using ZRANGEBYSCORE or ZRANGE
 - Use HGET/HMGET/HGETALL to fetch the actions for analysis

Combining structures for success (reading)

```
-- Many-caveats about using this, you probably don't want it (not cluster compatible)!
-- KEYS should be something like:
--     {'action:ts', 'action:'}
-- ARGV should be something like:
-- (gets all events from 2016-03-23 UTC)
--     {'1458691200', '1458777600'}
```

```
local ids = redis.call('ZRANGEBYSCORE', KEYS[1], unpack(ARGV))
```

```
local results = {}
for i, id in ipairs(ids) do
    local item = redis.call('HGETALL', KEYS[2]..id)
    if item then
        table.insert(results, item)
    end
end

return cjson.encode(results)
```

Combining structures for success (writing)

```
-- Many-caveats about using this, you probably don't want it (not cluster compatible)!
-- KEYS should be something like:
--     {'action:'}
-- ARGV should be something like:
--     {'<json-encoded data, with ts>'}
```

```
local new_id = redis.call('INCR', KEYS[1])
local dest = KEYS[1] .. new_id
```

```
local data = cjson.decode(ARGV[1])
data.id = new_id
```

```
for k, v in pairs(data) do
    redis.call('HSET', dest, k, v)
end
```

```
-- note:      ZADD          <key>          <score> <member>
redis.call('ZADD', KEYS[1] .. 'ts', data.ts, new_id)
```


Analyzing/segmenting events

We now have events, but want to (for example) count the number of events of different types in a time range, equivalent to:

```
SELECT type, COUNT(*)  
FROM actions  
WHERE ...  
GROUP BY type
```

- We can use Lua to scan over the items in the time range, counting them
- We can create additional ZSETs, one for each event type...
 - **actions:login:ts, actions:logout:ts, ...**
 - Stores (**<event_id>: <ts>**) pairs, like **actions:ts**
 - Can use ZCOUNT to get a count without scanning*

Analyzing/segmenting events (continued)

- Or if you needed to index events by user id
 - Add (**<event_id>**, **<user_id>**) member/score pairs to **actions:user**
 - Easy to get total count/user with ZCOUNT (**ZCOUNT actions:user <user_id> <user_id>**)
 - Can partially induce a secondary timestamp index over the same data; to fetch a specific time range for one user
 - Use ZRANGEBYSCORE and ZREVRANGEBYSCORE on **actions:ts** for start/end of id range for a pair of timestamps (assuming synchronized clocks, or using Redis 3.2-unreleased Lua side-effects propagation)
 - Either:
 - Use ZRANGEBYSCORE on **actions:user** to scan ids and filter
 - Use ZRANGEBYSCORE + ZREVRANGEBYSCORE + ZRANGE to bisect (speed optimization, requires left padding of members with zeros)

Combining structures for success (revisited)

- If you know you never need to extract individual fields, and you know that your events are unique...
 - Use a single ZSET with (**<encoded json event>**: **<ts>**) pairs
- If you are planning on bucketing based on time slot, only need a list of items
 - Use LISTS keyed as **actions:<time slice>**
 - Just R PUSH **<encoded json event>** onto these LISTS
- Approximate counts? HyperLogLog
- If you plan your metrics, you can explicitly count them in real time
 - INCR, DECR, INCRBY, INCRBYFLOAT, HINCRBY, HINCRBYFLOAT, ZINCRBY, all very fast
 - Lua scripting can be a huge win here

Combining structures for success (revisited)

- Redis doesn't need to be 100% of this
 - If you only need counts over arbitrary time ranges, just **action:ts** or the **action:login:ts** and related sorted sets may be enough
 - Redis can be your counters/aggregate/sorted set storage, even if it's not your event storage
- HASHes + ZSETs can solve just about anything, alternatives **can** offer:
 - Better performance
 - More convenient access to necessary data
 - Lower memory utilization

Next steps

Time series is a specific example of data modeling...

- My data modeling talk from Redisconf 2015
 - Get an idea of the general process involved
- Specific examples:
 - For indexing/search: chapter 7 of RiA, my Redisconf 2012 talk
 - Other data modeling examples: the rest of RiA, my blog
 - Survey of data modeling examples: my Python with Redis talk from July 2012
- Google will know more

Questions?

Josiah Carlson - @dr_josiah