

LinuxCon Europe 2016

# Control Groups (cgroups)

© 2016 Michael Kerrisk

man7.org Training and Consulting

<http://man7.org/training/>

@mkerrisk      mtk@man7.org

4 October 2016

Berlin, Germany

# Outline

---

- 1 Introduction
- 2 Cgroups v1: hierarchies and controllers
- 3 Cgroups v1: populating a cgroup
- 4 Cgroups v1: a survey of the controllers
- 5 Cgroups /proc files
- 6 Cgroups v2: background and introduction
- 7 Cgroups v2: enabling and disabling controllers
- 8 Cgroups v2: organizing cgroups and processes

# Who am I?

---

- Maintainer of Linux man-pages (since 2004)
  - Documents kernel-user-space + C library APIs
    - ~1000 manual pages
    - <http://www.kernel.org/doc/man-pages/>
- API review, testing, and documentation
  - API design and design review
  - Lots of testing, lots of bug reports, a few kernel patches
- “Day job”: programmer, trainer, writer
  - <http://man7.org/>

# Outline

---

## 1 Introduction

2 Cgroups v1: hierarchies and controllers

3 Cgroups v1: populating a cgroup

4 Cgroups v1: a survey of the controllers

5 Cgroups /proc files

6 Cgroups v2: background and introduction

7 Cgroups v2: enabling and disabling controllers

8 Cgroups v2: organizing cgroups and processes

- Cgroups is a big topic
  - Many controllers
  - V1 versus V2 interfaces
- Our goal: understand fundamental semantics of cgroup filesystem and interfaces
  - (“Containers are too high level for me”)
  - Useful from a programming perspective
    - How do I build container frameworks?
    - What else can I build with cgroups?
  - And useful from a system engineering perspective
    - What's going on underneath my container's hood?

- We'll focus on:
  - General principles of operation; goals of cgroups
  - The cgroup filesystem
  - Interacting with the cgroup filesystem using shell commands
  - Problems with cgroups v1, motivations for cgroups v2
  - Differences between cgroups v1 and v2
- We'll look **briefly** at some of the controllers

- Kernel Documentation files
  - Documentation/cgroup-v1/\*.txt
  - Documentation/cgroup-v2.txt
- *cgroups(7)* man page
- Neil Brown's excellent (2014) LWN.net series on Cgroups:  
*<https://lwn.net/Articles/604609/>*
  - Thought-provoking commentary on the meaning of grouping and hierarchy
- *<https://lwn.net/Articles/484254/>* – Tejun Heo's initial thinking about redesigning cgroups
- Other articles at *[https://lwn.net/Kernel/Index/#Control\\_groups](https://lwn.net/Kernel/Index/#Control_groups)*

- 2006/2007, “Process Containers”
  - Developed by engineers at Google
  - 2007: renamed “control groups” to avoid confusion with alternate meaning for “containers”
- January 2008: initial release in mainline kernel (Linux 2.6.24)
- Fast-forward a few years...
  - Many new resource controllers added
- Various problems arose from haphazard/uncoordinated development of cgroup controllers
  - “Design followed implementation” :-(



- Sep 2012: work begins on cgroups v2
  - In-kernel changes, but marked experimental
  - Changes were necessarily incompatible with cgroups v1
    - $\Rightarrow$  Create new/orthogonal filesystem interface for v2
- March 2016, Linux 4.5: cgroups version 2 becomes official
  - Older version (cgroups v1) remains
    - A.k.a. “legacy cgroups”, but not going away in a hurry
- Cgroups v2 work is ongoing
  - For now, some functionality remains available only via cgroups v1
    - Subject to some rules, can use both versions at same time

- Two principle components:
  - A **mechanism for hierarchically grouping** processes
  - A set of **controllers** that manage, control, or monitor processes in cgroups
    - (Resources such as CPU, memory, block I/O bandwidth)
- Interface is via a pseudo-filesystem
- Cgroup manipulation takes form of filesystem operations
  - E.g., can use shell commands

# What do cgroups allow us to do?

---

- **Limit** resource usage of group
  - E.g., limit percentage of CPU available to group
- **Prioritize** group for resource allocation
  - E.g., some group might get greater proportion of CPU
- **Resource** accounting
  - Measure resources used by processes
- Freeze a group
  - Freeze, restore, and checkpoint a group
- And more...

- **Control group**: group of **processes bound** to set of parameters or limits
- **(Resource) controller**: kernel component that controls or monitors processes in a cgroup
  - E.g., memory controller limits memory usage; cpublacc accounts for CPU usage
  - Also known as **subsystem**
    - (But that term is rather ambiguous)
- Cgroups for each controller can be arranged in a **hierarchy**
  - Child cgroups may **inherit attributes** from parent

- Cgroup filesystem **directory structure defines cgroups + cgroup hierarchy**
  - I.e., use *mkdir(2)* / *rmdir(2)* (or equivalent shell commands) to create cgroups
- Each **subdirectory contains automagically created files**
  - Some files are used to **manage the cgroup** itself
  - Other files are **controller-specific**
- Files in cgroup are used for purposes such as:
  - **Defining/displaying membership** of cgroup
  - **Controlling behavior** of processes in cgroup
  - **Exposing information** about processes in cgroup (e.g., resource usage stats)

## Example: the pids controller (cgroups v1)

---

- pids (“process number”) controller allows us to limit number of PIDs in cgroup
  - Prevent *fork()* bombs!
- Use *mount* to attach pids controller to cgroup filesystem:

```
# mkdir -p /sys/fs/cgroup/pids # Create mount point
# mount -t cgroup -o pids none /sys/fs/cgroup/pids
```

- ⚠ May not be necessary
- Some systems automatically mount filesystems with controllers attached
  - E.g., *systemd* mounts the v1 controllers under subdirectories of */sys/fs/cgroup*, a *tmpfs* filesystem mounted via:

```
# mount -t tmpfs tmpfs /sys/fs/cgroup
```

## Example: the pids controller (cgroups v1)

---

- Create new cgroup, and place shell's PID in that cgroup:

```
# mkdir /sys/fs/cgroup/pids/g1
# echo $$
17273
# echo $$ > /sys/fs/cgroup/pids/g1/cgroup.procs
```

- cgroup.procs defines/displays PIDs in cgroup
- Which processes are in cgroup?

```
# cat /sys/fs/cgroup/pids/g1/cgroup.procs
17273
20591
```

- Where did PID 20591 come from?
    - PID 20591 is *cat* command, created as a child of shell
      - Child processes inherit parent's cgroup membership(s)

## Example: the pids controller (cgroups v1)

---

- Limit number of processes in cgroup, and show effect:

```
# echo 20 > /sys/fs/cgroup/pids/g1/pids.max
# for a in $(seq 1 20); do sleep 20 & done
[1] 20938
...
[18] 20955
bash: fork: retry: Resource temporarily unavailable
```

- `pids.max` defines/exposes limit on number of PIDs in cgroup



# Outline

---

- 1 Introduction
- 2 **Cgroups v1: hierarchies and controllers**
- 3 Cgroups v1: populating a cgroup
- 4 Cgroups v1: a survey of the controllers
- 5 Cgroups /proc files
- 6 Cgroups v2: background and introduction
- 7 Cgroups v2: enabling and disabling controllers
- 8 Cgroups v2: organizing cgroups and processes

- **Cgroup** == collection of processes
- **cgroup hierarchy** == hierarchical arrangement of cgroups
  - Implemented via a **cgroup pseudo-filesystem**
- Structure and membership of cgroup hierarchy is defined by:
  - ① **Mounting** a cgroup filesystem
  - ② **Creating a subdirectory structure** that reflects desired cgroup hierarchy
  - ③ **Moving processes within hierarchy** by writing their PIDs to special files in cgroup subdirectories

# Attaching a controller to a hierarchy

---

- A controller is attached to a hierarchy by mounting a cgroup filesystem:

```
# mkdir -p /sys/fs/cgroup/mem    # Create mount point
# mount -t cgroup -o memory none /sys/fs/cgroup/mem
```

- Here, memory controller was mounted
- none can be replaced by any suitable mnemonic name
  - Not interpreted by system, but appears in /proc/mounts

# Attaching a controller to a hierarchy

---

- To see which cgroup filesystems are mounted and their attached controllers:

```
# mount | grep cgroup
none on /sys/fs/cgroup/mem type cgroup (rw,memory)
# cat /proc/mounts | grep cgroup
none /sys/fs/cgroup/mem cgroup rw,relatime,memory 0 0
```

- Unmounting filesystem detaches the controller:

```
# umount /sys/fs/cgroup/mem
```

- But..., filesystem will remain (invisibly) mounted if it contains child cgroups
  - I.e., must move all processes to root cgroup, and remove child cgroups, to truly unmount

# Attaching controllers to hierarchies

---

- A controller can be **attached to only one hierarchy**
  - Mounting same controller at different mount point simply creates second view of same hierarchy
- **Multiple** controllers can be attached to same hierarchy:

```
# mkdir -p /sys/fs/cgroup/mem_cpu
# mount -t cgroup -o memory,cpu none \
        /sys/fs/cgroup/mem_cpu
```

- In effect, resources associated with those controllers are being managed together

# Creating cgroups

---

- When a new hierarchy is created, all **tasks** on system are part of **root cgroup** for that hierarchy
- New cgroups are **created** by creating subdirectories under cgroup mount point:

```
# mkdir /sys/fs/cgroup/mem/g1
```

- Relationships between cgroups are reflected by creating nested (arbitrarily deep) subdirectory structure
  - Meaning of hierarchical relationship depends on controller

# Destroying cgroups

---

- An **empty cgroup** can be **destroyed** by removing directory
  - **Empty** == last process in cgroup terminates or migrates to another cgroup **and** last child cgroup is removed
  - Not necessary (or possible) to delete attribute files inside cgroup directory before deleting it

# Outline

---

- 1 Introduction
- 2 Cgroups v1: hierarchies and controllers
- 3 Cgroups v1: populating a cgroup**
- 4 Cgroups v1: a survey of the controllers
- 5 Cgroups /proc files
- 6 Cgroups v2: background and introduction
- 7 Cgroups v2: enabling and disabling controllers
- 8 Cgroups v2: organizing cgroups and processes



# Files for managing cgroup membership

---

- To manage cgroup membership, each subdirectory in a hierarchy includes two automatically created files:
  - `cgroup.procs`
  - `tasks`

# Tasks?

---

- Cgroups v1 draws distinction between **process** and **task**
- **Task** == kernel scheduling entity
  - From scheduler's perspective, "processes" and "threads" are pretty much the same thing....
  - (Threads just share more state than processes)
- Multithreaded (MT) process == set of tasks with same **thread group ID** (TGID)
  - TGID == PID!
  - Each thread has unique **thread ID** (TID)
- Here, TID means **kernel thread ID**
  - I.e., value returned by *clone(2)* and *gettid(2)*
  - Not same as POSIX threads `pthread_t`
    - (But there is 1:1 relationship in NPTL implementation...)

# Placing a process in a cgroup

---

- To move a **process** to a cgroup, write its PID to cgroup.procs file in corresponding subdirectory

```
# echo $$ > /sys/fs/cgroup/mem/g1/cgroup.procs
```

# Viewing cgroup membership

---

- **To see PIDs in cgroup**, read `cgroup.procs` file
  - PIDs are newline-separated
- ⚠ List is **not guaranteed to be sorted or free of duplicates**
  - PID might be moved out and back into cgroup or recycled while reading list

# Placing a thread (task) in a cgroup

---

- Writing a PID to **cgroup.procs** **moves all threads in thread group** to a cgroup
- Each cgroup directory also has a tasks file...
  - Writing a TID to **tasks** **moves that thread** to cgroup
    - This feature goes away in cgroups v2...
  - Reading tasks shows all TIDs in cgroup

# Cgroup membership details

---

- Within a hierarchy, a **task can be member of just one cgroup**
  - That association defines attributes / parameters that apply to the task
- Adding a task to a different cgroup automatically removes it from previous cgroup
- A task can be a member of multiple cgroups, each of which is in a different hierarchy
- On *fork()*, **child inherits cgroup memberships** of parent
  - Afterward, cgroup memberships of parent and child can be independently changed

- Consider the following scenario:
  - We create a cgroup subdirectory
  - Some processes are moved into cgroup
  - Eventually, all of those processes terminate
- Who cleans up/gets notified when last process leaves cgroup?
  - We might want cgroup subdirectory to be removed
  - Manager process might want to know when all workers have terminated

- `release_agent` in cgroup **root** directory
  - Contains pathname of binary/script that is executed **when cgroup becomes empty**
    - E.g., this program might remove cgroup subdirectory
  - Release agent gets one **command-line argument**:  
pathname of cgroup subdirectory that has become empty
- `notify_on_release` in each cgroup subdirectory
  - Should `release_agent` be run when cgroup becomes empty? (0 == no, 1 == yes)
  - Initial setting for this file is inherited from cgroup parent



# Mounting a *named* hierarchy with no controller

---

- Can mount a *named* hierarchy with no attached controller:

```
# mount -t cgroup cgroup -o none,name=somename \  
    /some/mount/point
```

- Named hierarchies can be used to organize and track processes
  - E.g., PIDs can be moved into `cgroup.procs`, and will automatically disappear on process termination
    - (And we can use `release_agent`, etc.)
  - *systemd* creates such a hierarchy for its management of processes
    - Mounted at `/sys/fs/cgroup/systemd`
- Cgroups v1 only

# Exercises

---

- ① In this exercise, we create a cgroup, place a process in the cgroup, and then migrate that process to a different cgroup.

- If the memory cgroup is not already mounted, mount it:

```
# cat /proc/mounts | grep cgroup      # Is cgroup mounted?
# mkdir -p /sys/fs/cgroup/memory
# mount -t cgroup -o memory none /sys/fs/cgroup/memory
# cd /sys/fs/cgroup/memory
```

- Note: some systems (e.g., Debian) provide a patched kernel that disables the memory controller by default. If you find that you can't mount the memory controller, it may be necessary to reboot the kernel with the `cgroup_enable=memory` command-line option.
- Create two subdirectories, `m1` and `m2`, in the memory cgroup root directory.
- Execute the following command, and note the PID assigned to the resulting process:

```
# sleep 300 &
```

- Write the PID of the process created in the previous step into the file `m1/cgroup.procs`, and verify by reading the file contents.
- Now write the PID of the process into the file `m2/cgroup.procs`.
- Is the PID still visible in the file `m1/cgroup.procs`? Explain.

# Outline

---

- 1 Introduction
- 2 Cgroups v1: hierarchies and controllers
- 3 Cgroups v1: populating a cgroup
- 4 Cgroups v1: a survey of the controllers**
- 5 Cgroups /proc files
- 6 Cgroups v2: background and introduction
- 7 Cgroups v2: enabling and disabling controllers
- 8 Cgroups v2: organizing cgroups and processes

- Each of following controllers is selectable via a **kernel configuration option**
  - And there is an overall option, `CONFIG_CGROUPS`
- For each controller, there are controller-specific files in each cgroup directory
  - Names are prefixed with controller-specific string
    - E.g., `cpuacct.stat`, `pids.max`, `freezer.state`
- Individual **documentation** files for most controllers can be found in `Documentation/cgroup-v1`
  - $\Rightarrow$  Following slides give just a brief picture...

- `cpu` (2.6.24): control distribution of CPU cycles to cgroups
  - `cpu.cfs_period_us`: measurement period for CFS scheduler ( $\mu$ s; default: 100000)
  - `cpu.cfs_quota_us`: allowed run-time within period ( $\mu$ s; default: -1 [no limit])
  - Constraints propagate into child cgroups
- `cpuacct` (2.6.24): expose CPU usage of cgroup
  - `cpuacct.usage`: CPU usage by this cgroup (nanoseconds)
  - `cpuacct.stat`: user vs system mode CPU time (measured in USER\_HZ [centiseconds])
  - Statistics include CPU consumed in descendant cgroups

- memory (2.6.25): control memory usage of cgroups
  - Limit memory usage per cgroup
    - Soft limits influence page reclaim under memory pressure
    - Hard limits trigger per-cgroup OOM killer
  - Memory-usage accounting (optionally hierarchical)
  - Disable knob for OOM killer
  - Kernel-to-user-space notification for low-memory and OOM situations
    - E.g., instead of OOM killing, freeze processes, notify user space, remedy situation, thaw processes
  - And more; see `Documentation/cgroup-v1/memory.txt`
    - (but “this document is hopelessly outdated”)

- freezer (2.6.28): freeze (suspend) and resume processes in a cgroup
  - Gets round some limitations of using SIGSTOP/SIGCONT for this purpose
    - SIGSTOP is observable by waiting/ptracing parent
    - SIGCONT can be caught by application!
  - Cgroup is frozen / resumed by writing FROZEN / THAWED to `freezer.state`
    - Operations propagate to child cgroups
- blkio (2.6.33): limit I/O on block devices
  - HDDs, SSDs, USB, etc.
  - Policies:
    - Proportional-weight division of device bandwidth
    - Throttling/upper-limit

- `pids` (4.3): limit number of tasks in a cgroup
  - Prevent fork bombs
  - `pids.max`: maximum number of tasks in cgroup (and cgroup descendants)
    - Writing “max” into this file means no limit
    - Limit affects *fork()*/*clone()*
    - Doesn't affect attempts to move processes into cgroup
  - `pids.current`: number of PIDs currently in cgroup
  - ⚠ `pids.current` & `pids.max` **count tasks** not processes
  - Limit on a cgroup == most stringent limit on any ancestor cgroup (and descendants)



## Other cgroups v1 controllers

---

- `cpuset` (2.6.24): assign CPUs & memory nodes to cgroups
- `devices` (2.6.26): whitelist controller to permit/deny access to device by members of cgroup
- `perf_event` (2.6.39): carry out *perf* monitoring per cgroup
- `net_cls` (2.6.29), `net_prio` (3.3): traffic shaping and priority control of cgroup's network traffic
- `hugetlb` (3.6): limit hugeTLB usage per cgroup

- ① The freezer controller can be used to suspend and resume execution of all of the processes in a cgroup hierarchy. Create a cgroup hierarchy containing two child cgroups (thus three cgroups in total) as follows:

```
# mkdir /sys/fs/cgroup/freezer/mfz
# mkdir /sys/fs/cgroup/freezer/mfz/sub1
# mkdir /sys/fs/cgroup/freezer/mfz/sub2
```

Then run four separate instances of the `timers/cpu_burner.c` program, and place two of the resulting processes in the `mfz/sub1` cgroup, and one each of the remaining processes in `mfz` and `mfz/sub2`. Observe what happens to these processes as the following commands are executed.

Freeze the processes in the `mfz/sub1` cgroup:

```
# echo FROZEN > /sys/fs/cgroup/freezer/mfz/sub1/freezer.state
```

Freeze all of the processes in all cgroups under the `mfz` subtree:

```
# echo FROZEN > /sys/fs/cgroup/freezer/mfz/freezer.state
```

# Exercises

---

Thaw all of the processes in the `mfz` subtree, so that they resume execution:

```
# echo THAWED > /sys/fs/cgroup/freezer/mfz/freezer.state
```

Once more freeze the entire subtree, and then try thawing just the processes in the `mfz/sub1` cgroup:

```
# echo FROZEN > /sys/fs/cgroup/freezer/mfz/freezer.state  
# echo THAWED > /sys/fs/cgroup/freezer/mfz/sub1/freezer.state
```

Do the processes in the `mfz/sub1` cgroup resume execution? Why not? For a clue, view the status of the cgroup parent of this cgroup using the following command:

```
# cat /sys/fs/cgroup/freezer/mfz/sub1/freezer.parent_freezing
```

# Outline

---

- 1 Introduction
- 2 Cgroups v1: hierarchies and controllers
- 3 Cgroups v1: populating a cgroup
- 4 Cgroups v1: a survey of the controllers
- 5 Cgroups /proc files**
- 6 Cgroups v2: background and introduction
- 7 Cgroups v2: enabling and disabling controllers
- 8 Cgroups v2: organizing cgroups and processes

- /proc/cgroups describes controllers available on system

#subsys_name	hierarchy	num_cgroups	enabled
cpuset	4	1	1
cpu	8	1	1
cpuacct	8	1	1
blkio	6	1	1
memory	3	1	1
devices	10	84	1
freezer	7	1	1
net_cls	9	1	1
perf_event	5	1	1
net_prio	9	1	1
hugetlb	0	1	0
pids	2	1	1

- ① Controller name
- ② Unique hierarchy ID (0 for v2 hierarchy)
  - Multiple controllers may be bound to same hierarchy
- ③ Number of cgroups in hierarchy
- ④ Controller enabled? 1 == yes, 0 == no
  - Kernel `cgroup_disable` boot parameter

- /proc/PID/cgroup shows cgroup memberships of PID

```
3:cpu,cpuacct:/memgrp3
2:freezer:/
0::/grp1
```

- ① Hierarchy ID (0 for v2 cgroup)
  - Can be matched to hierarchy ID in /proc/cgroups
- ② Comma-separated list of controllers bound to the hierarchy
  - Field is empty for v2 cgroup
- ③ Pathname of cgroup to which this process belongs
  - Pathname is relative to cgroup root directory

# Outline

---

- 1 Introduction
- 2 Cgroups v1: hierarchies and controllers
- 3 Cgroups v1: populating a cgroup
- 4 Cgroups v1: a survey of the controllers
- 5 Cgroups /proc files
- 6 Cgroups v2: background and introduction**
- 7 Cgroups v2: enabling and disabling controllers
- 8 Cgroups v2: organizing cgroups and processes

- Designed to address perceived problems with cgroups v1
  - Section “R” of Documentation/cgroup-v2.txt details the problems
- Cgroups v2 officially released in Linux 4.5 (March 2016)
  - After extended experimental development phase...
- Both cgroups v1 and cgroups v2 can be used on same system
  - **But** can't mount same controller in both filesystems



- V2 currently implements only a subset of equivalents of v1 controllers
  - Work in progress...
- Documentation/cgroup-v2.txt documents v2 controllers
  - memory: control distribution of memory
    - Successor of v1 memory controller
  - io: regulate distribution of I/O resources
    - Successor of v1 blkio controller
  - pids: control number of processes
    - Exactly the same as v1 pids controller
  - ⚠ cpu: documented in Documentation/cgroup-v2.txt, but not yet merged (as at Linux 4.8)
  - freezer: work in progress (late 2016/early 2017?)

- **V1 hierarchy scheme** was supposed to allow great flexibility
  - V1: arbitrary number of hierarchies, with each hierarchy hosting any number of controllers
- But, that flexibility was **less useful than originally envisaged**

## Problems with the v1 hierarchy scheme:

- ☹ Utility controllers (e.g., freezer) that might be useful in all hierarchies could be used in only one
- ☹ Controllers bound to same hierarchy were forced to have same hierarchical view
  - Could not vary granularity according to controller
- These problems meant apps commonly put most controllers on separate, but highly similar, hierarchies
  - ☹ Same hierarchical management operations needed to be repeated on multiple hierarchies
  - ☹ Cooperation between controllers becomes complex
- ⇒ **v2 uses single hierarchy for all controllers**
  - Establish common resource domain across different resource types, so controllers (e.g., memory and io) can cooperate

Allowing **thread-granularity** for cgroup membership proved problematic

- Didn't make sense for some controllers
  - E.g., memory controller (all threads share memory...)
- Writing TIDs to tasks file is a **system-level activity**, but only **applications** well understand their thread topology
- $\Rightarrow$  **v2 allows only process-granularity** membership

- There may yet be some backtracking on process-vs-thread granularity for cpu controller
  - Some users are pushing back strongly for thread granularity
- Further info
  - “Resource groups”; <https://lwn.net/Articles/656115/>  
<https://lwn.net/Articles/679940/>  
<https://lkml.org/lkml/2016/1/5/366>  
  
<https://lwn.net/Articles/697369/> (“[Documentation] State of CPU controller in cgroup v2”, Aug 2016)  
<https://lwn.net/Articles/697366/> “The case of the stalled CPU controller”

- Allowing a cgroup to contain both tasks and child cgroups is problematic
  - Two different types of entities—*tasks* and *groups* of tasks—compete for distribution of same resources
    - Different controllers dealt with this in differing ways...
    - which could cause difficulties if trying to generically combine multiple controllers on same hierarchy
  - ⇒ **In v2, only leaf cgroups can contain processes**
    - (The story is a little more subtle...)

- **Inconsistencies** between controllers (“design followed implementation”)
  - In some hierarchies, new cgroups inherit parent’s attributes; in others, they get defaults
  - Some controllers have controller-specific interfaces in root cgroup; others don’t
  - v2: **consistent names and values** for interface files, **consistent inheritance rules** for all controllers
    - With some clearly documented guidelines!
- V1 cgroup release mechanism (firing up a process) has problems:
  - Firing up a process is expensive
  - Can’t delegate release handling to process inside a container
  - $\Rightarrow$  v2 has a lightweight solution that supports delegation

# Outline

---

- 1 Introduction
- 2 Cgroups v1: hierarchies and controllers
- 3 Cgroups v1: populating a cgroup
- 4 Cgroups v1: a survey of the controllers
- 5 Cgroups /proc files
- 6 Cgroups v2: background and introduction
- 7 Cgroups v2: enabling and disabling controllers**
- 8 Cgroups v2: organizing cgroups and processes



# Mounting the cgroups v2 filesystem

---

- To use cgroups v2, we mount new filesystem type:

```
# mount -t cgroup2 none /path/to/mount
```

- All **v2 controllers are automatically available** under single hierarchy
  - No need to explicitly bind controllers to mount point
    - cgroup2 filesystem allows/needs no `-o` mount options

# The cgroup.controllers file

---

- Each v2 cgroup has a cgroup.controllers file, which lists **available controllers** this cgroup can enable
- But, if we look in cgroups v2 root directory, we might find cgroup.controllers is empty:

```
# mkdir /mnt/cgroup2
# mount -t cgroup2 none /mnt/cgroup2
# cd /mnt/cgroup2
# cat cgroup.controllers
# wc -l cgroup.controllers
0 cgroup.controllers
```

# A v2 controller is available only if not mounted in v1

---

- V2 controller is available only if not bound in v1 hierarchy

```
# cat /proc/mounts | grep pids
cgroup /sys/fs/cgroup/pids cgroup rw,...,pids 0 0
pids
```

- That's why we didn't see pids in v2 cgroup.controllers
- ⇒ May need to unmount controller in v1 hierarchy to have it available in v2 hierarchy:

```
# umount /sys/fs/cgroup/pids
# cat cgroup.controllers
pids
```

- (Since Linux 4.6) kernel boot parameter, cgroup\_no\_v1:
  - cgroup\_no\_v1=all, to disable all v1 controllers
  - cgroup\_no\_v1=controller,..., to disable selected v1 controllers

# Enabling and disabling controllers

---

- Controllers are enabled/disabled by writing some subset of available controllers to `cgroup.subtree_control`

```
# echo "+pids -memory" > cgroup.subtree_control
```

- `+`  $\Rightarrow$  enable controller, `-`  $\Rightarrow$  disable controller

# Enabling and disabling controllers

---

- Enabling a controller in `cgroup.subtree_control`:
  - Allows resource to be **controlled in child cgroups**
  - **Creates controller-specific attribute files in each child directory**
- ⚠ ⚠ Attribute files in child cgroups are **used by process managing parent cgroup** to manage resource allocation across child cgroups
  - Different from v1...

# Example: enabling a controller

---

- In the cgroup root directory, list available controllers:

```
# cat cgroup.controllers
io memory pids
```

- Create a child cgroup; see what files are in subdirectory:

```
# mkdir grp1
# ls grp1
cgroup.controllers  cgroup.events  cgroup.procs
cgroup.subtree_control
```

- Enable pids controller for child cgroups; new control files have been created in child cgroup:

```
# echo '+pids' > cgroup.subtree_control
# ls grp1
cgroup.controllers  cgroup.procs      pids.current
cgroup.events       cgroup.subtree_control  pids.max
```

## Example: enabling a controller

---

- In grp1 cgroup, only available controller is pids:

```
# cat grp1/cgroup.controllers
pids
```

- In child of grp1, we can enable pids controller:

```
# mkdir grp1/sub
# echo '+pids' > grp1/cgroup.subtree_control
# cat grp1/cgroup.subtree_control
pids
```

- But io controller is not available:

```
# echo '+io' > grp1/cgroup.subtree_control
sh: echo: write error: No such file or directory
```

# Top-down constraints

---

- Child cgroups are always subject to any resource constraints established by controllers in ancestor cgroups
  - $\Rightarrow$  Descendant cgroups can't relax constraints imposed by ancestor cgroups
- If a controller is disabled in a cgroup (i.e., not written to `cgroup.subtree_control` in parent cgroup), it cannot be enabled in any descendants of the cgroup



- ① This exercise demonstrates that resource constraints apply in a top-down fashion, using the cgroups v2 pids controller.
  - Mount the cgroup2 filesystem if it is not already mounted and check that the pids controller is visible in the cgroup root `cgroup.controllers` file. If it is not, unmount the cgroup v1 pids filesystem. (See the steps at the start of this section.)
    - In some cases, unmounting the cgroup v1 pids filesystem may not be enough, since the controller is in use (e.g., by *systemd*). Therefore, it may be necessary to reboot the system with the `cgroup_no_v1=pids` kernel boot parameter.
  - To simplify the following steps, change your current directory to the cgroup root directory (i.e., the location where the cgroup2 filesystem is mounted).

# Exercises

---

- Create a child and grandchild directory in the cgroup filesystem and enable the PIDs controller in the root directory and the first subdirectory:

```
# mkdir xxx
# mkdir xxx/yyy
# echo '+pids' > cgroup.subtree_control
# echo '+pids' > xxx/cgroup.subtree_control
```

- Set an upper limit of 10 tasks in the child cgroup, and an upper limit of 20 tasks in the grandchild cgroup:

```
# echo '10' > xxx/pids.max
# echo '20' > xxx/yyy/pids.max
```

- In another terminal, use the supplied `cgroups/fork_bomb.c` program with the following command line, which will cause the program to first sleep 60 seconds and then create 30 children:

# Exercises

---

```
$ ./fork_bomb 30 60
```

- The parent process in the `fork_bomb` program prints its PID before sleeping. While it is sleeping, return to the first terminal and place the parent process in the grandchild `pids` cgroup:

```
# echo parent-PID > xxx/yyy/cgroup.procs
```

- When the parent finishes sleeping, how many children does it successfully create?

# Outline

---

- 1 Introduction
- 2 Cgroups v1: hierarchies and controllers
- 3 Cgroups v1: populating a cgroup
- 4 Cgroups v1: a survey of the controllers
- 5 Cgroups /proc files
- 6 Cgroups v2: background and introduction
- 7 Cgroups v2: enabling and disabling controllers
- 8 Cgroups v2: organizing cgroups and processes

# Organizing cgroups and processes

---

Broadly similar to cgroups v1:

- Hierarchy organized as set of subdirectories
- All processes initially in root cgroup
- Move process into group by writing PID into `cgroup.procs`
- Read `cgroup.procs` to discover process membership
  - ⚠ Returned list is not sorted
  - ⚠ List may contain duplicate PIDs
    - E.g., if PID moved out and then back into cgroup, or PID recycled, while reading
- Child of *fork()* inherits parent's cgroup membership
- Cgroup directory with no process members or child cgroups can be removed

# Organizing cgroups and processes

---

Differences between v1 and v2:

- Cgroup can't both control cgroup children and have member processes
  - $\Rightarrow$  Place member processes in leaf nodes
- No tasks file
  - Granularity for cgroup membership is process
  - Writing TID of any thread to `cgroup.procs` moves all of process's threads to cgroup
- Root cgroup does not contain controller interface files

## “Only leaf nodes can have member process”

---

- Earlier statement: cgroup can't have both child cgroups and member processes
- Let's refine that...
- A cgroup can't both:
  - distribute a resource to child cgroups, **and**
  - have child processes
  - (**Note:** root cgroup is an exception to this rule)
- Conversely (1):
  - A cgroup **can** have member processes and child cgroups...
  - **iff** it does not enable controllers for child cgroups
- Conversely (2):
  - If cgroup has child cgroups and processes, the processes must be moved elsewhere before enabling controllers
    - E.g., processes could be moved to child cgroups

# Cgroup (un)populated notification

---

- Cgroups v1: firing up a process is an expensive way of get notification of an empty cgroup!
- Cgroups v2: dispenses with `release_agent` and `notify_on_release` files
- Instead, each (non-root) cgroup has a file, `cgroup.events`, with a populated field:

```
# cat grp1/cgroup.events
populated 1
```

- 1 == subhierarchy contains live processes
  - I.e., live process in any descendant cgroup
- 0 == no live processes in subhierarchy



# Cgroup (un)populated notification

---

- Can monitor `cgroup.events` file, to get notification of transition between populated and unpopulated states
  - *inotify*: transitions generate `IN_MODIFY` events
  - *poll()*: transitions generate `POLLPRI` events
- One process can monitor multiple `cgroup.events` files
  - Much cheaper notification!
  - **Notification can be delegated** per container
    - I.e., one process can monitor all `cgroup.events` files in a subhierarchy

For the following exercises, you'll need to mount the `cgroup2` filesystem if it is not already mounted and check that the `pids` controller is visible in the `cgroup` root `cgroup.controllers` file. If it is not, unmount the `cgroup v1 pids` filesystem. (Details can be found at the start of section *Cgroups v2: enabling and disabling controllers*.) In the exercises below, we assume that the `cgroup2` filesystem is mounted at `/mnt/cgroup2`.

- ① This exercise demonstrates what happens if we try to enable a controller in a `cgroup` that has member processes.
  - Under the `cgroup2` mount point, create a new `cgroup`, and enable the `pids` controller in the root `cgroup`:

```
# cd /mnt/cgroup2
# mkdir child
# echo '+pids' > cgroup.subtree_control
```

# Exercise

---

- Start a process running *sleep*, and place it into the child cgroup:

```
# sleep 1000 &  
# echo $! > child/cgroup.procs
```

- What happens if we now try to enable the pids controller in the child cgroup via the following command?

```
# echo '+pids' > child/cgroup.subtree_control
```

# Thanks!

mtk@man7.org    @mkerrisk  
Slides at <http://man7.org/conf/>

Linux System Programming,  
System Programming for Linux Containers,  
and other training at <http://man7.org/training/>  
The Linux Programming Interface, <http://man7.org/tlpi/>

