



KVM Architecture Overview

2015 Edition

Stefan Hajnoczi <stefanha@redhat.com>

Introducing KVM virtualization

KVM hypervisor runs virtual machines on Linux hosts

- Mature on x86, recent progress on ARM and ppc

Most popular and best supported hypervisor on
OpenStack

- <https://wiki.openstack.org/wiki/HypervisorSupportMatrix>

Built in to Red Hat Enterprise Linux


- Qumranet startup created KVM, joined Red Hat in 2008



Virtualization goals

Efficiently and securely running virtual machines on a Linux host

- Linux, Windows, etc guest operating systems
- Access to networking and storage in a controlled fashion



Where does KVM fit into the stack?

Management for
datacenters and clouds

OpenStack

RHEV

Management for
one host

libvirt

Emulation for
one guest

QEMU

Guest

QMP

Host hardware access
and resource mgmt

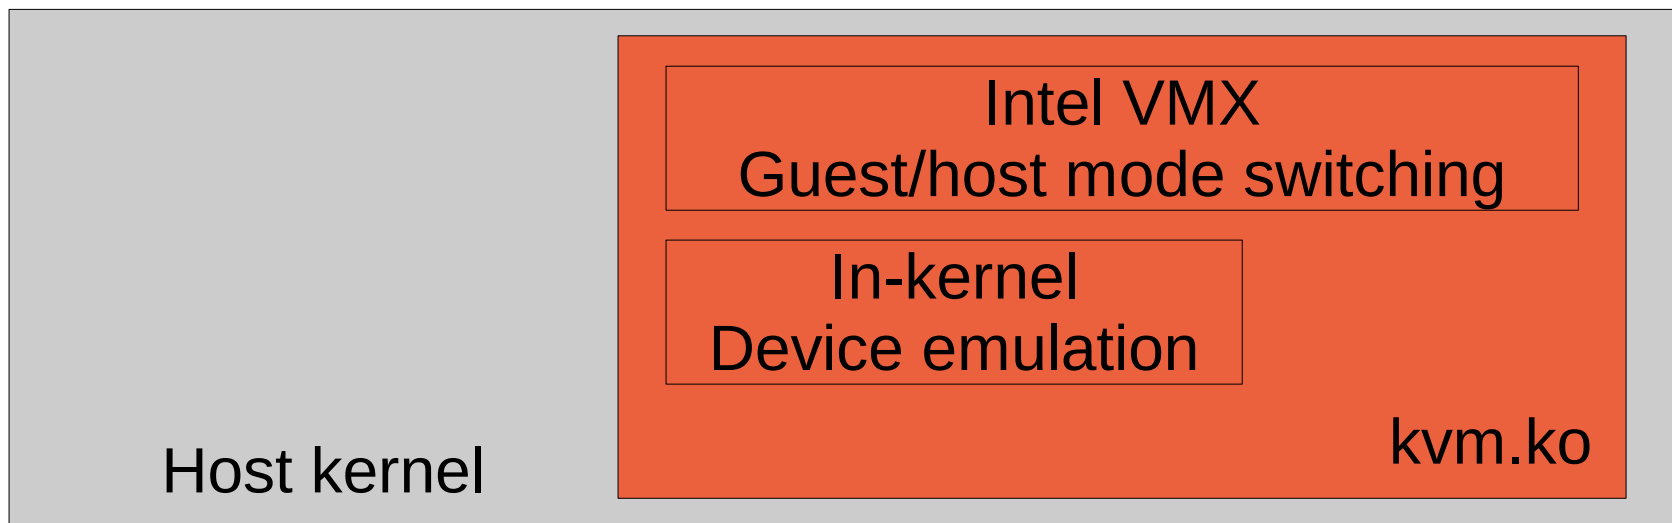

Host kernel

kvm.ko



More on QEMU and kvm.ko


Virtualization features



Hardware virtualization support with Intel **VMX**


Allows safe guest code execution at native speed

- Certain operations trap out to the hypervisor



Memory virtualization with Intel **EPT**

Extended Page Tables (EPT) add a level of address translation for guest physical memory.



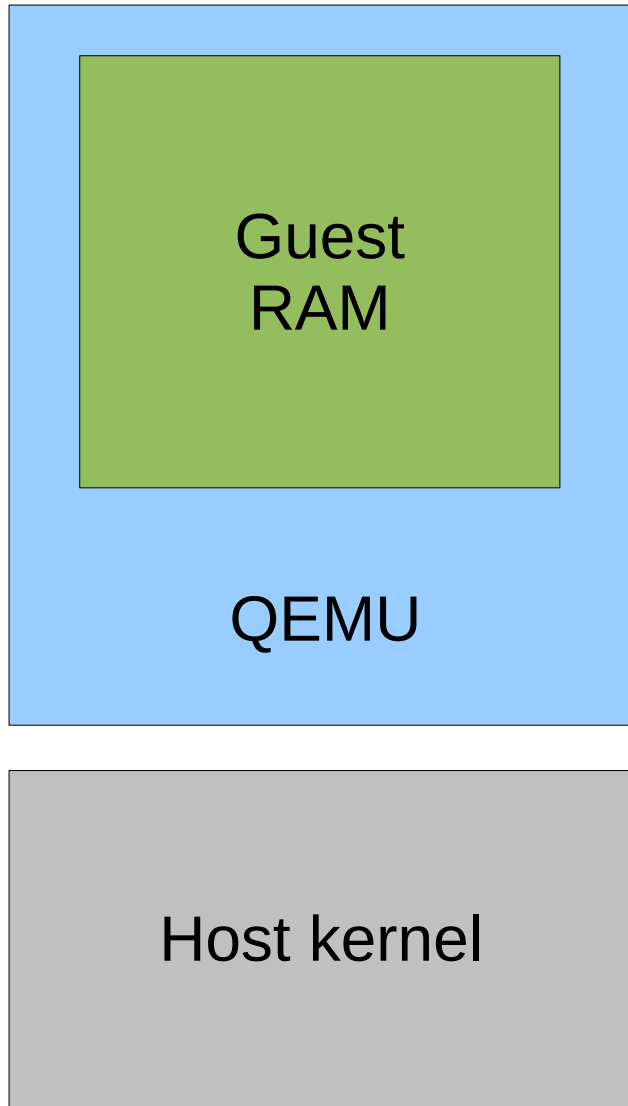
How QEMU uses kvm.ko

QEMU userspace process uses kvm.ko driver to execute guest code:

```
open( "/dev/kvm" )
ioctl(KVM_CREATE_VM)
ioctl(KVM_CREATE_VCPU)
for (;;) {
    ioctl(KVM_RUN)
    switch (exit_reason) {
    case KVM_EXIT_IO: /* ... */
    case KVM_EXIT_HLT: /* ... */
    }
}
```



QEMU process model



QEMU is a userspace process

Unprivileged and isolated using SELinux for security

Each KVM vCPU is a thread

Host kernel scheduler decides when vCPUs run



Linux concepts apply to QEMU/KVM

Since QEMU is a userspace process, the usual Linux tools work:

ps(1), top(1), etc see QEMU processes and threads

tcpdump(8) sees tap network traffic

blktrace(8) sees disk I/O requests

SystemTap and perf see QEMU activity

etc



Architecture: Event-driven multi-threaded

Event loops are used for timers, file descriptor monitoring, etc

- Non-blocking I/O
- Callbacks or coroutines

Multi-threaded architecture but with big lock

- **VCPU** threads execute in parallel
- Specific tasks that would block event loop are done in threads, e.g. remote display encoding, RAM live migration work, virtio-blk dataplane, etc
- Rest of QEMU code runs under global mutex



Architecture: Emulated and pass-through devices

Guest sees CPU, RAM, disk, etc like on real machines

- Unmodified operating systems can run
- Paravirtualized devices for better performance

Most devices are emulated and not real

- Isolation from host for security
- Sharing of resources between guests


Pass-through PCI adapters, disks, etc also possible

- Dedicated hardware



Architecture: Host/guest device emulation split

Guest device – device model visible to guest



Decouples
hardware
emulation from
I/O mechanism

Host device – performs I/O on behalf of guest



Architecture: virtio devices

KVM implements virtio device models

- net, blk, scsi, serial, rng, balloon
- See <http://docs.oasis-open.org/virtio/> for specs

Open standard for paravirtualized I/O devices

Red Hat contributes to Linux and Windows guest drivers



Architectural exception: **vhost** in-kernel devices

Most device emulation is best done in userspace

- Some APIs or performance features only available in host kernel

vhost drivers emulate virtio devices in host kernel


- vhost_net.ko high-performance virtio-net emulation takes advantage of kernel-only zero-copy and interrupt handling features
- Other devices could be developed in theory, but usually userspace is a better choice



Storage in QEMU

Block drivers fall in two categories:

Formats – image file formats (qcow2, vmdk, etc)

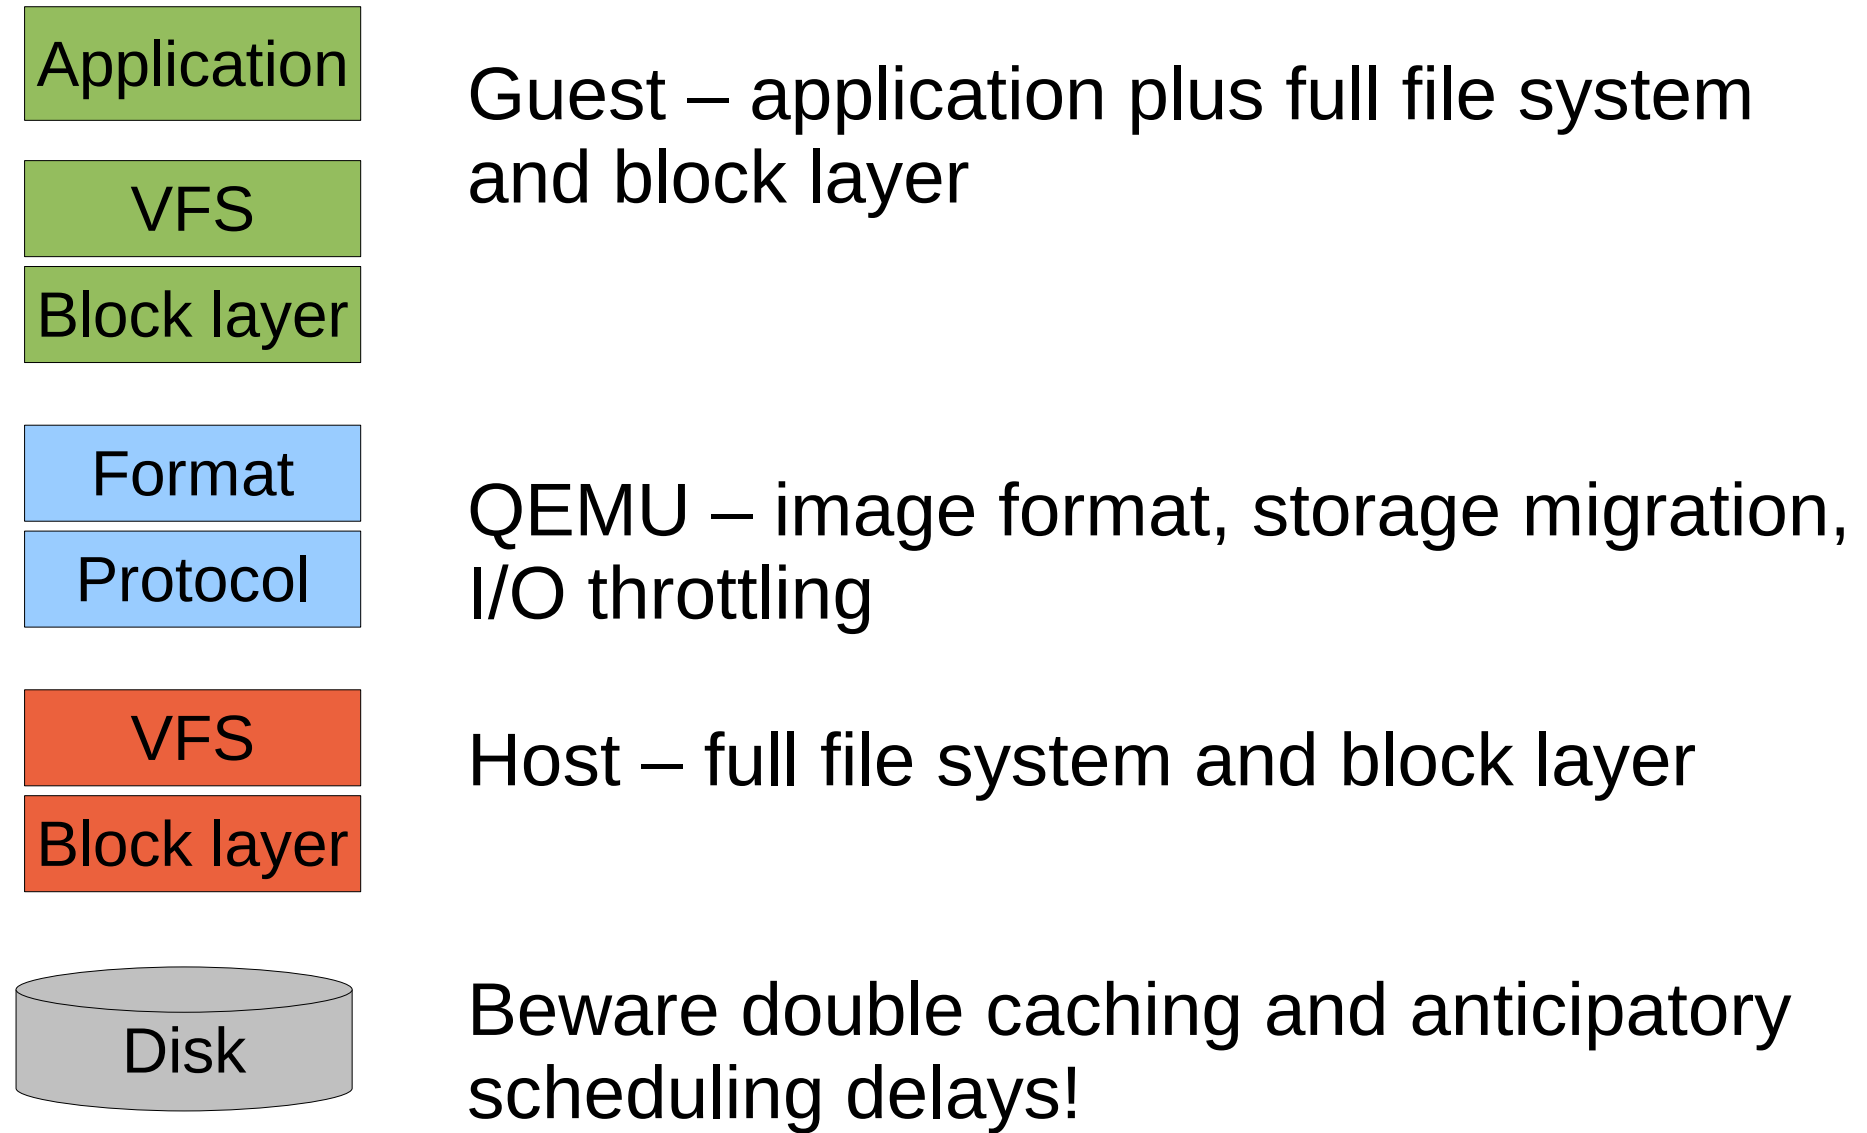


Protocols – I/O transports (POSIX file, rbd/Ceph, etc)

Plus additional block drivers that interpose like quorum, blkdebug, blkverify

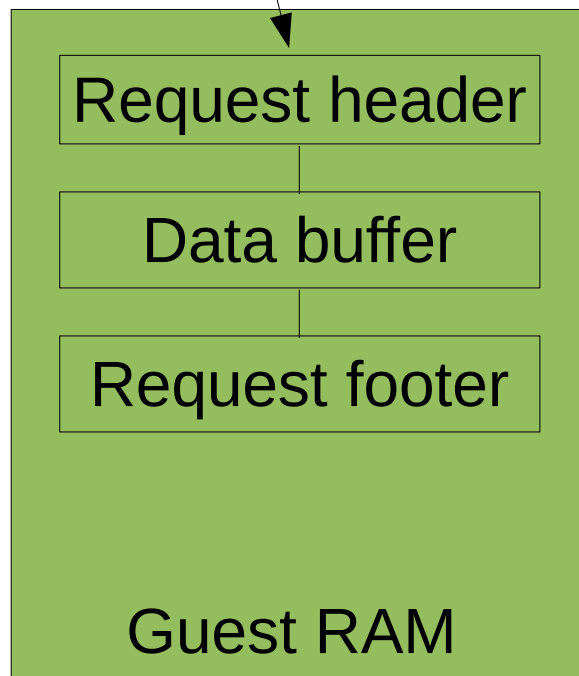


Storage stack




Walkthrough: virtio-blk disk read request (Part 1)

1. Guest fills in request descriptors

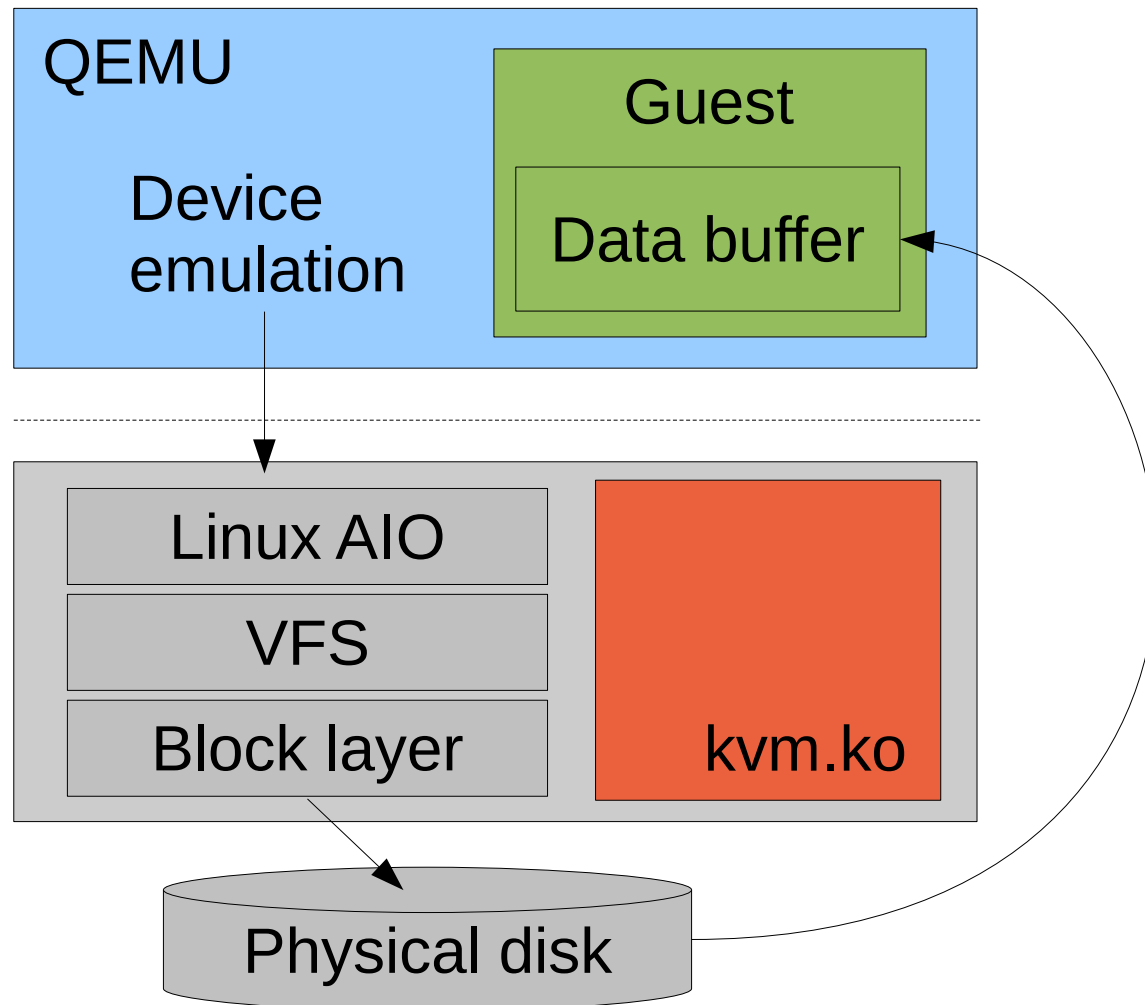


2. Guest writes to virtio-blk virtqueue notify register




Walkthrough: virtio-blk disk read request (Part 2)

3. QEMU issues I/O request on behalf of guest



Walkthrough: virtio-blk disk read request (Part 3)

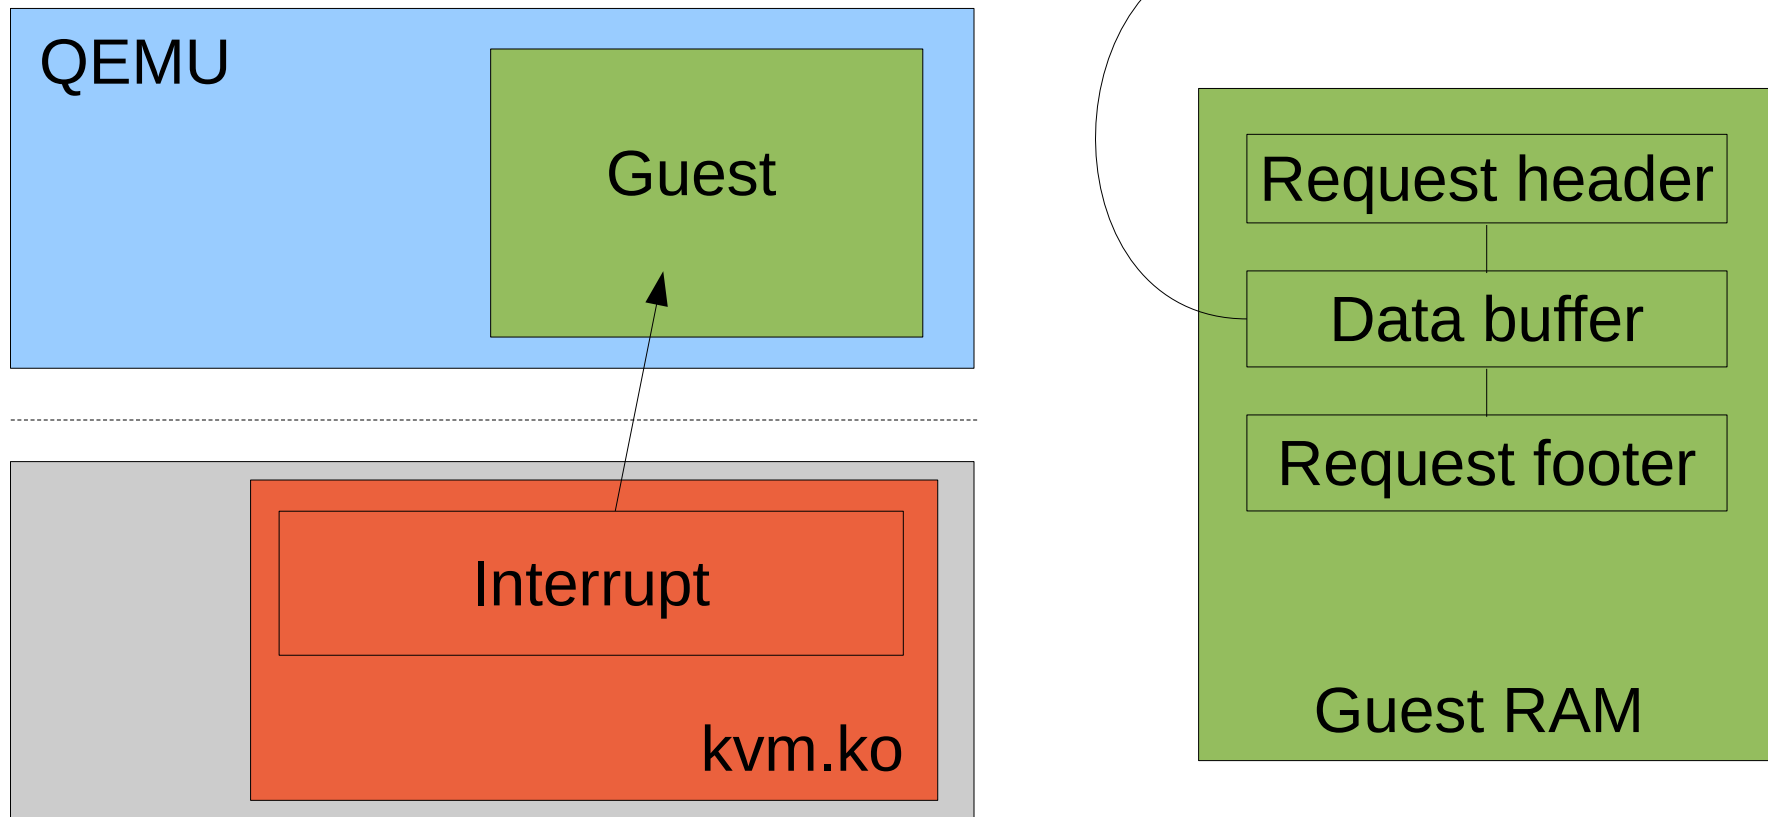
4. QEMU fills in request footer and injects completion interrupt



Walkthrough: virtio-blk disk read request (Part 4)

5. Guest receives interrupt and executes handler

6. Guest reads data from buffer



Thank you!

Technical discussion: qemu-devel@nongnu.org

IRC

- #qemu on irc.oftc.net
- #kvm on chat.freenode.net

<http://qemu-project.org/>

<http://linux-kvm.org/>

More on my blog: <http://blog.vmsplice.net/>

