

Test neutron network performance locally

Table of Contents

Linux bridge.....	2
Host to virtual machine	3
VM1 to VM2.....	4
OVS bridge with Linux bridge connected via veth.....	4
Host to virtual machine	6
VM1 to VM2.....	7
OVS bridge with Linux bridge connected via OVS internal port.....	8
Host to virtual machine	9
VM1 to VM2.....	10
OVS bridge	11
Host to virtual machine	12
VM1 to VM2.....	14
Linux bridge with OVS bridge connected by OVS internal port.....	14
Host to virtual machine	16
VM1 to VM2.....	17
Devstack test	18
Host to virtual machine	19
VM1 to VM2.....	20
Devstack test with OVS internal port	20
Host to virtual machine	21
VM1 to VM2.....	22
Conclusion	23

For a long time I want to do testing about the performance of the networks created by neutron. In particular, when I see the *nine* devices that the virtual machine's traffic must pass described by [Neutron administrator guide](#), I have a stronger motivation to do it. This post records the whole testing process I have done.

In the front parts of this post, Two virtual machines are launched with 'virsh' manually. They are attached to Linux bridges, OpenvSwitch bridges or hybrid bridges. The hybrid way is to simulate the nova's LibvirtHybridOVSBridgeDriver vif driver. After that, I test the LibvirtHybridOVSBridgeDriver driver with devstack. And then I come up with a new vif driver to avoid the performance of the LibvirtHybridOVSBridgeDriver.

The testing process is organized into seven different cases. Each test case has its own network fabric. I test two kind of traffics in each test case. One is from host to virtual machine, the other is from a virtual machine to the other virtual machine. The testing is done with iperf tool, and the data is collected. The result is shown at the end of this post.

Linux bridge

In this network fabric, the virtual machines are attached to one Linux bridge:


```
$ sudo brctl show
bridge name bridge id          STP enabled interfaces
br100          8000.fe5400a04c09  no          vnet0

vnet1
```

The virtual machine's interface definition in its domain file is like:

```
<interface type='bridge'>
  <mac address='52:54:00:a0:4c:09'>
  <source bridge='br100'>
  <model type='virtio'>
  <address type='pci' domain='0x0000' bus='0x00' slot='0x03' function='0x0'>
</interface>
```

The network fabric is shown as following figure:



Host to virtual machine

Start iperf as server on virtual machine VM1 and wait for connection, the iperf will print the bandwidth after testing:

```
# iperf -s
```

Server listening on TCP port 5001

TCP window size: 85.3 KByte (default)

[4] local 10.0.1.2 port 5001 connected with 10.0.1.1 port 58596

[ID]	Interval	Transfer	Bandwidth
-------	----------	----------	-----------

[4]	0.0-30.0 sec	95.6 GBytes	27.4 Gbits/sec
------	--------------	-------------	----------------

Start iperf as client to connect to virtual machine one. the iperf will run for 30 seconds to push data to server and print the bandwidth:

```
# iperf -c 10.0.1.2 -i 2 -t 30
```

Client connecting to 10.0.1.2, TCP port 5001

TCP window size: 16.0 KByte (default)

[3] local 10.0.1.1 port 58596 connected with 10.0.1.2 port 5001

[ID]	Interval	Transfer	Bandwidth
-------	----------	----------	-----------

[3]	0.0- 2.0 sec	6.00 GBytes	25.8 Gbits/sec
------	--------------	-------------	----------------

[3]	2.0- 4.0 sec	6.24 GBytes	26.8 Gbits/sec
------	--------------	-------------	----------------

[3]	4.0- 6.0 sec	6.85 GBytes	29.4 Gbits/sec
------	--------------	-------------	----------------

[3]	6.0- 8.0 sec	6.69 GBytes	28.7 Gbits/sec
------	--------------	-------------	----------------

```
[ 3] 8.0-10.0 sec 6.24 GBytes 26.8 Gbits/sec
[ 3] 10.0-12.0 sec 6.78 GBytes 29.1 Gbits/sec
```

VM1 to VM2

```
vm1# iperf -c 10.0.1.3 -i 2 -t 30
```

```
-----
Client connecting to 10.0.1.3, TCP port 5001
```

```
TCP window size: 23.5 KByte (default)
-----
```

```
[ 3] local 10.0.1.2 port 40280 connected with 10.0.1.3 port 5001
```

[ID]	Interval	Transfer	Bandwidth
[3]	0.0- 2.0 sec	5.62 GBytes	24.1 Gbits/sec
[3]	2.0- 4.0 sec	5.51 GBytes	23.7 Gbits/sec
[3]	4.0- 6.0 sec	5.89 GBytes	25.3 Gbits/sec
[3]	6.0- 8.0 sec	6.12 GBytes	26.3 Gbits/sec

```
vm2#iperf -s
```

```
-----
Server listening on TCP port 5001
```

```
TCP window size: 85.3 KByte (default)
-----
```

```
[ 4] local 10.0.1.3 port 5001 connected with 10.0.1.2 port 40280
```

[ID]	Interval	Transfer	Bandwidth
[4]	0.0-30.0 sec	83.6 GBytes	23.9 Gbits/sec

OVS bridge with Linux bridge connected via veth

In this network fabric, the virtual machines are attached to one Linux bridge irrespectively:

```
$ sudo brctl show
```

bridge name	bridge id	STP enabled	interfaces
br1	8000.e2e3ea921a4b	no	qvb1 vnet0
br2	8000.daede0593691	no	qvb2 vnet1

And then the Linux bridges are connected to a OVS bridge via veth devices qvb1 ↔ qvo1 and qvb2 ↔ qvo2:

```
$ sudo ovs-vsctl show
```

```
e5e2f9b3-a938-48cf-8754-6c70f7e251bc
```

```
Bridge br-int-test
```

```
    Port "qvo2"
```

```
        Interface "qvo2"
```

```
    Port br-int-test
```

```
        Interface br-int-test
```

```
            type: internal
```

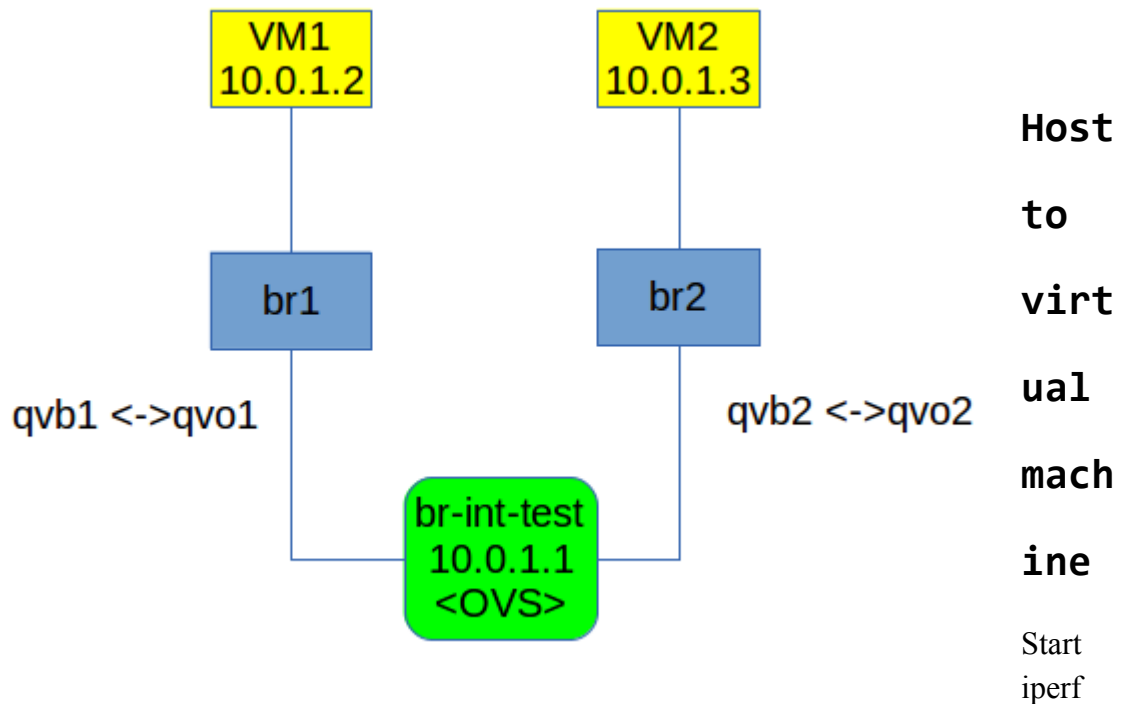
```
    Port "qvo1"
```

```
        Interface "qvo1"
```

The virtual machine's interface definition in its domain file is like:

```
<interface type='bridge'>
  <mac address='52:54:00:a0:4c:09'/>
  <source bridge='br1'/>
  <model type='virtio'/>
  <address type='pci' domain='0x0000' bus='0x00' slot='0x03' function='0x0'/>
</interface>
```

The network fabric is like this:



server on virtual machine and wait for connection:

```
# iperf -s
```

```
-----  
Server listening on TCP port 5001
```

```
TCP window size: 85.3 KByte (default)  
-----
```

```
[ 4] local 10.0.1.3 port 5001 connected with 10.0.1.1 port 38999
```

```
[ 4]  0.0-30.0 sec   3.94 GBytes   1.13 Gbits/sec
```

Start the iperf client on host:

```
$ iperf -c 10.0.1.3 -i 2 -t 30
```

```
-----  
Client connecting to 10.0.1.3, TCP port 5001
```

```
TCP window size: 22.9 KByte (default)  
-----
```

```
[ 3] local 10.0.1.1 port 38999 connected with 10.0.1.3 port 5001
```

```
[ ID] Interval          Transfer    Bandwidth
```

```
[ 3]  0.0- 2.0 sec    286 MBytes  1.20 Gbits/sec
```

[3]	2.0- 4.0 sec	247 MBytes	1.03 Gbits/sec
[3]	4.0- 6.0 sec	253 MBytes	1.06 Gbits/sec
[3]	6.0- 8.0 sec	238 MBytes	1.00 Gbits/sec

VM1 to VM2:

Start iperf server on VM2:

```
# iperf -s
```

Server listening on TCP port 5001

TCP window size: 85.3 KByte (default)

[4] local 10.0.1.3 port 5001 connected with 10.0.1.2 port 58664

[ID]	Interval	Transfer	Bandwidth
[4]	0.0-30.0 sec	12.0 GBytes	3.43 Gbits/sec

Start iperf client on VM2:

```
# iperf -c 10.0.1.3 -i 2 -t 30
```

Client connecting to 10.0.1.3, TCP port 5001

TCP window size: 23.5 KByte (default)

[3] local 10.0.1.2 port 58664 connected with 10.0.1.3 port 5001

[ID]	Interval	Transfer	Bandwidth
[3]	0.0- 2.0 sec	662 MBytes	2.78 Gbits/sec
[3]	2.0- 4.0 sec	809 MBytes	3.39 Gbits/sec
[3]	4.0- 6.0 sec	792 MBytes	3.32 Gbits/sec
[3]	6.0- 8.0 sec	807 MBytes	3.38 Gbits/sec
[3]	8.0-10.0 sec	833 MBytes	3.50 Gbits/sec

```
[ 3] 10.0-12.0 sec    817 MBytes    3.43 Gbits/sec
[ 3] 12.0-14.0 sec    836 MBytes    3.51 Gbits/sec
[ 3] 14.0-16.0 sec    906 MBytes    3.80 Gbits/sec
```

OVS bridge with Linux bridge connected via OVS internal port

This fabric is like previous one. However, the Linux bridge is connected to OVS bridge via OVS internal port, instead of the veth devices.

```
# ovs-vsctl add-port br-int-test vm2 -- set interface vm2 type=internal
# ovs-vsctl add-port br-int-test vm1 -- set interface vm1 type=internal
# ovs-vsctl show
e5e2f9b3-a938-48cf-8754-6c70f7e251bc
```

```
    Bridge br-int-test
```

```
        Port br-int-test
```

```
            Interface br-int-test
```

```
                type: internal
```

```
        Port "vm1"
```

```
            Interface "vm1"
```

```
                type: internal
```

```
        Port "vm2"
```

```
            Interface "vm2"
```

```
                type: internal
```


```
    ovs_version: "1.11.0"
```

```
# brctl show
```

bridge name	bridge id	STP enabled	interfaces
br1	8000.62ad95cf4b15	no	vm1
			vnet0
br2	8000.fe5400f8fa26	no	vm2

vnet1

The fabric is like:



Host to virtual machine

Start iperf server on virtual machine VM2:

```
# iperf -s
```

Server listening on TCP port 5001

TCP window size: 85.3 KByte (default)

[4] local 10.0.1.3 port 5001 connected with 10.0.1.1 port 42314

[ID]	Interval	Transfer	Bandwidth
-------	----------	----------	-----------

[4]	0.0-30.0 sec	91.1 GBytes	26.1 Gbits/sec
------	--------------	-------------	----------------

and start iperf client on host:

```
$ iperf -c 10.0.1.3 -i 2 -t 30
```

Client connecting to 10.0.1.3, TCP port 5001

TCP window size: 22.9 KByte (default)

[3] local 10.0.1.1 port 42314 connected with 10.0.1.3 port 5001

[ID]	Interval	Transfer	Bandwidth
[3]	0.0- 2.0 sec	6.28 GBytes	27.0 Gbits/sec
[3]	2.0- 4.0 sec	5.68 GBytes	24.4 Gbits/sec
[3]	4.0- 6.0 sec	6.44 GBytes	27.7 Gbits/sec
[3]	6.0- 8.0 sec	5.51 GBytes	23.7 Gbits/sec
[3]	8.0-10.0 sec	5.94 GBytes	25.5 Gbits/sec
[3]	10.0-12.0 sec	5.76 GBytes	24.7 Gbits/sec
[3]	12.0-14.0 sec	6.25 GBytes	26.8 Gbits/sec
[3]	14.0-16.0 sec	6.13 GBytes	26.3 Gbits/sec
[3]	16.0-18.0 sec	6.25 GBytes	26.9 Gbits/sec
[3]	18.0-20.0 sec	6.53 GBytes	28.0 Gbits/sec

VM1 to VM2

Start iperf server on VM2:

iperf -s

Server listening on TCP port 5001

TCP window size: 85.3 KByte (default)

[4] local 10.0.1.3 port 5001 connected with 10.0.1.2 port 44556

[ID]	Interval	Transfer	Bandwidth
[4]	0.0-30.0 sec	74.0 GBytes	21.2 Gbits/sec

Start iperf client on VM1:

```
# iperf -c 10.0.1.3 -i 2 -t 30
```

```
-----  
Client connecting to 10.0.1.3, TCP port 5001
```

```
TCP window size: 23.5 KByte (default)  
-----
```

```
[  3] local 10.0.1.2 port 44556 connected with 10.0.1.3 port 5001
```

[ID]	Interval	Transfer	Bandwidth
[3]	0.0- 2.0 sec	4.73 GBytes	20.3 Gbits/sec
[3]	2.0- 4.0 sec	4.76 GBytes	20.5 Gbits/sec
[3]	4.0- 6.0 sec	4.79 GBytes	20.6 Gbits/sec
[3]	6.0- 8.0 sec	4.64 GBytes	19.9 Gbits/sec
[3]	8.0-10.0 sec	5.03 GBytes	21.6 Gbits/sec
[3]	10.0-12.0 sec	4.93 GBytes	21.2 Gbits/sec
[3]	12.0-14.0 sec	5.08 GBytes	21.8 Gbits/sec
[3]	14.0-16.0 sec	4.80 GBytes	20.6 Gbits/sec
[3]	16.0-18.0 sec	4.86 GBytes	20.9 Gbits/sec

OVS bridge

In this fabric, the virtual machines are connected to OVS bridge directly. The interface definition of VMs' domain files is like:

```
<interface type='bridge'>  
  <mac address='52:54:00:a0:4c:09'>  
  <source bridge='br-int-test'>  
  <virtualport type='openvswitch'>  
    <parameters interfaceid='6a7bbe68-f4fe-f8fb-cf09-f2bb933e005f'>  
  </virtualport>  
  <target dev='vnet0'>  
  <model type='virtio'>  
  <alias name='net0'>  
  <address type='pci' domain='0x0000' bus='0x00' slot='0x03' function='0x0'>
```

</interface>

When the virtual machines are started, we have an OVS bridge like the following:

```
# ovs-vsctl show
```

```
e5e2f9b3-a938-48cf-8754-6c70f7e251bc
```

```
Bridge "br-int-test"
```

```
Port "vnet0"
```

```
Interface "vnet0"
```

```
Port "vnet1"
```

```
Interface "vnet1"
```


```
Port "br-int-test"
```

```
Interface "br-int-test"
```

```
type: internal
```

```
ovs_version: "1.11.0"
```

They are connected like the following figure:



Host to virtual machine

On VM2, we run:

```
# iperf -s
```

```
-----  
Server listening on TCP port 5001
```

```
TCP window size: 85.3 KByte (default)  
-----
```

```
[  4] local 10.0.1.2 port 5001 connected with 10.0.1.1 port 58905
```

```
[ ID] Interval      Transfer    Bandwidth
```

```
[  4]  0.0-30.0 sec   102 GBytes  29.1 Gbits/sec
```

On host, we run:

```
$ iperf -c 10.0.1.2 -i 2 -t 30
```

```
-----  
Client connecting to 10.0.1.2, TCP port 5001
```

```
TCP window size: 22.9 KByte (default)  
-----
```

```
[  3] local 10.0.1.1 port 58905 connected with 10.0.1.2 port 5001
```

```
[ ID] Interval      Transfer    Bandwidth
```

```
[  3]  0.0- 2.0 sec   6.90 GBytes  29.6 Gbits/sec
```

```
[  3]  2.0- 4.0 sec   6.46 GBytes  27.8 Gbits/sec
```

```
[  3]  4.0- 6.0 sec   6.74 GBytes  29.0 Gbits/sec
```

```
[  3]  6.0- 8.0 sec   7.18 GBytes  30.8 Gbits/sec
```

```
[  3]  8.0-10.0 sec   6.61 GBytes  28.4 Gbits/sec
```

```
[  3] 10.0-12.0 sec   7.13 GBytes  30.6 Gbits/sec
```

```
[  3] 12.0-14.0 sec   6.81 GBytes  29.2 Gbits/sec
```

```
[  3] 14.0-16.0 sec   5.95 GBytes  25.5 Gbits/sec
```

```
[  3] 16.0-18.0 sec   7.03 GBytes  30.2 Gbits/sec
```

```
[  3] 18.0-20.0 sec   7.14 GBytes  30.7 Gbits/sec
```

```
[  3] 20.0-22.0 sec   6.94 GBytes  29.8 Gbits/sec
```

```
[  3] 22.0-24.0 sec   7.12 GBytes  30.6 Gbits/sec
```

VM1 to VM2

On VM2:

```
# iperf -s
```

```
-----  
Server listening on TCP port 5001
```

```
TCP window size: 85.3 KByte (default)  
-----
```

```
[  4] local 10.0.1.3 port 5001 connected with 10.0.1.2 port 34248
```

[ID]	Interval	Transfer	Bandwidth
-------	----------	----------	-----------

[4]	0.0-30.0 sec	80.7 GBytes	23.1 Gbits/sec
-------	--------------	-------------	----------------

On VM1:

```
# iperf -c 10.0.1.3 -i 2 -t 30
```

```
-----  
Client connecting to 10.0.1.3, TCP port 5001
```

```
TCP window size: 23.5 KByte (default)  
-----
```

```
[  3] local 10.0.1.2 port 34248 connected with 10.0.1.3 port 5001
```

[ID]	Interval	Transfer	Bandwidth
-------	----------	----------	-----------

[3]	0.0- 2.0 sec	5.37 GBytes	23.1 Gbits/sec
-------	--------------	-------------	----------------

[3]	2.0- 4.0 sec	5.87 GBytes	25.2 Gbits/sec
-------	--------------	-------------	----------------

[3]	4.0- 6.0 sec	5.22 GBytes	22.4 Gbits/sec
-------	--------------	-------------	----------------

[3]	6.0- 8.0 sec	4.79 GBytes	20.6 Gbits/sec
-------	--------------	-------------	----------------

[3]	8.0-10.0 sec	5.29 GBytes	22.7 Gbits/sec
-------	--------------	-------------	----------------

[3]	10.0-12.0 sec	5.80 GBytes	24.9 Gbits/sec
-------	---------------	-------------	----------------

[3]	12.0-14.0 sec	5.65 GBytes	24.3 Gbits/sec
-------	---------------	-------------	----------------

Linux bridge with OVS bridge connected by OVS

internal port

This fabric connects each virtual machine to one OVS bridge irrespectively, and then add each OVS bridge's internal port into a Linux bridge.

The interface definition of virtual machine is like:

```
<interface type='bridge'>
  <mac address='52:54:00:a0:4c:09'/>
  <source bridge='br-int-1'/>
  <virtualport type='openvswitch'>
    <parameters interfaceid='6a7bbe68-f4fe-f8fb-cf09-f2bb933e005f'/>
  </virtualport>
  <model type='virtio'/>
  <address type='pci' domain='0x0000' bus='0x00' slot='0x03' function='0x0'/>
</interface>
```

And the interfaces of OVS and Linux bridge are like this:

```
# ovs-vsctl show
```

```
e5e2f9b3-a938-48cf-8754-6c70f7e251bc
```

```
Bridge "br-int-2"
```

```
Port "vnet1"
```

```
Interface "vnet1"
```

```
Port "br-int-2"
```

```
Interface "br-int-2"
```

```
type: internal
```

```
Bridge "br-int-1"
```

```
Port "br-int-1"
```

```
Interface "br-int-1"
```

```
type: internal
```

```
Port "vnet0"
```


```
Interface "vnet0"
```

```
ovs_version: "1.11.0"
```

```
# brctl show
```

bridge name	bridge id	STP enabled	interfaces
br100	8000.a236a5523a46	no	br-int-1 br-int-2

This network structure is shown as following:



Host to virtual machine

Run on VM2:

```
# iperf -s
```

Server listening on TCP port 5001

TCP window size: 85.3 KByte (default)

[4] local 10.0.1.3 port 5001 connected with 10.0.1.1 port 42231

[4] 0.0-30.0 sec 90.2 GBytes 25.8 Gbits/sec

Run on host:

```
$ iperf -c 10.0.1.3 -i 2 -t 30
```

Client connecting to 10.0.1.3, TCP port 5001

TCP window size: 22.9 KByte (default)

[3] local 10.0.1.1 port 42231 connected with 10.0.1.3 port 5001

[ID]	Interval	Transfer	Bandwidth
[3]	0.0- 2.0 sec	5.69 GBytes	24.4 Gbits/sec
[3]	2.0- 4.0 sec	6.24 GBytes	26.8 Gbits/sec
[3]	4.0- 6.0 sec	6.16 GBytes	26.5 Gbits/sec
[3]	6.0- 8.0 sec	5.99 GBytes	25.7 Gbits/sec

VM1 to VM2

Start iperf server on VM2:

```
# iperf -s
```

Server listening on TCP port 5001

TCP window size: 85.3 KByte (default)

[4] local 10.0.1.3 port 5001 connected with 10.0.1.2 port 42533

[ID]	Interval	Transfer	Bandwidth
[4]	0.0-30.0 sec	70.6 GBytes	20.2 Gbits/sec

Start client on VM1:

```
# iperf -c 10.0.1.3 -i 2 -t 30
```

Client connecting to 10.0.1.3, TCP port 5001

TCP window size: 23.5 KByte (default)

[3] local 10.0.1.2 port 42533 connected with 10.0.1.3 port 5001

[ID]	Interval	Transfer	Bandwidth
[3]	0.0- 2.0 sec	4.78 GBytes	20.5 Gbits/sec
[3]	2.0- 4.0 sec	4.56 GBytes	19.6 Gbits/sec
[3]	4.0- 6.0 sec	4.53 GBytes	19.4 Gbits/sec
[3]	6.0- 8.0 sec	5.04 GBytes	21.7 Gbits/sec
[3]	8.0-10.0 sec	4.74 GBytes	20.3 Gbits/sec
[3]	10.0-12.0 sec	4.65 GBytes	20.0 Gbits/sec
[3]	12.0-14.0 sec	4.81 GBytes	20.7 Gbits/sec
[3]	14.0-16.0 sec	4.90 GBytes	21.0 Gbits/sec
[3]	16.0-18.0 sec	4.75 GBytes	20.4 Gbits/sec

Devstack test

In this case, ML2 plugin and OVS agent is used. The fabric is similar to one in the case “OVS bridge with Linux bridge connected via veth”.

We customize the localrc:

```
disable_service n-net
enable_service q-svc
enable_service q-agt
enable_service q-dhcp
enable_service q-l3
enable_service q-meta
enable_service neutron
Q_PLUGIN=ml2
Q_AGENT=openvswitch
DATABASE_PASSWORD=root
RABBIT_PASSWORD=password
SERVICE_TOKEN=token
SERVICE_PASSWORD=password
ADMIN_PASSWORD=password
SCREEN_LOGDIR=$DEST/logs/screen
```

After running the stack.sh, run following commands to open ports of VMs:

```
$ neutron security-group-rule-create --protocol icmp --direction ingress default
$ neutron security-group-rule-create --protocol tcp --port-range-min 0 --port-range-max 65000
--direction ingress default
```

Host to virtual machine

Run on VM2:

```
# iperf -s
-----
Server listening on TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[  4] local 10.0.0.4 port 5001 connected with 10.0.0.1 port 49297
[ ID] Interval      Transfer    Bandwidth
[  4]  0.0-30.0 sec  3.53 GBytes  1.01 Gbits/sec
```

Run iperf client on host's router namespace:

```
$ sudo ip netns exec qrouter-63b8359f-a774-4e2d-83ff-5f0172644b2b iperf -c 10.0.0.4 -i 2 -t
30
-----
Client connecting to 10.0.0.4, TCP port 5001
TCP window size: 22.9 KByte (default)
-----
[  5] local 10.0.0.1 port 49297 connected with 10.0.0.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[  5]  0.0- 2.0 sec   224 MBytes   941 Mbits/sec
[  5]  2.0- 4.0 sec   218 MBytes   916 Mbits/sec
[  5]  4.0- 6.0 sec   228 MBytes   957 Mbits/sec
[  5]  6.0- 8.0 sec   235 MBytes   985 Mbits/sec
[  5]  8.0-10.0 sec   260 MBytes  1.09 Gbits/sec
[  5] 10.0-12.0 sec   280 MBytes  1.18 Gbits/sec
[  5] 12.0-14.0 sec   226 MBytes   946 Mbits/sec
[  5] 14.0-16.0 sec   243 MBytes  1.02 Gbits/sec
[  5] 16.0-18.0 sec   238 MBytes   999 Mbits/sec
[  5] 18.0-20.0 sec   247 MBytes  1.03 Gbits/sec
```

```
[ 5] 20.0-22.0 sec    250 MBytes    1.05 Gbits/sec
```

VM1 to VM2

Run iperf server on VM2:

```
# iperf -s
```

```
-----  
Server listening on TCP port 5001  
TCP window size: 85.3 KByte (default)  
-----
```

```
[ 4] local 10.0.0.3 port 5001 connected with 10.0.0.4 port 37903  
[ ID] Interval      Transfer    Bandwidth  
[ 4]  0.0-30.0 sec  7.85 GBytes 2.25 Gbits/sec
```

Run iperf client on VM1:

```
# iperf -c 10.0.0.3 -i 2 -t 30
```

```
-----  
Client connecting to 10.0.0.3, TCP port 5001  
TCP window size: 23.5 KByte (default)  
-----
```

```
[ 3] local 10.0.0.4 port 37903 connected with 10.0.0.3 port 5001  
[ ID] Interval      Transfer    Bandwidth  
[ 3]  0.0- 2.0 sec   548 MBytes 2.30 Gbits/sec  
[ 3]  2.0- 4.0 sec   542 MBytes 2.27 Gbits/sec  
[ 3]  4.0- 6.0 sec   536 MBytes 2.25 Gbits/sec  
[ 3]  6.0- 8.0 sec   516 MBytes 2.16 Gbits/sec  
[ 3]  8.0-10.0 sec   552 MBytes 2.32 Gbits/sec  
[ 3] 10.0-12.0 sec   510 MBytes 2.14 Gbits/sec  
[ 3] 12.0-14.0 sec   530 MBytes 2.22 Gbits/se
```

Devstack test with OVS internal port

It is obvious the veth connection solution is bad. So I tried to write a new vif driver named LibvirtHybridOVSBridgeDriverViaInternalPort which uses the OVS internal port as connection between Linux bridge and OVS bridge.

Add the following line in the localrc file used by previous case:

NOVA_VIF_DRIVER=nova.virt.libvirt.vif.LibvirtHybridOVSBridgeDriverViaInternalPort

This case's network structure is the same to the one in the case “OVS bridge with Linux bridge connected via OVS internal port”.

Host to virtual machine

On virtual machine, run iperf server:

```
# iperf -s
```

```
-----  
Server listening on TCP port 5001
```

```
TCP window size: 85.3 KByte (default)  
-----
```

```
[  4] local 10.0.0.5 port 5001 connected with 10.0.0.1 port 57274
```

```
[ ID] Interval      Transfer    Bandwidth
```

```
[  4]  0.0-30.0 sec  95.4 GBytes 27.3 Gbits/sec
```

In the host's route namespace run:

```
$ sudo ip netns exec qrouter-1021848e-d727-4085-8883-de3ddca677eb iperf -c 10.0.0.5 -i 2 -t 30
```

```
-----  
Client connecting to 10.0.0.5, TCP port 5001
```

```
TCP window size: 22.9 KByte (default)  
-----
```

```
[  5] local 10.0.0.1 port 57274 connected with 10.0.0.5 port 5001
```

```
[ ID] Interval      Transfer    Bandwidth
```

```
[  5]  0.0- 2.0 sec  6.83 GBytes 29.3 Gbits/sec
```

```
[  5]  2.0- 4.0 sec  6.14 GBytes 26.4 Gbits/sec
```

```
[  5]  4.0- 6.0 sec  6.11 GBytes 26.3 Gbits/sec
```

```
[  5]  6.0- 8.0 sec  6.76 GBytes 29.0 Gbits/sec
```

```
[  5]  8.0-10.0 sec  6.29 GBytes 27.0 Gbits/sec
```

```
[  5] 10.0-12.0 sec  6.88 GBytes 29.5 Gbits/sec
```

VM1 to VM2

On VM2 run:

```
# iperf -s
```

```
-----  
Server listening on TCP port 5001
```

```
TCP window size: 85.3 KByte (default)  
-----
```

```
[  4] local 10.0.0.5 port 5001 connected with 10.0.0.3 port 49414
```

[ID]	Interval	Transfer	Bandwidth
[4]	0.0-30.0 sec	69.2 GBytes	19.8 Gbits/sec

On VM1 run:

```
# iperf -c 10.0.0.5 -i 2 -t 30
```

```
-----  
Client connecting to 10.0.0.5, TCP port 5001
```

```
TCP window size: 23.5 KByte (default)  
-----
```


```
[  3] local 10.0.0.3 port 49414 connected with 10.0.0.5 port 5001
```

[ID]	Interval	Transfer	Bandwidth
[3]	0.0- 2.0 sec	4.15 GBytes	17.8 Gbits/sec
[3]	2.0- 4.0 sec	4.63 GBytes	19.9 Gbits/sec
[3]	4.0- 6.0 sec	4.48 GBytes	19.3 Gbits/sec
[3]	6.0- 8.0 sec	5.05 GBytes	21.7 Gbits/sec
[3]	8.0-10.0 sec	4.98 GBytes	21.4 Gbits/sec
[3]	10.0-12.0 sec	4.94 GBytes	21.2 Gbits/sec

Conclusion

Collect the bandwidth data of all the “iperf -s”, we can have a table and a figure below.

		Bandwidth output of iperf -s	
		Host to VM2	VM1 to VM2
1	Linux bridge	27.4	23.9
2	OVS bridge with Linux bridge connected via veth	1.13	3.43
3	OVS bridge with Linux bridge connected via OVS port	26.1	21.2
4	OVS bridge	29.1	23.1
5	Linux bridge with OVS bridge connected via OVS port	25.8	20.2
6	devstack test with hybrid OVS driver	1.01	2.25
7	devstack test with hybrid OVS driver via OVS port	27.3	19.8



We can get these points:

1. The network bandwidth is not influenced much by the so-called *nine devices passage. We can prove it from the fact that the bandwidth of cases 3, 5 and 7 are near to the bandwidth of cases 1 and 4;
2. The veth pair is reducing the bandwidth a lot. The cases 2 and 6 are using veth, where the bandwidth is almost ten times worse than other solution;
3. The OVS bridge is as good as Linux bridge which can be seen by comparing case 1 & case 4.