



# **Appendix A**


Authors: John Hennessy & David Patterson




**Figure A.1 Operand locations for four instruction set architecture classes.** The arrows indicate whether the operand is an input or the result of the arithmetic-logical unit (ALU) operation, or both an input and result. Lighter shades indicate inputs, and the dark shade indicates the result. In (a), a Top Of Stack register (TOS) points to the top input operand, which is combined with the operand below. The first operand is removed from the stack, the result takes the place of the second operand, and TOS is updated to point to the result. All operands are implicit. In (b), the Accumulator is both an implicit input operand and a result. In (c), one input operand is a register, one is in memory, and the result goes to a register. All operands are registers in (d) and, like the stack architecture, can be transferred to memory only via separate instructions: push or pop for (a) and load or store for (d).




**Figure A.7 Summary of use of memory addressing modes (including immediates).** These major addressing modes account for all but a few percent (0% to 3%) of the memory accesses. Register modes, which are not counted, account for one-half of the operand references, while memory addressing modes (including immediate) account for the other half. Of course, the compiler affects what addressing modes are used; see Section A.8. The memory indirect mode on the VAX can use displacement, autoincrement, or autodecrement to form the initial memory address; in these programs, almost all the memory indirect references use displacement mode as the base. Displacement mode includes all displacement lengths (8, 16, and 32 bits). The PC-relative addressing modes, used almost exclusively for branches, are not included. Only the addressing modes with an average frequency of over 1% are shown.




**Figure A.8 Displacement values are widely distributed.** There are both a large number of small values and a fair number of large values. The wide distribution of displacement values is due to multiple storage areas for variables and different displacements to access them (see Section A.8) as well as the overall addressing scheme the compiler uses. The  $x$ -axis is  $\log_2$  of the displacement, that is, the size of a field needed to represent the magnitude of the displacement. Zero on the  $x$ -axis shows the percentage of displacements of value 0. The graph does not include the sign bit, which is heavily affected by the storage layout. Most displacements are positive, but a majority of the largest displacements (14+ bits) are negative. Since these data were collected on a computer with 16-bit displacements, they cannot tell us about longer displacements. These data were taken on the Alpha architecture with full optimization (see Section A.8) for SPEC CPU2000, showing the average of integer programs (CINT2000) and the average of floating-point programs (CFP2000).




**Figure A.9 About one-quarter of data transfers and ALU operations have an immediate operand.** The bottom bars show that integer programs use immediates in about one-fifth of the instructions, while floating-point programs use immediates in about one-sixth of the instructions. For loads, the load immediate instruction loads 16 bits into either half of a 32-bit register. Load immediates are not loads in a strict sense because they do not access memory. Occasionally a pair of load immediates is used to load a 32-bit constant, but this is rare. (For ALU operations, shifts by a constant amount are included as operations with immediate operands.) The programs and computer used to collect these statistics are the same as in Figure A.8.



**Figure A.10 The distribution of immediate values.** The  $x$ -axis shows the number of bits needed to represent the magnitude of an immediate value—0 means the immediate field value was 0. The majority of the immediate values are positive. About 20% were negative for CINT2000, and about 30% were negative for CFP2000. These measurements were taken on an Alpha, where the maximum immediate is 16 bits, for the same programs as in Figure A.8. A similar measurement on the VAX, which supported 32-bit immediates, showed that about 20% to 25% of immediates were longer than 16 bits. Thus, 16 bits would capture about 80% and 8 bits about 50%.




**Figure A.11 Distribution of data accesses by size for the benchmark programs.** The double-word data type is used for double-precision floating point in floating-point programs and for addresses, since the computer uses 64-bit addresses. On a 32-bit address computer the 64-bit addresses would be replaced by 32-bit addresses, and so almost all double-word accesses in integer programs would become single-word accesses.




**Figure A.14 Breakdown of control flow instructions into three classes: calls or returns, jumps, and conditional branches.** Conditional branches clearly dominate. Each type is counted in one of three bars. The programs and computer used to collect these statistics are the same as those in Figure A.8.





**Figure A.15 Branch distances in terms of number of instructions between the target and the branch instruction.** The most frequent branches in the integer programs are to targets that can be encoded in 4 to 8 bits. This result tells us that short displacement fields often suffice for branches and that the designer can gain some encoding density by having a shorter instruction with a smaller branch displacement. These measurements were taken on a load-store computer (Alpha architecture) with all instructions aligned on word boundaries. An architecture that requires fewer instructions for the same program, such as a VAX, would have shorter branch distances. However, the number of bits needed for the displacement may increase if the computer has variable-length instructions to be aligned on any byte boundary. The programs and computer used to collect these statistics are the same as those in Figure A.8.



**Figure A.17 Frequency of different types of compares in conditional branches.** Less than (or equal) branches dominate this combination of compiler and architecture. These measurements include both the integer and floating-point compares in branches. The programs and computer used to collect these statistics are the same as those in Figure A.8.

|                                  |                        |                    |     |                          |                      |
|----------------------------------|------------------------|--------------------|-----|--------------------------|----------------------|
| Operation and<br>no. of operands | Address<br>specifier 1 | Address<br>field 1 | ... | Address<br>specifier $n$ | Address<br>field $n$ |
|----------------------------------|------------------------|--------------------|-----|--------------------------|----------------------|

(a) Variable (e.g., Intel 80x86, VAX)

|           |                    |                    |                    |
|-----------|--------------------|--------------------|--------------------|
| Operation | Address<br>field 1 | Address<br>field 2 | Address<br>field 3 |
|-----------|--------------------|--------------------|--------------------|

(b) Fixed (e.g., Alpha, ARM, MIPS, PowerPC, SPARC, SuperH)

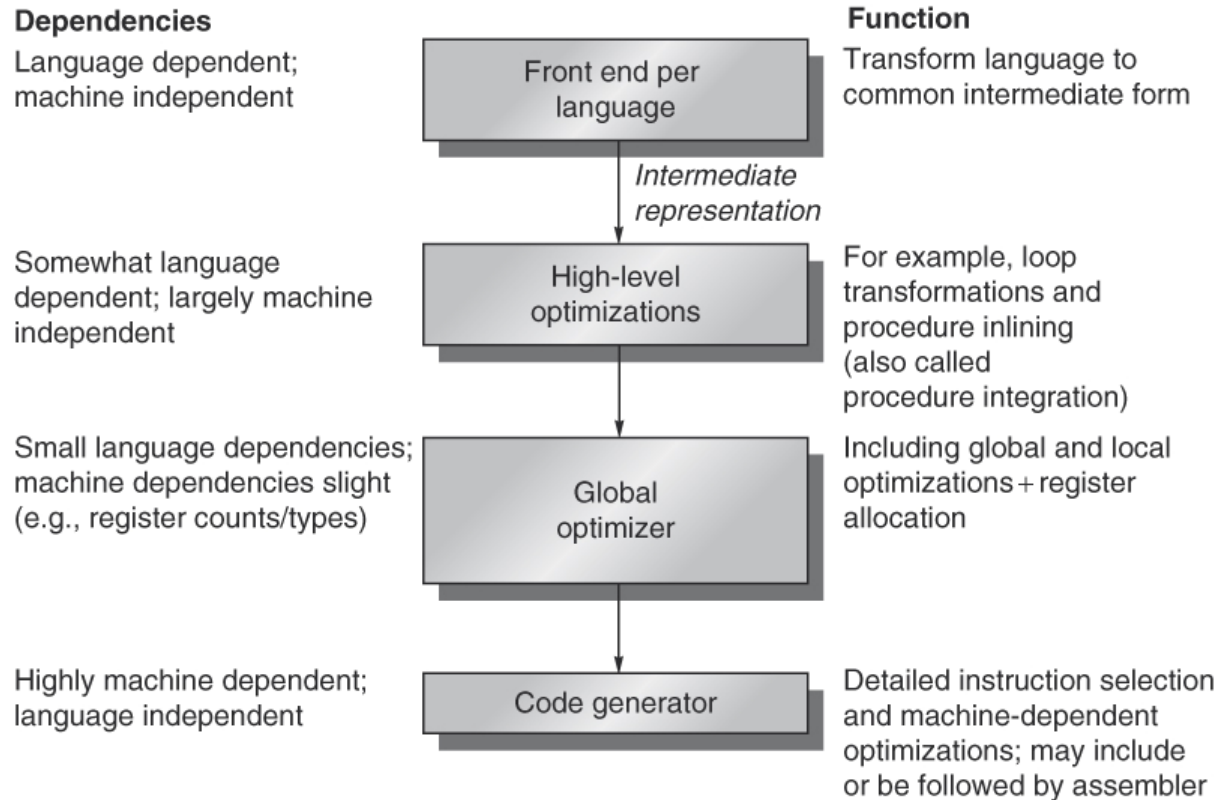
|           |                      |                  |
|-----------|----------------------|------------------|
| Operation | Address<br>specifier | Address<br>field |
|-----------|----------------------|------------------|

|           |                        |                        |                  |
|-----------|------------------------|------------------------|------------------|
| Operation | Address<br>specifier 1 | Address<br>specifier 2 | Address<br>field |
|-----------|------------------------|------------------------|------------------|


|           |                      |                    |                    |
|-----------|----------------------|--------------------|--------------------|
| Operation | Address<br>specifier | Address<br>field 1 | Address<br>field 2 |
|-----------|----------------------|--------------------|--------------------|

(c) Hybrid (e.g., IBM 360/370, MIPS16, Thumb, TI TMS320C54x)

**Figure A.18 Three basic variations in instruction encoding: variable length, fixed length, and hybrid.** The variable format can support any number of operands, with each address specifier determining the addressing mode and the length of the specifier for that operand. It generally enables the smallest code representation, since unused fields need not be included. The fixed format always has the same number of operands, with the addressing modes (if options exist) specified as part of the opcode. It generally results in the largest code size. Although the fields tend not to vary in their location, they will be used for different purposes by different instructions. The hybrid approach has multiple formats specified by the opcode, adding one or two fields to specify the addressing mode and one or two fields to specify the operand address.



**Figure A.19 Compilers typically consist of two to four passes, with more highly optimizing compilers having more passes.** This structure maximizes the probability that a program compiled at various levels of optimization will produce the same output when given the same input. The optimizing passes are designed to be optional and may be skipped when faster compilation is the goal and lower-quality code is acceptable. A *pass* is simply one phase in which the compiler reads and transforms the entire program. (The term *phase* is often used inter-changeably with *pass*.) Because the optimizing passes are separated, multiple languages can use the same optimizing and code generation passes. Only a new front end is required for a new language.



**Figure A.21 Change in instruction count for the programs lucas and mcf from the SPEC2000 as compiler optimization levels vary.** Level 0 is the same as unoptimized code. Level 1 includes local optimizations, code scheduling, and local register allocation. Level 2 includes global optimizations, loop transformations (software pipelining), and global register allocation. Level 3 adds procedure integration. These experiments were performed on Alpha compilers.

### I-type instruction



Encodes: Loads and stores of bytes, half words, words, double words. All immediates ( $rt \leftarrow rs \text{ op immediate}$ )

Conditional branch instructions (rs is register, rd unused)

Jump register, jump and link register

(rd=0, rs=destination, immediate=0)

### R-type instruction



Register-register ALU operations:  $rd \leftarrow rs \text{ funct } rt$

Function encodes the data path operation: Add, Sub, . . .

Read/write special registers and moves


### J-type instruction



Jump and jump and link

Trap and return from exception

**Figure A.22 Instruction layout for MIPS.** All instructions are encoded in one of three types, with common fields in the same location in each format.



**Figure A.29 Data reference size of four programs from SPEC2000.** Although you can calculate an average size, it would be hard to claim the average is typical of programs.